BUS13-1

Marketing Fundamentals

Where to Play, How to Win, How to Allocate

Marketing is the discipline of choosing what not to do.

Albert School — 22 hours

Preface

You arrive at this course already fluent in marketing — or so the world has told you. You have seen ten thousand advertisements, scrolled past sponsored posts, noticed which logos feel premium and which feel cheap. It is tempting to conclude that marketing is the craft of producing those artefacts: the ad, the campaign, the clever tagline.

This course asks you to unlearn that. Marketing, as a serious business discipline, is not the production of artefacts. It is a sequence of **decisions** — and like all consequential decisions, the hardest part is not what you choose to do but what you deliberately refuse to do. A strategy that forbids nothing commits to nothing. A brand that serves everyone is salient to no one. A budget that funds every tactic optimises none.

We will organise the entire course around three irreversible decisions that every marketing strategy must make, whether consciously or by accident:

Where to Play **How to Win** **How to Allocate**

These come from the classic strategy doctrine of Roger Martin and A.G. Lafley (*Playing to Win*), and we will treat them not as a slogan but as a working diagnostic instrument. By the end you should be able to look at any business situation — a startup choosing its first customers, an incumbent under attack, a product launch — and say precisely which of these three decisions is actually at stake, and which are settled.

No formal prerequisites are assumed. What is assumed is intellectual honesty: the willingness to write down a choice, state its cost in money or attention, and defend why the thing you gave up was worth giving up. That habit — costing your trade-offs explicitly — is the spine of everything that follows.

Contents

Preface . . . 2
1 Marketing as a Decision Discipline . . . 4
1.1 A concrete situation first . . . 4
1.2 The three irreversible decisions . . . 4
1.3 Business strategy vs marketing strategy . . . 5
1.4 Marketing as demand-side allocation in the AI era . . . 6
2 Where to Play I — Segmentation as Strategic Hypothesis . . . 6
2.1 The personas-on-a-wall failure . . . 6
2.2 Demand heterogeneity: the right basis for cutting the market . . . 7
2.3 Jobs-to-be-done . . . 7
2.4 Stated vs revealed preferences . . . 8
3 Where to Play II — Targeting as Capital Allocation . . . 8
3.1 Choosing a segment is spending money you cannot get back . . . 8
3.2 Broad vs narrow: the central tension . . . 9
3.3 The Ehrenberg-Bass perspective . . . 9
3.4 Category entry points . . . 9
3.5 Defending the decision to ignore a segment . . . 10
4 How to Win I — Positioning as Constraint . . . 11
4.1 A position is a promise you let constrain you . . . 11
4.2 Competitive frames of reference . . . 11
4.3 The four stress tests . . . 11
4.4 Differentiation vs distinctiveness . . . 12

5 How to Win II — Brand as Strategic Asset 13
5.1 Why brand is an asset, not an expense 13
5.2 Mental availability vs brand awareness 13
5.3 Distinctive brand assets 14
5.4 The long and the short of it: Binet & Field 14
6 How to Allocate I — The Marketing Mix as a Trade-off Framework 16
6.1 The mix is a set of trade-offs, not a checklist 16
6.2 Allocating a fixed budget across the 4Ps 17
7 How to Allocate II — Pricing as Strategic Signal 17
7.1 Price is the loudest message you send 17
7.2 Three ways to set a price 18
7.3 Price as positioning signal 18
7.4 Trading off margin, volume, and signal 18
8 Competitive Landscapes — Where Not to Fight 19
8.1 Strategy is partly the art of avoidance 19
8.2 Direct and indirect competition 19
8.3 Writing a "where not to fight" recommendation from public data 20
9 Auditing and Quantifying the Marketing Mix 21
9.1 Coherence is the thing you audit for 21
9.2 Quantifying the impact: a simplified P&L 21
10 B2B vs B2C — Adapting the Doctrine 22
10.1 The doctrine is invariant; the application is not 22
10.2 Closing 23

1 Marketing as a Decision Discipline

1.1 A concrete situation first

Imagine two founders, Maya and Tom, who have each built a meal-kit delivery service. Their products are nearly identical: chilled boxes of pre-portioned ingredients with recipe cards. They have the same launch budget, €200,000.

Maya spends her budget the way most people imagine marketing works. She hires a designer, runs Instagram ads to “everyone who likes cooking”, offers a 50% first-box discount, and A/B-tests her subject lines. She is busy, her dashboards are full, and her cost-per-click is falling. Eighteen months later she is out of business.

Tom spends three weeks **before** touching a single ad. He decides his service is not for people who like cooking — it is for dual-income parents of young children who have abandoned cooking on weeknights out of exhaustion, not lack of skill. He decides he will **not** serve students, foodie hobbyists, or singles, even though they would happily buy. He prices at a premium because his promise is “dinner solved by 7 p.m., no thinking required,” and a discount would contradict that promise. Only then does he advertise — narrowly, to his chosen segment. Eighteen months later he is profitable.

Maya did more marketing **activity** than Tom and lost. Tom made three **decisions** before any activity and won. This is the entire thesis of the course in one story: the activity is the cheap, visible part; the decisions are the expensive, invisible part, and they are where the value is created or destroyed.

1.2 The three irreversible decisions

Definition 1 (The strategy doctrine) Every marketing strategy, explicit or not, takes a position on three decisions:

- **Where to Play** — which demand you will compete for, and which you will deliberately ignore (segmentation and targeting).
- **How to Win** — why the demand you chose should pick you rather than a rival (positioning and brand).
- **How to Allocate** — how you split finite resources across the levers of the marketing mix (the 4Ps, pricing, brand-vs-activation).

We call them **irreversible** not because they can never be changed, but because changing them is slow and costly: a brand built around premium positioning cannot become a discounter on Tuesday, and a company that has targeted families cannot credibly re-aim at students overnight. Decisions with that property deserve disproportionate care. The order matters too — **Where to Play** logically precedes **How to Win**, which precedes **How to Allocate** — because you cannot choose how to win demand you have not yet chosen to pursue.

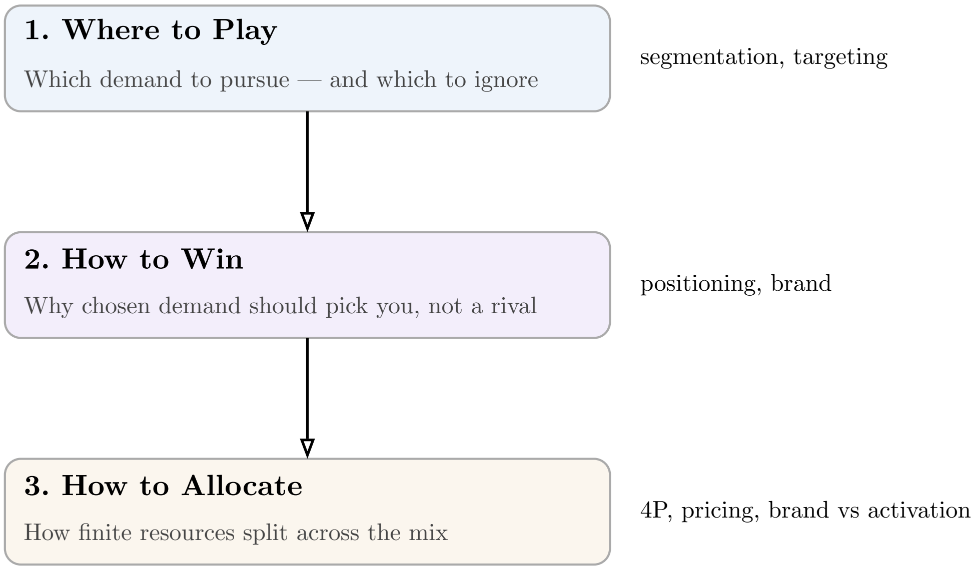

Figure 1: The three irreversible decisions, in logical order. Each constrains the one below it: you cannot decide how to win demand you have not decided to pursue.

> **Worked example — Diagnosing which decision is at stake.** A profitable regional gym chain sees membership growth stall. The CEO asks for "better marketing." A diagnosis using the doctrine:
>
> - If the problem is that the gym appeals to no one in particular against boutique studios and budget chains, the decision at stake is **How to Win** — positioning, not more ads.
> - If the problem is that it has saturated its current segment and must enter a new one (e.g. older members, or corporate accounts), the decision is **Where to Play**.
> - If positioning and segment are sound but spend is scattered across nine half-funded channels, the decision is **How to Allocate**.
>
> "Better marketing" is not an answer. Naming the decision is.

> **Common pitfall** The most common managerial error is to treat every problem as a **How to Allocate** problem — to reach for more spend, more channels, more tactics — when the real failure is upstream in **Where to Play** or **How to Win**. No budget can rescue a strategy that competes for the wrong demand with no reason to be chosen.

1.3 Business strategy vs marketing strategy

Marketing is a powerful lever, but it is not always the lever. Confusing the two wastes money on campaigns when the company's problem lies elsewhere.

> **Definition 2 (When marketing is the lever)** Marketing is the right lever when the binding constraint on the business is on the **demand side**: customers do not know you, do not understand why to choose you, or are not being reached. Marketing is **not** the lever when the constraint is on the **supply side** — a broken product, an uncompetitive cost structure, a capacity ceiling, a regulatory wall. There, marketing only accelerates the underlying problem.

> **Worked example — A demand problem and a supply problem.** Two restaurants have empty tables. Restaurant A serves excellent food but almost no one in the neighbourhood

> has heard of it — a **demand** problem; marketing (awareness, distinctive signage, reviews) is exactly the lever. Restaurant B is well known but the food is mediocre and reviews say so — a **supply** problem; advertising would merely bring more people to a bad meal, accelerating the negative word of mouth. The visible symptom (empty tables) is identical; the correct lever is opposite.

This distinction recurs throughout the course in a sharper, modern form.

1.4 Marketing as demand-side allocation in the AI era

A useful one-line definition: **marketing is the allocation of finite resources against heterogeneous demand to build an asset (the brand) that compounds.** Three words in that sentence carry the weight.

- **Finite** — attention, money, and managerial focus are scarce; spending on one segment, position, or channel is spending withheld from another. This is why marketing is fundamentally about trade-offs, the theme we cost explicitly throughout.
- **Heterogeneous** — demand is not a uniform mass; it varies in needs, in willingness to pay, in what job people are hiring the product to do. If demand were homogeneous, segmentation would be pointless.
- **Compounds** — done well, marketing builds a brand whose value grows and persists, lowering the cost of every future sale. Done as pure short-term activation, it builds nothing and must be re-bought each quarter.

The phrase "in the AI era" is not decoration. As AI collapses the cost of **producing** marketing artefacts — copy, images, variants, micro-targeted messages — toward zero, the production of artefacts ceases to be a source of advantage. What remains scarce, and therefore valuable, is exactly the part AI does not do for you: the **judgement** in the three irreversible decisions. When everyone can generate a thousand ads instantly, the firm that wins is the one that knew which demand to pursue, why it deserves to be chosen, and where to concentrate. The decisions become **more** decisive, not less.

> **Remark** This reframing is why we spend the course on judgement, not tools. A tool you can learn in an afternoon and so can your competitor. The discipline of costing a trade-off and committing to a constraint is what does not commoditise.

Chapter takeaways. Marketing is a decision discipline, not an execution toolkit. Its three irreversible decisions — Where to Play, How to Win, How to Allocate — run in that logical order. Marketing is the right lever only for demand-side constraints. As AI commoditises the production of artefacts, the judgement inside these decisions becomes the scarce asset.

2 Where to Play I — Segmentation as Strategic Hypothesis

2.1 The personas-on-a-wall failure

Walk into many marketing departments and you will find laminated "personas": "Marketing Mary, 34, drinks oat lattes, loves yoga." They decorate a wall and change nothing. The test of a segmentation is brutally simple and almost never applied:

Definition 3 (The allocation test for a segment) A segmentation is **strategic** only if knowing which segment a customer belongs to would **change what you do** — which customers you pursue, what you build, what you charge, where you advertise. If two "segments" would receive identical treatment, they are not two segments. They are decoration.

Segmentation is therefore not a demographic sorting exercise. It is a **hypothesis about why demand is heterogeneous** — and like any hypothesis, it is useful only if it makes different predictions that lead to different actions.

2.2 Demand heterogeneity: the right basis for cutting the market

Demand varies along many axes, and the art is choosing the axis that actually drives different buying behaviour.

> **Worked example — Two ways to cut the toothpaste market.** Cut by demographics (age, income) and you predict little: a wealthy 20-year-old and a poor 60-year-old may buy the same tube. Cut by **what people want the product to do** and the market splits cleanly: the "fresh breath" buyer, the "white teeth" buyer, the "sensitive teeth" buyer, the "kids won't fuss" buyer, the "cheapest that works" buyer. These segments want different products, respond to different claims, and will pay different prices. The second cut changes your allocation; the first does not.

The lesson generalises: demographic and firmographic descriptors are convenient because they are easy to measure, but they are usually **proxies** — and often bad ones — for the thing that truly varies, which is what the customer is trying to accomplish.

2.3 Jobs-to-be-done

That "what the customer is trying to accomplish" has a name, due to Clayton Christensen: the **job to be done**.

Definition 4 (Jobs-to-be-done (JTBD)) A **job to be done** is the progress a customer is trying to make in a given circumstance. Customers do not "buy products"; they "hire" products to do a job, and "fire" them when something does the job better. The unit of analysis is the job, not the customer and not the product.

> **Worked example — The milkshake.** Christensen's fast-food client wanted to sell more milkshakes and segmented by demographics — fruitlessly. Studying **when** milkshakes were bought revealed a surprising job: morning commuters "hired" a thick milkshake to make a long, boring drive less dull and stave off mid-morning hunger, one-handed and mess-free. Its competitors were not other milkshakes but bananas, bagels, and boredom. The job — not the demographic — told the company what to improve (make it thicker, move it to a pre-pay queue) and what **not** to (it would ruin the product for the afternoon "treat for my child" job, an entirely different segment).

The milkshake story also teaches the **competitive frame**: your real competitors are whatever else gets the job done, which may not be in your category at all — a theme we return to in competitive landscapes.

2.4 Stated vs revealed preferences

Segmentation built on what customers **say** is treacherous. People misreport, flatter themselves, and rationalise.

Definition 5 (Stated vs revealed preference) **Stated preferences** are what customers say they want (surveys, focus groups, interviews). **Revealed preferences** are what their actual behaviour shows they want (purchases, usage, what they pay, what they abandon). When the two conflict, revealed preference is the truth and stated preference is the story.

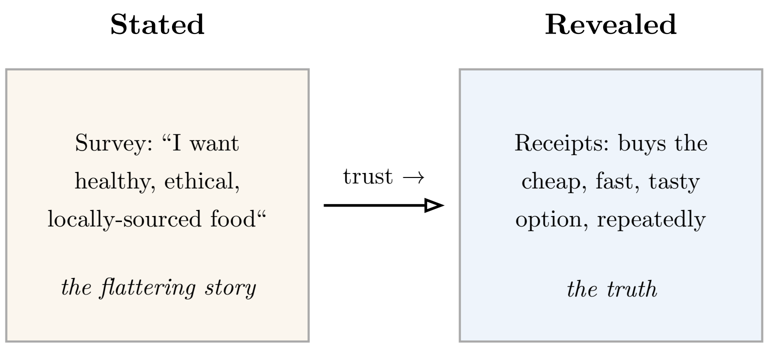

Figure 2: When stated and revealed preferences diverge, build the segmentation on behaviour. Surveys capture the aspirational self; receipts capture the buying self.

Common pitfall Whole product lines have died because their segmentation was built on survey-stated virtue ("consumers want healthy / sustainable / premium") that evaporated at the checkout, where the same consumers revealed they wanted cheap and convenient. Always ask: would this segment survive contact with a credit-card statement?

Exercises

1. Take a category you know well (coffee, gyms, streaming). Propose one demographic segmentation and one jobs-to-be-done segmentation of the same market. For each, state explicitly what you would do differently for two of the segments. Which segmentation passes the allocation test?

2. Find one product whose marketing clearly rests on a **stated** preference that its buyers' **revealed** behaviour contradicts. Describe the contradiction.

3. Rewrite "Marketing Mary, 34, loves yoga" as a job-to-be-done statement of the form "When [circumstance], I want to [progress], so I can [outcome]."

Chapter takeaways. A segment is real only if it changes your allocation (the allocation test). Cut the market by demand heterogeneity — ideally by the job to be done — not by easy-to-measure demographic proxies. Trust revealed over stated preference whenever they diverge.

3 Where to Play II — Targeting as Capital Allocation

3.1 Choosing a segment is spending money you cannot get back

Segmentation describes how demand is heterogeneous; **targeting** is the decision of which segments to spend against. Treat that decision as what it is — a capital allocation. Aiming your product, message, and money at one segment is withholding them from every other. The relevant question is never "is this segment attractive?" (many are) but "is this segment more deserving of our finite resources than the alternatives we are thereby refusing?"

3.2 Broad vs narrow: the central tension

There is a genuine, unresolved tension here, and good marketers hold both sides.

Definition 6 (The targeting tension)

- **Narrow targeting** concentrates limited resources where they bite hardest, lets you tailor the product and message precisely, and is usually right for challengers and new entrants who cannot afford to be everywhere.
- **Broad targeting** reaches more of the category's buyers and is supported by evidence (below) that brands grow mainly by **acquiring many light buyers** across the whole category, not by deepening loyalty within a niche.

These pull in opposite directions. The resolution is usually **narrow in execution, broad in ambition**: concentrate now to establish a beachhead, while designing the brand so it can scale across the category later.

3.3 The Ehrenberg-Bass perspective

The narrow-targeting instinct collides with one of the most robust empirical findings in marketing, associated with the Ehrenberg-Bass Institute (Byron Sharp, *How Brands Grow*).

Definition 7 (How brands actually grow (Ehrenberg-Bass)) Across categories, large brands do not have dramatically more loyal customers than small brands; they have **more customers**, most of whom buy **infrequently** (the **law of double jeopardy**: small brands suffer twice — fewer buyers, who also buy slightly less often). Growth comes overwhelmingly from increasing **penetration** — acquiring many light and non-buyers across the whole category — not from squeezing more loyalty out of a heavy-buying niche.

The practical consequence is uncomfortable for the narrow-targeting purist: mass reach and broad availability matter more than most challenger marketers want to believe. A brand that talks only to its devoted niche caps its own growth, because its future revenue lives disproportionately in the casual, barely-interested buyers it is ignoring.

Common pitfall Do not over-correct. Ehrenberg-Bass argues against **narrow loyalty-only** strategies for **established** brands seeking growth; it does not license a cash-strapped startup to advertise to "everyone." Reconcile the two by sequence: start narrow to earn a foothold and a reason to be chosen, then broaden reach as resources allow. The error is mistaking the startup phase for the permanent strategy.

3.4 Category entry points

Ehrenberg-Bass also reframes **how** to target broadly without wasting money on the genuinely irrelevant.

Definition 8 (Category entry points (CEPs)) **Category entry points** are the cues — situations, needs, moments, occasions — that bring a category to mind. ("I need a quick lunch near the office"; "I want a gift that looks expensive"; "the kids are bored on a rainy afternoon.") Brands grow by becoming mentally linked to **more** of the category's entry points, so they are recalled in more buying situations.

CEPs are the bridge between jobs-to-be-done and growth: each entry point is roughly a job in a buying moment, and the more of them your brand owns, the more of the category's demand you can capture. This is why we target jobs and occasions, not just demographic groups — occasions are where buying actually gets triggered.

3.5 Defending the decision to ignore a segment

The hardest, most valuable skill in targeting is **saying no with evidence**. Anyone can list attractive segments; a strategist can defend ignoring one and quantify what that costs.

Definition 9 (Costing the segment you refuse) To defend ignoring a segment, you must (1) **size** it — how many customers, spending how much, i.e. its market value; (2) state the **opportunity cost** — the profit you forgo by not serving it; and (3) justify why serving it would **damage** your position with the segment you chose (diluted message, conflicting product requirements, margin erosion) by more than that forgone profit.

Worked example — Why a premium brand ignores the bargain segment. A premium audio brand sizes the budget-buyer segment at, say, €40M of annual category spend — real money. It refuses it anyway, and can defend the refusal: serving budget buyers would require a cheap product line whose mere existence signals "this brand is affordable," eroding the premium price its core segment pays. If that signal shaved 10% off €300M of premium revenue (€30M) to chase a thin-margin €40M, the refusal is correct. The number — not taste — justifies the no.

Common pitfall Sizing a segment is not the same as targeting it. A large, attractive segment can be exactly the one to ignore if serving it contradicts your **How to Win**. The common mistake is to treat every quantified opportunity as an obligation.

Exercises

4. For a challenger brand of your choice, write a one-paragraph targeting decision that is "narrow in execution, broad in ambition." Name the beachhead segment and the eventual broad ambition, and the trigger to broaden.

5. List five category entry points for instant coffee. For each, say whether the market leader is mentally linked to it, and where a challenger could enter.

6. Pick a segment a real company deliberately does **not** serve. Estimate its size with public data, then write the three-part defence (size, opportunity cost, why serving it would damage the core) for ignoring it.

Chapter takeaways. Targeting is capital allocation: choosing a segment withholds resources from all others. Hold the broad/narrow tension by sequencing — narrow foothold, broad ambition. Ehrenberg-Bass shows growth comes from penetration and category entry points, not niche loyalty. The signature skill is defending an ignored segment by sizing it and costing the refusal.

4 How to Win I — Positioning as Constraint

4.1 A position is a promise you let constrain you

Having chosen **where to play**, you must give the chosen demand a reason to pick you. That reason is your **positioning**. The defining property of a good position is counterintuitive: it is valuable precisely because it **forbids** things.

> **Definition 10 (Positioning)** A **position** is the specific, defensible place a brand occupies in the customer's mind relative to alternatives — the answer to "for [segment] who [job/need], we are the [frame] that [differentiator], unlike [competitors]." A position is only real if it **constrains future decisions**: if it rules some actions in and others out. A statement that permits everything positions nothing.

> **Worked example — A position that constrains vs one that does not.** "We make high-quality products with great customer service at a fair price" constrains nothing — every competitor would sign it, and it forbids no decision. Contrast Volvo's historical "the safest car": it **constrains**. It rules in spending on crash engineering and safety storytelling; it rules **out** marketing the car as a thrilling sports machine, even when that segment looks tempting. A position you would never let veto a tempting opportunity is not a position.

4.2 Competitive frames of reference

A position is meaningless without a **frame**: the set of alternatives the customer mentally compares you against. Change the frame and the same product is strong or weak.

> **Worked example — Reframing.** A €4 energy bar looks expensive **framed as a snack** (against a €1 chocolate bar) and cheap **framed as a meal replacement** (against a €9 lunch). Choosing the frame — "this competes with lunch, not with sweets" — is half the positioning decision, because it sets the reference price and the relevant differentiators.

4.3 The four stress tests

A positioning statement is cheap to write and easy to fool yourself with. Subject every draft to four tests; a position that fails any of them is not yet a position.

1. Could a competitor copy it?
If yes → not differentiating.

2. Does it constrain decisions?
If no → it's a slogan, not a position.

3. Would anyone argue the opposite?
If no → it's a platitude ("we value quality").

4. Could it be any company?
If yes → it's generic, not yours.

Figure 3: The four stress tests. A positioning statement must pass all four; the failure mode beside each is the trap it catches.

Definition 11 (The four stress tests)

1. **Could a competitor copy it?** If a rival could truthfully claim the same thing, it does not differentiate you.
2. **Does it constrain decisions?** If it never tells you what **not** to do, it is a slogan, not a strategy.
3. **Would anyone sane argue the opposite?** If no one would ever champion the reverse ("we offer poor service"), you have stated a platitude, not a choice.
4. **Could it be any company?** If you could swap in a competitor's name and the statement still reads true, it is generic.

Worked example — Running the tests. Draft: "We are the meal-kit for busy people who value convenience."

- Copy? Every meal-kit claims convenience for busy people — **fails 1**.
- Constrain? It would never stop us doing anything — **fails 2**.
- Opposite? No one targets relaxed people who hate convenience — **fails 3**.

Revised: "We are the only meal-kit that guarantees dinner on the table in 15 minutes for exhausted parents of young kids — so we refuse recipes over 15 minutes, however gourmet." Now competitors with 40-minute gourmet kits cannot copy it (1), it forbids slow recipes (2), and a rival could sincerely argue the opposite — "real cooking takes time" (3). It passes.

4.4 Differentiation vs distinctiveness

Here lies one of the most important — and most misunderstood — distinctions in the course, due to Jenni Romaniuk and Byron Sharp.

Definition 12 (Differentiation vs distinctiveness (Romaniuk)) **Differentiation** is being **meaningfully different** — offering a reason to choose you (safer, faster, cheaper). **Distinctiveness** is being **recognisable** — owning assets (colours, logos, characters, sounds) that let buyers identify you instantly and effortlessly, without necessarily being different at all. Differentiation answers "why choose us?"; distinctiveness answers "which one are we?"

The misunderstanding is to chase differentiation alone. Romaniuk's evidence is that in most mature categories, products are far more similar than marketers admit, real differentiation is small and fleeting (competitors copy it), and what actually drives buying is **distinctiveness** —

being effortlessly noticed and recognised at the moment of choice. You need some differentiation to justify existing, but you need distinctiveness to be **chosen amid clutter**.

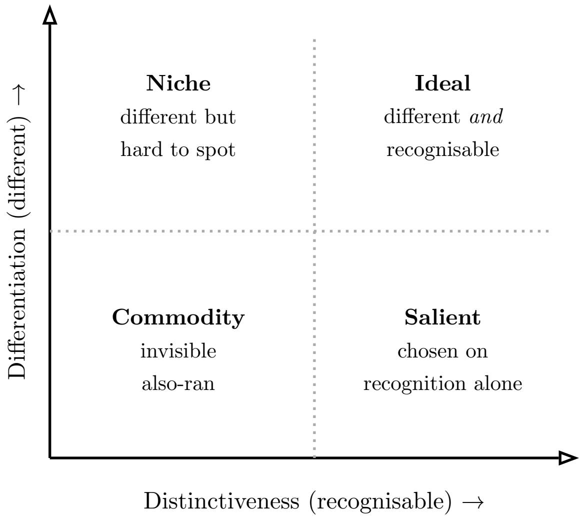

Figure 4: Differentiation and distinctiveness are independent axes. Marketers obsess over the vertical axis; Romaniuk's evidence says the horizontal axis (being recognised) often does more of the work in driving choice.

Common pitfall The classic blunder is a "rebrand" that improves differentiation (a cleverer value proposition) while **destroying** distinctiveness — throwing away the recognisable colour, logo, or jingle the brand spent years building. You can be the most differentiated brand in the category and still lose if shoppers no longer recognise you on the shelf. Guard distinctive assets like capital.

Chapter takeaways. A position is valuable because it constrains — it forbids tempting actions. State it within an explicit competitive frame, and stress-test every draft against the four tests. Distinguish differentiation (why choose us) from distinctiveness (which one are we); both matter, but distinctiveness is chronically underrated and easily destroyed.

5 How to Win II — Brand as Strategic Asset

5.1 Why brand is an asset, not an expense

Accountants book marketing as an expense. Strategists treat the brand it builds as an **asset** — something that, once built, lowers the cost of every future sale and persists between campaigns. This chapter makes that claim precise with two frameworks, both empirical rather than mystical.

5.2 Mental availability vs brand awareness

Definition 13 (Mental availability (Sharp)) **Brand awareness** asks: do people recognise the name? **Mental availability** asks the harder question: does the brand come to mind, **quickly and in the right buying situations**, across many category entry points? A brand can be widely **aware** yet poorly **available** — everyone knows it, but no one thinks of it when actually buying. Mental availability, not bare awareness, predicts sales.

> **Worked example — Aware but not available.** Almost everyone recognises the name of some old insurance or airline brand — high awareness. But if, at the moment of buying travel insurance, it never **comes to mind** (no link to the entry point "booking a risky trip"), its awareness is commercially inert. A challenger with lower awareness but a strong mental link to that exact buying moment will out-sell it. Awareness is a vanity metric; availability is the asset.

Mental availability is built by linking the brand to **many category entry points** (Chapter 3) and by repetition through **distinctive assets** (Chapter 4) so the brand is recognised instantly when a buying cue fires. This is where the "asset" language becomes concrete: each link, once built, keeps working.

5.3 Distinctive brand assets

> **Definition 14 (Distinctive brand assets)** **Distinctive brand assets** are the non-name elements that trigger the brand in memory: colours (Tiffany blue, Coca-Cola red), shapes (the Coke bottle), logos, characters (the Michelin Man), sounds (the Intel chime), typefaces, taglines. Their value is measured on two dimensions: **fame** (what share of buyers link the asset to the brand) and **uniqueness** (whether they link it to **only** you, or also to rivals). An asset that is famous but not unique is a liability — you are advertising the category, or worse, a competitor.

> **Worked example — Auditing distinctive assets.** Take a brand you know and list candidate assets: colour, logo, character, sound, slogan. For each, ask the two questions. A purple wrapper that 80% of buyers link to **this** chocolate brand and to no other is a prime asset (famous
> 1. unique). A "swoosh-like" tick that buyers confuse with three other sportswear
>
> logos is famous but not unique — a warning sign. This audit is exactly Learning Outcome 5.

5.4 The long and the short of it: Binet & Field

The final and most consequential allocation framework concerns **time**. Les Binet and Peter Field, analysing hundreds of campaigns, separated two distinct jobs marketing does.

> **Definition 15 (Brand building vs sales activation (Binet & Field))** **Sales activation** (the "short") drives immediate response — promotions, performance ads, "buy now." Its effect is large but **spikes and decays** quickly; it captures existing demand. **Brand building** (the "long") grows mental availability and pricing power; its effect is **slow to appear but accumulates and persists**, creating future demand. Both are necessary, and they work on different timescales.

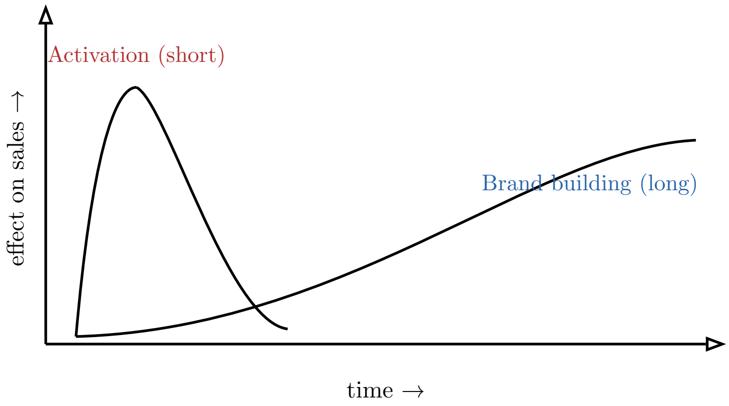

Figure 5: The two effects on different timescales. Activation spikes and decays; brand building accumulates and persists. Fund only activation and you must re-buy every sale; fund only brand and you starve this quarter's revenue.

Definition 16 (The 60/40 benchmark) Across many categories, Binet & Field find the long-run-optimal split is roughly **60% of budget to brand building, 40% to activation**. It is a **benchmark, not a law**: the right split shifts with category, business model, and maturity — shorter for low-consideration impulse goods and discount retail, longer for considered, high-margin, or subscription businesses. The value of 60/40 is as a reference against which to justify deviations.

Worked example — Using 60/40 against a real spend split. Suppose a company spends 85% on performance ads and 15% on brand — heavily short. Against the 60/40 benchmark this flags under-investment in the long effect, predicting rising customer-acquisition cost over time as it exhausts existing demand without replenishing future demand. The recommendation is not "move to 60/40 because the book says so," but "shift toward brand **because** your category is considered and high-margin, where the long effect compounds" — the category-specific reasoning Learning Outcome 11 demands.

Common pitfall Activation is seductive because its results are **immediate and measurable**, while brand building is slow and hard to attribute. Dashboards therefore bias firms toward over-spending on the short term — quietly cannibalising the long-term brand equity that makes future activation cheap. Beware optimising what is easy to measure over what matters.

Exercises

7. Audit one real brand's distinctive assets. List five candidates and rate each on fame and uniqueness. Identify the strongest asset and one that is famous-but-not-unique.

8. Find a brand with high awareness but, in your judgement, low mental availability for a specific buying occasion. Explain the gap using category entry points.

9. Estimate a real company's brand-vs-activation split from public signals (its ads, promotions, channels). Compare to 60/40 and recommend an adjustment with category-specific reasoning — not a generic "move to 60/40."

Chapter takeaways. The brand is an asset that lowers the cost of future sales. Mental availability — coming to mind in real buying moments — beats bare awareness, and is built via category entry points and distinctive assets (rated on fame and uniqueness). Binet & Field's long/short split (≈60/40, adjusted by category) governs the trade-off between building future demand and capturing present demand.

6 How to Allocate I — The Marketing Mix as a Trade-off Framework

6.1 The mix is a set of trade-offs, not a checklist

The **marketing mix** — the 4Ps (Product, Price, Place, Promotion), extended in services to 7Ps (adding People, Process, Physical evidence) — is usually taught as a checklist to fill in. That framing is actively harmful. The 4Ps are not independent boxes to optimise one at a time; they are **coupled levers**, and pushing one constrains the others. Allocation means choosing a **coherent combination**, accepting that strengthening one P typically weakens another.

Definition 17 (The marketing mix as allocation) The 4Ps — **Product**, **Price**, **Place**, **Promotion** (plus services' **People**, **Process**, **Physical evidence**) — are the levers through which strategy is executed. The skill is not to maximise each lever but to set them so they **reinforce one position**. Every setting trades off against the others: a richer product raises cost (price), a premium price narrows place (distribution), a promotion-heavy push can undermine the product's premium signal.

Worked example — How optimising one P breaks another. A coffee brand wants the **best** on every P: finest beans (Product), lowest price (Price), available everywhere including discounters (Place), and frequent coupons (Promotion). The combination is incoherent. Premium beans force a higher cost that the low price cannot cover; selling through discounters contradicts the premium product; constant coupons train customers to wait for deals, destroying the price. "Best at everything" is not a strategy — it is four optimisations that cancel out. A coherent mix instead might be: premium product, premium price, selective place (specialist cafés), brand-building promotion rather than coupons — each P paying for and reinforcing the others.

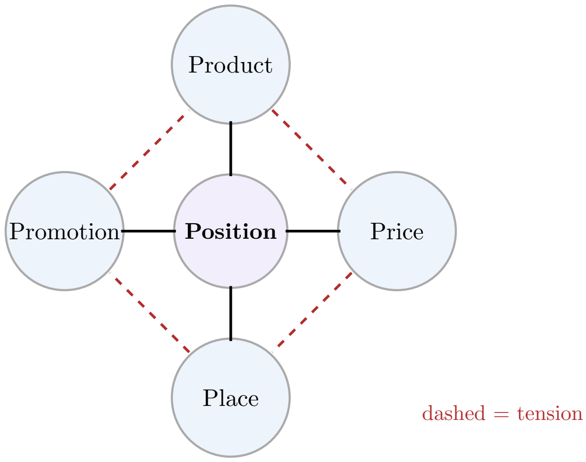

Figure 6: The 4Ps must all serve one position (centre). Solid spokes: each lever expresses the strategy. Dashed links: adjacent levers pull against each other — optimising one constrains its neighbours.

Common pitfall Beware the "best at everything" trap: a plan that calls for the highest quality, the lowest price, the widest distribution, and the heaviest promotion is not ambitious — it is incoherent, and usually reveals that no real **Where to Play** or **How to Win** choice was made upstream.

6.2 Allocating a fixed budget across the 4Ps

Because the levers are coupled and the budget is fixed, allocation forces written justification: every euro to one P is a euro withheld from another, and the test is whether the combination reinforces the position.

> **Worked example — A fixed-budget allocation with justified trade-offs.** A new premium tea brand has €500k. It might allocate: €200k Product (sourcing, packaging that signals premium), €50k Price (no discounting — instead, spend on margin-protecting value communication), €100k Place (a few high-end stockists, not mass retail), €150k Promotion (brand building over activation). The justification for **each** line references the position: heavy Product because the premium promise must be real; thin Place because wide distribution would cheapen the brand; promotion skewed long because the asset must compound. Move €150k from Product into mass-retail Place and you would reach more people but contradict the premium position — a quantified trade-off, which is Learning Outcome 6.

Chapter takeaways. The 4Ps/7Ps are coupled levers, not a checklist. A good mix sets them to reinforce a single position; optimising one in isolation constrains the others. Allocating a fixed budget means writing the justification for every trade-off and rejecting the incoherent "best at everything" plan.

7 How to Allocate II — Pricing as Strategic Signal

7.1 Price is the loudest message you send

Price deserves its own chapter because it is the only P that directly captures value rather than creating cost, **and** because it is read by customers as a **signal** of quality and positioning. A price is never just a number; it is a sentence about what kind of brand you are. This is why pricing belongs to **How to Win** as much as to **How to Allocate**.

7.2 Three ways to set a price

Definition 18 (Three pricing logics)

- **Cost-plus**: take your unit cost and add a target margin. Simple and safe, but ignores both the customer's willingness to pay and competitors — routinely leaves money on the table or prices you out.
- **Competitive**: price relative to rivals (match, undercut, or premium-to). Keeps you in the market's frame but cedes the decision to competitors and invites price wars.
- **Value-based**: price to the value the customer perceives and the job's worth to them, largely independent of your cost. Hardest to do — it requires understanding the customer deeply — and usually the most profitable.

Worked example — The same coffee, three prices. A cup of coffee costs €0.40 to make. **Cost-plus** at a 150% margin → €1.00. **Competitive**, matching the café next door → €2.50. **Value-based**, recognising that in a business-district café the customer is buying a place to meet a client for an hour → €4.50, and the customer does not blink. Same liquid; the pricing **logic**, not the cost, sets the number.

7.3 Price as positioning signal

Common pitfall Because price signals quality, **the cheapest option is not automatically the most attractive**, and a discount is not a free lever. Cutting price on a premium brand does not just sacrifice margin — it tells the market the brand is worth less, can be hard to reverse, and trains customers to wait for the next discount. A premium position and habitual discounting are contradictory; one must give. This contradiction is among the most common audit findings (next chapter).

7.4 Trading off margin, volume, and signal

Every pricing recommendation must reconcile three things that pull apart, and the right answer depends on the strategy, not on a formula.

Definition 19 (The pricing trade-off) A price simultaneously sets **margin** (per-unit profit), **volume** (how many you sell — lower price usually sells more), and **positioning signal** (what the price says about the brand). A low price may win volume but cut margin **and** signal cheapness; a high price protects margin and signals premium but caps volume. The recommendation must state which two you are prioritising and why, given the position chosen upstream.

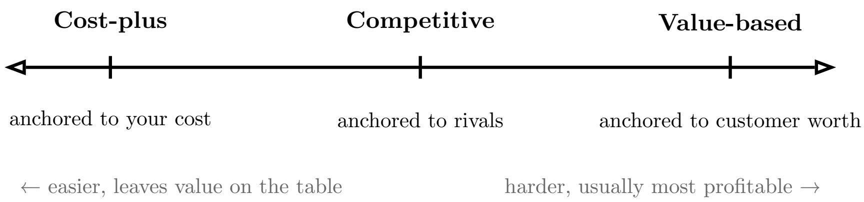

Figure 7: The pricing spectrum by what the price is anchored to. Moving right ties the price to customer-perceived value rather than your costs or your rivals' — harder to estimate, but where margin and positioning power live.

> **Worked example — A pricing recommendation with explicit trade-offs.** For a premium skincare launch: recommend **value-based** pricing at the high end. Trade-off stated explicitly — we prioritise **margin** and **signal** over **volume**, accepting fewer units sold, because the position is premium and a low price would both erode margin and contradict the promise. We reject competitive matching (it would cede our premium frame) and cost-plus (it would under-price the perceived value). This is the shape Learning Outcome 7 asks for: a recommendation that names the trade-off, not just the number.

Chapter takeaways. Price both captures value and signals quality, so it is a strategic decision, not an arithmetic one. Choose among cost-plus, competitive, and value-based logics — value-based is hardest and usually most profitable. Every price trades off margin, volume, and signal; a recommendation must say which two it prioritises and why, consistent with the position.

8 Competitive Landscapes — Where Not to Fight

8.1 Strategy is partly the art of avoidance

A competitive analysis that only asks "how do we beat them head-on?" misses the more profitable question: "which fights should we decline?" Just as targeting is about which segments to ignore, competitive strategy is substantially about which battles to **avoid** — choosing terrain where you are strong and the incumbent is weak or unwilling to follow.

> **Definition 20 (Strategic avoidance)** **Strategic avoidance** is deliberately competing where strong rivals are absent, complacent, or constrained from responding — rather than attacking their strength. A challenger rarely wins by being a slightly better version of the leader on the leader's own terms; it wins by redefining the terrain so the leader's advantages do not apply.

8.2 Direct and indirect competition

> **Definition 21 (Direct vs indirect competition)** **Direct** competitors offer the same kind of solution (Pepsi vs Coke). **Indirect** competitors satisfy the same job by different means (for "stay alert at work": energy drinks, coffee, a nap, a walk). Recall the milkshake — the real competitor was bananas and boredom. Map competition by the **job**, or you will defend against the wrong rivals and be blindsided by substitutes outside your category.

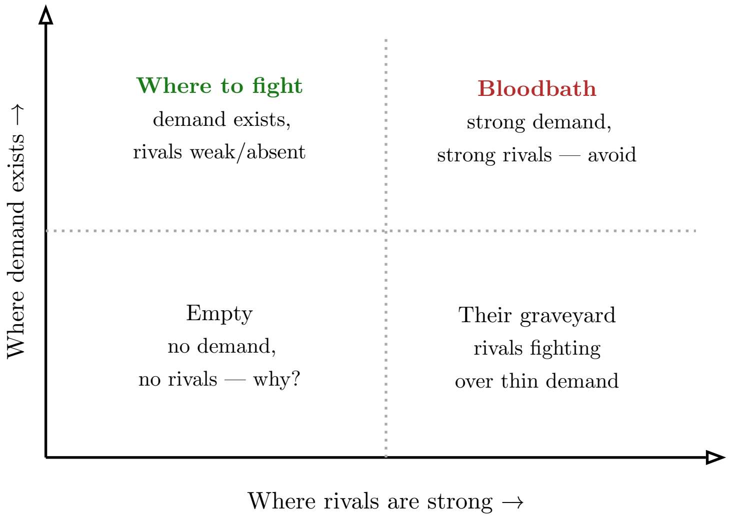

Figure 8: A "where not to fight" map. Seek the upper-left: real demand where rivals are weak or absent. The upper-right — strong demand defended by strong rivals — is where challengers bleed.

8.3 Writing a "where not to fight" recommendation from public data

The capstone competitive skill (Learning Outcome 8) is to map a landscape from **public** evidence and recommend where **not** to compete.

Definition 22 (Reading a landscape from public signals) Public data is richer than it seems: **financials** (revenue, margins, R&D and marketing spend) reveal where a rival is committed and where it is starving; **web traffic and app rankings** reveal demand and momentum; **pricing pages** reveal positioning and segments served; **messaging** reveals the position they have committed to — and therefore the moves they **cannot** make without contradicting themselves. A rival's stated position is also a map of the ground it has **ceded**.

Worked example — From signals to a one-page recommendation. Suppose the category leader's pricing page shows only enterprise tiers, its messaging is all "security and compliance," and its financials show fat margins it will protect. The reading: it is **structurally unwilling** to serve small customers cheaply (that would threaten its margins and enterprise signal). The "where not to fight" recommendation writes itself — do **not** attack enterprise security head-on; instead serve the small, self-serve segment the leader has ceded and cannot chase without cannibalising its premium. Evidence (the pricing page, the margin) backs every sentence.

Common pitfall Do not confuse "a rival is big" with "a rival will respond." The strategically decisive fact is often what a competitor **cannot or will not** do because of its own commitments — its position, cost structure, or shareholders. The ground a strong rival has **ceded** is frequently the best place to build.

Chapter takeaways. Competitive strategy is largely avoidance: fight where rivals are weak, absent, or constrained, not where they are strong. Map competition by the job (direct **and** indirect substitutes). Read landscapes from public financials, traffic, pricing, and messaging — and exploit the ground a committed rival cannot defend without contradicting itself.

9 Auditing and Quantifying the Marketing Mix

9.1 Coherence is the thing you audit for

We can now close the loop. The three decisions — Where to Play, How to Win, How to Allocate — must form a **coherent whole**, and the job of an audit is to find where a real company's choices contradict each other. A marketing audit is not a hunt for "things we could do more of"; it is a hunt for **contradictions** between stated position and actual mix.

> **Definition 23 (The coherence audit)** To audit a company's marketing mix, state its **claimed position**, then test each element of its actual mix against that claim and flag contradictions. Typical flags: premium positioning undermined by frequent discounting; a job-to-be-done the targeted segment does not actually have; short-term activation cannibalising long-term brand equity; distribution that contradicts the price signal; a distinctive asset being discarded. Each flag must be **concrete** — pointing at a specific decision, not a vibe.

> **Worked example — An audit finding.** A brand positions itself as "affordable luxury" yet (a) runs a 40%-off promotion every month — training customers never to pay full price, contradicting "luxury"; (b) sells through deep-discount marketplaces — contradicting the premium signal; and (c) spends 90% on performance ads — starving the brand equity "luxury" requires. Three concrete contradictions, each tied to a specific decision. This is Learning Outcome 9.

9.2 Quantifying the impact: a simplified P&L

Auditing tells you **what** is incoherent; quantification tells you what it **costs** and whether fixing it pays. The instrument is a simplified profit-and-loss scenario.

> **Definition 24 (A simplified marketing P&L)**
>
> $$\text{Profit} = (\text{Price} - \text{Unit cost}) \times \text{Volume} - \text{Marketing spend}$$
>
> Changing any marketing-mix lever ripples through this identity: a price cut lowers the first bracket but may raise Volume; a brand investment raises Marketing spend now but should raise Price (pricing power) and Volume (mental availability) later. To compare two mix decisions, project each through this P&L and compare the resulting profit — making every trade-off a number, not an opinion.

> **Worked example — Comparing two mix decisions on the P&L.** Current: Price €50, Unit cost €20, Volume 100,000, Marketing €1.0M.
>
> $$\text{Profit} = (50 - 20) \times 100\{,\}000 - 1\{,\}000\{,\}000 = €2.0\text{M}.$$
>
> **Decision A — discount to €40** to chase volume. Suppose volume rises to 140,000:
>
> $$(40 - 20) \times 140\{,\}000 - 1\{,\}000\{,\}000 = €1.8\text{M}.$$
>
> Despite 40% more units, profit **falls** — the margin cut outweighs the volume gain (and the discount may damage positioning, a cost not even in the number).
>
> **Decision B — hold price, add €0.5M brand building**, lifting volume to 120,000 **and** supporting a later price rise:

$$(50-20) \times 120\{,\}000 - 1\{,\}500\{,\}000 = €2.1\text{M},$$

with a stronger asset for next year. The P&L turns "should we discount or build the brand?" into a comparison of €1.8M vs €2.1M — Learning Outcomes 10 and 6 made quantitative.

Common pitfall A simplified P&L is a decision aid, not the whole truth: it omits brand effects that land **next** year, competitor responses, and the long-run cost of signalling cheapness. Use it to **quantify** trade-offs, but never let a tidy one-period number override a positioning argument whose costs fall outside the period. The number informs the judgement; it does not replace it.

Exercises

10. Pick a real company. State its claimed position in one sentence, then audit its actual mix (product, price, place, promotion) against it. Flag at least two **concrete** contradictions.
11. Build a simplified P&L for a product with assumed price, cost, volume, and spend. Compute the profit impact of (a) a 15% price cut that lifts volume 25%, and (b) holding price while adding 30% to brand spend for a 10% volume lift. Which wins, and what does the P&L **omit**?
12. For the audit in the first exercise, recommend one change and estimate its P&L impact, stating explicitly which effects your number cannot capture.

Chapter takeaways. An audit hunts for **contradictions** between stated position and actual mix, flagged concretely. A simplified P&L — (Price − Cost) × Volume − Spend — turns mix trade-offs into comparable numbers, exposing traps like volume-chasing discounts. But the P&L omits next-year brand effects and signalling costs; it informs judgement rather than replacing it.

10 B2B vs B2C — Adapting the Doctrine

10.1 The doctrine is invariant; the application is not

A reassuring closing message: nothing in the three irreversible decisions changes between consumer (B2C) and business (B2B) markets. You still choose Where to Play, How to Win, How to Allocate. What changes is the **texture** of each decision — texture that, ignored, quietly breaks an imported B2C playbook.

Definition 25 (What differs in B2B) Relative to B2C, B2B markets typically feature: **fewer, larger buyers** (so each account matters and mass reach matters less); **buying groups, not individuals** (procurement, users, finance, and an executive sponsor — each with a different job); **longer, considered sales cycles** (favouring the long brand effect and relationship over impulse activation); **higher switching costs** and stickier relationships; and **price set by negotiation and total-cost-of-ownership**, not a shelf tag. The frameworks hold; their **settings** change.

Worked example — Same decision, different texture. Segmentation: in B2C you may segment by job and occasion across millions; in B2B you segment by firmographics **and** by

> the buying group's differing jobs — the engineer hires your product to "ship faster," the CFO to "cut cost," the CISO to "not get fired for a breach." One position rarely satisfies all three, so How to Win in B2B often means tailoring the **message** per role around one coherent core position.
>
> **Brand**: it is tempting to think B2B needs no brand — buyers are "rational." But mental availability matters **more** when the buying cycle is long: you must be the brand that comes to mind when a need finally crystallises months later, on a shortlist you were added to before the budget existed. The long effect dominates.

> **Common pitfall** The classic cross-over error is importing a B2C activation playbook into B2B — heavy short-term lead-gen, no brand building — and then wondering why deals stall and acquisition costs climb. With long cycles and buying groups, under-investing in the long brand effect is even more damaging than in B2C. Adapt the **settings**; do not abandon the **doctrine**.

10.2 Closing

You began with the belief, common and comfortable, that marketing is the production of ads. You end, I hope, with a harder and more useful belief: that marketing is the disciplined making of three irreversible decisions under scarcity — choosing which demand to pursue and which to refuse, why you deserve to be chosen, and where to concentrate finite resources — and that the value lies in the choices you are willing to **forbid** yourself. The tools will keep changing; AI will keep cheapening the artefacts. The judgement to choose what **not** to do is the part that lasts, and the part that is now yours to practise.

> **Chapter takeaways.** The three decisions are invariant across B2C and B2B; only their settings change — fewer/larger buyers, buying groups with conflicting jobs, long considered cycles, negotiated pricing. Long-run brand building matters **more** in B2B, not less. Adapt the application; keep the doctrine.