

BUS33-2

Competitive Strategy

Five Forces, VRIO, generic strategies, Blue Ocean, competitive dynamics — and the diagnosis that ties them together

Reasoning strategy against strategy, under ambiguity and time pressure.

Albert School — 22 hours

Preface

You do not arrive empty-handed. From **Microeconomics** (BUS13-4) you can already read a market structure — perfect competition, monopolistic competition, oligopoly, monopoly — and you know that profit above the cost of capital is, in the long run, a **disequilibrium** that something must be holding in place. From **Business Models & Value Creation** (BUS23-2) you can decode how a single firm creates value, captures some of it, and defends that capture; you can compute its unit economics — CAC, LTV, gross margin — in your sleep. From **Financial Diagnosis** (BUS23-3) you can measure whether an advantage actually pays, through ROIC and free cash flow. And from **Macroeconomics** (BUS23-4) you know that none of this happens in a vacuum: rates, cycles, and policy set the weather.

This course takes those tools and points them at a harder target. So far you have analysed firms essentially **one at a time** — a business model, a set of accounts, a market structure. Competitive strategy is what happens when you put **two thinking firms in the same room** and let them react to each other. The central difficulty is no longer “is this a good business?” but “is this a good business **given that a rival, just as clever as you, is trying to take what you have?**” Advantage is not a property a firm owns; it is a position it **holds against pressure**, and the whole question is whether the position is defensible and for how long.

This is the **Compete** stage of the four-stage Strategy arc — Build (BUS23-2) → Compete (this course) → Evolve (Tech & Platform Strategies, S6) → Allocate (Strategic Allocation, S6). It hands you a toolkit — Five Forces, VRIO, the value chain, generic strategies, the Blue Ocean canvas, strategic groups, the logic of competitive dynamics — but it insists on something most strategy courses quietly skip: **every tool has boundary conditions**. Five Forces breaks in platform markets. VRIO flatters a resource that the environment is about to make worthless. A generic strategy that is brilliant in one market structure is suicide in another. A course that teaches these frameworks as checklists produces students who can fill a template and conclude nothing. We will instead teach them as instruments with a known range — and drill them through **timed case methodology** (MECE decomposition, hypothesis-led reasoning), because in the real diagnosis you will never have all the data and never have enough time.

The arc ends where it must: you take a **real** company, run the external analysis (Five Forces plus strategic groups), run the internal analysis (VRIO plus value chain), and defend a **quantified, falsifiable recommendation** before a jury that may include a practitioner who knows the industry better than you do. By then the reflex this course exists to build should be automatic: never to recite a framework, but to ask of any advantage — **what holds this in place, who is trying to take it, and what happens when they do?**

Contents

Preface	2
1 Industry analysis with Five Forces	3
1.1 Two restaurants, two destinies	3
1.2 The five forces	3
1.3 Quantify, do not recite	4
1.4 Where the framework breaks	5
2 Resource-based view and VRIO	6
2.1 Why two firms in the same industry earn different returns	6
2.2 VRIO is a durability test, not a strength test	6
2.3 The mistake that VRIO is designed to prevent	7
3 Value chain and activity systems	7

3.1	Why you cannot copy Ryanair by copying one thing	7
3.2	From a list of activities to a system	8
3.3	Static endowments versus dynamic capabilities	8
4	Competitive positioning	9
4.1	Three ways to not be like everyone else	9
4.2	Porter's generic strategies	9
4.3	Blue Ocean: making the competition irrelevant	10
4.4	Strategic groups: who you actually fight	11
5	Competitive dynamics	11
5.1	Strategy as a move in a game the rival also plays	11
5.2	Why a threat must be credible to work	12
5.3	When ignoring a competitor is the right move	12
6	Strategic case methodology	13
6.1	Why the framework is not the hard part	13
6.2	Hypothesis-led reasoning: answer first, then test	14
7	Bridges: market structure, unit economics, and ROIC	15
7.1	When the strategy is brilliant and the spreadsheet says no	15
8	Strategic failure modes	17
8.1	How good firms lose	17
9	Integrated strategic diagnosis	18
9.1	The test of whether the tools have become judgement	18
9.2	A protocol for strategic diagnosis	18

1 Industry analysis with Five Forces

1.1 Two restaurants, two destinies

Open a fine-dining restaurant in a wealthy city and an independent pharmacy in the same city. Run both well — same talent, same discipline, same hours. Ten years later the pharmacy owner is comfortably wealthy and the restaurateur, despite a Michelin star, is exhausted and barely solvent. The difference is not effort or skill. It is that the two firms sit in industries with utterly different **structures**: restaurants face near-zero barriers to entry, fickle customers with thousands of substitutes, and ruinous rivalry; pharmacies are licensed, geographically protected, and sell something with few substitutes to buyers who cannot easily shop on price. The **industry**, before any single decision the owner makes, has already set the ceiling on what the firm can earn.

Michael Porter's Five Forces is the framework that makes this ceiling legible. It answers one question: **how much profit does the structure of this industry permit an average competent incumbent to earn, and why?**

1.2 The five forces

Definition 1 (the five forces) Industry profitability is determined by five structural forces. The stronger a force, the more profit it drains away from incumbents.

1. **Rivalry among existing competitors** — the intensity of competition among firms already in the industry (price wars, advertising battles, capacity races).
2. **Threat of new entrants** — how easily outsiders can enter and compete away profit, governed by **barriers to entry**.
3. **Bargaining power of suppliers** — suppliers' ability to raise input prices or cut quality, capturing margin from the industry.

4. **Bargaining power of buyers** — buyers' ability to force prices down or demand more for the same price.
5. **Threat of substitutes** — alternatives, **outside** the industry, that meet the same underlying need (video calls substitute for airline business travel).




Figure 1: Porter's Five Forces. Rivalry sits at the centre; the four peripheral forces each press inward, raising the intensity of competition and lowering the profit the structure permits. The diagnostic task is not to label all five but to find the **one** that binds hardest.

1.3 Quantify, do not recite

The beginner fills five boxes with “high / medium / low” and stops. That is a checklist, not an analysis. Each force must be tied to **measurable indicators** for the specific sector, so that the verdict is defensible to someone who knows the industry.

Indicators that quantify each force

- **Rivalry**: number of competitors and concentration (CR4, the four-firm concentration ratio, or the Herfindahl index); industry growth rate; fixed-cost intensity; exit barriers.
- **Entry**: minimum efficient scale; capital required; the size of scale economies incumbents enjoy; regulatory or licensing barriers; access to distribution.
- **Supplier power**: supplier concentration relative to the industry; share of the industry's cost the input represents; switching costs; threat of forward integration.
- **Buyer power**: buyer concentration; share of the buyer's spend; product standardisation; switching costs; threat of backward integration.
- **Substitutes**: the **price–performance ratio** of the best substitute and the buyer's switching cost to it.

Definition 2 (the dominant force) The **dominant force** is the single force that most constrains profitability in the sector under study. Identifying it is the whole point of the analysis: it tells you where strategic effort must go. A firm that pours money into differentiating its product when the binding constraint is **supplier power** is solving the wrong problem with the right energy.

Worked example — Five Forces on commercial aviation. Why do airlines collectively destroy capital decade after decade? Run the forces quantitatively. **Rivalry:** many carriers, near-identical product, huge fixed costs (planes fly whether full or not) and high exit barriers (leased fleets, slots) — so any downturn triggers fare wars to fill seats. **Buyer power:** price-transparent comparison sites turned a once-captive flyer into a ruthless shopper; buyer power is high. **Supplier power:** two airframe makers (Boeing, Airbus), a handful of engine makers, unionised pilots, and monopoly airports — supplier power is high on every input. **Substitutes:** rail and video-conferencing on key routes. **Entry:** deregulation cut barriers. Four of five forces are strong; the **dominant** ones are intense rivalry fused with high supplier power. The structure, not management incompetence, is why the average airline earns less than its cost of capital. Compare this with the airport that serves them — a local monopoly — and you see the structure, not the sector glamour, decides who gets rich.

1.4 Where the framework breaks

Common pitfall Five Forces was built for the industrial economy of the late 1970s, and it carries three assumptions that fail in modern markets: that the **industry is well-defined**, that firms, suppliers, and buyers are **distinct roles**, and that advantage comes from **industry structure** rather than from a firm's own assets. In **platform, network-effect, and winner-takes-most** markets all three break:

- the same actor is simultaneously supplier and buyer (a driver and a rider, a host and a guest — **multi-sided markets**);
- value depends on the **size of the user base**, not on industry structure, so increasing returns concentrate the market on one or two winners;
- “the industry” has no stable boundary — is Amazon in retail, logistics, cloud, or advertising?

In these contexts a sixth force, **complementors** (firms whose product makes yours more valuable), often matters more than any of the original five. **State the boundary explicitly before drawing any recommendation** — a clean Five Forces analysis of a platform business is worse than no analysis, because it is confidently wrong.

The discipline, then, is two-step: run the forces quantitatively to find the dominant one, **and** state where the framework stops describing your sector before you let it drive a decision.

Exercises

1. Choose a real sector (e.g. budget gyms, smartphone handsets, grocery retail). Score each of the five forces against named indicators, and identify the dominant force with the evidence that makes it dominant.
2. Take a platform business (ride-hailing, a marketplace, an app store). Show precisely which of the three Five-Forces assumptions break, and which sixth force (complementors) you would add.
3. Using your two analyses, formulate an entry-or-not recommendation for each sector, and state explicitly the structural reason behind each verdict.

2 Resource-based view and VRIO

2.1 Why two firms in the same industry earn different returns

Five Forces explains why the average airline is poor — but Southwest, Ryanair, and later a few others were, for long stretches, **consistently profitable** in that same brutal industry. If structure set the ceiling, how did they break through it? Because profitability has **two** sources: the industry you are in (the **outside** view of Five Forces) and the resources you control that rivals do not (the **inside** view). The **resource-based view** (RBV) is that inside view, and VRIO is its test.

Definition 3 (resources and capabilities) Under the **resource-based view**, a **resource** is a tangible or intangible asset a firm controls (a brand, a patent, a location, a data set, cash). A **capability** is the firm's capacity to deploy resources to an end (the ability to design products customers love, or to run a supply chain at low cost). Competitive advantage comes from resources and capabilities that rivals cannot cheaply acquire or replicate.

2.2 VRIO is a durability test, not a strength test

The natural question — “is this resource good?” — is the wrong one. The strategic question is whether the advantage it confers **survives contact with competitors**. VRIO asks four questions in a fixed order, and the order matters: each one is a gate the resource must pass to reach the next.

Definition 4 (VRIO) For a given resource or capability, ask:

- **Valuable** — does it let the firm exploit an opportunity or neutralise a threat (raise willingness-to-pay or lower cost)?
- **Rare** — is it controlled by only one or few competitors?
- **Inimitable** — is it costly for others to imitate, substitute, or buy (because of history, causal ambiguity, social complexity, or legal protection)?
- **Organised** — is the firm organised — structure, processes, culture — to actually capture the resource's value?

Only a resource that is valuable, rare, inimitable, **and** exploited by the organisation yields a **sustained** competitive advantage.




Figure 2: The VRIO ladder. A resource is run through four gates in order. Failing a gate caps the advantage at the level shown below it; only a resource that clears all four delivers **sustained** advantage. Note that **Inimitable** is the gate that separates a temporary advantage from a durable one — it is where most resources die.

2.3 The mistake that VRIO is designed to prevent

Common pitfall The two recurring errors are both failures of the **durability** lens. First, students stop at **Valuable** and **Rare** and declare victory — but a valuable, rare resource that rivals can copy gives only a **temporary** advantage (the gate that matters most, Inimitability, is the one they skipped). Second, they treat VRIO as a permanent verdict. A resource that is inimitable **today** can be made worthless overnight by a technological or regulatory shift: Kodak's film-chemistry mastery was deeply inimitable right up to the moment digital sensors made the **Valuable** gate fail. VRIO is a snapshot of durability **under current conditions**; re-run it whenever the environment moves.

Worked example — VRIO on a coffee chain's real estate. A coffee chain's most-cited asset is its brand — but apply VRIO and the brand scores **valuable, rare, but imitable** (rivals build brands too): a temporary, erodible edge. Now test a less glamorous asset: its portfolio of **long-term leases on prime corner locations** signed years ago at low rents. Valuable (footfall and cost), rare (there is one best corner per neighbourhood and it is taken), inimitable (a rival literally cannot acquire a site that is already leased for fifteen years — **physical scarcity plus history**), and organised (the firm's site-selection and operations capture it). The real sustained advantage is the boring lease book, not the celebrated brand. VRIO's value is exactly this: it redirects attention from what looks impressive to what is actually durable, and tells you the next move — **keep locking up irreplaceable sites**.

Exercises

4. Pick a firm you admire. List its three most-cited resources, run each through the four VRIO gates, and identify which one (if any) is a source of **sustained** advantage.
5. Find a resource that passed VRIO a decade ago but fails a gate today. Name the gate that flipped and the environmental change that flipped it.
6. For the firm's strongest resource, recommend the next strategic move that would **deepen** the gate it is most at risk of failing.

3 Value chain and activity systems

3.1 Why you cannot copy Ryanair by copying one thing

Suppose you want to beat Ryanair. You buy the same planes, fly the same routes, and match its fares. You lose money and they do not. Why? Because Ryanair's low cost does not come from any **single** decision you can imitate — it comes from **hundreds of decisions that reinforce each other**: one aircraft type (cheap maintenance and training), secondary airports (low fees, fast turnarounds), no free meals, no seat allocation, point-to-point routes, online-only sales, ancillary fees. Copy one and you get nothing; the saving lives in the **fit between them**. To see this you must take the firm apart activity by activity — Porter's **value chain** — and then see how the parts lock together into an **activity system**.

Definition 5 (the value chain) Porter's **value chain** decomposes a firm into the discrete activities through which it creates value, so the source of cost or differentiation can be located precisely.

- **Primary activities:** inbound logistics, operations, outbound logistics, marketing & sales, and service — the activities that physically create, sell, and support the product.
- **Support activities:** firm infrastructure, human-resource management, technology development, and procurement — which underpin the primary ones.

The gap between the total value the chain produces and the total cost of the activities is the firm's **margin**.




Figure 3: Porter's value chain. Support activities (top) underpin the sequence of primary activities (bottom) that create, deliver, and service the product. The **margin** is what survives after every activity's cost — and advantage means doing some activities at lower cost, or in a more valued way, than rivals.

3.2 From a list of activities to a system

Definition 6 (activity system and strategic fit) An **activity system** is the set of a firm's activities together with the **fit** — the mutual reinforcement — between them. Fit makes a strategy hard to imitate: a rival cannot copy one high-performing activity and get the result, because each activity's contribution depends on all the others. The system is defensible not because any single piece is unique but because **the whole interlocking set must be copied at once**, and copying it halfway often makes a rival worse off than not copying at all (**the straddle penalty**).

Worked example — IKEA's interlocking system. IKEA's low prices are not one decision. Flat-pack design lowers shipping and storage **and** shifts assembly labour to the customer **and** enables self-service warehouses **and** removes the need for sales staff **and** allows huge out-of-town stores with cheap land **and** the famous one-way layout that turns a furniture trip into a half-day outing that raises basket size. Each choice makes the others pay off more. A traditional furniture retailer that “adds flat-pack” to a downtown, full-service, salesperson-staffed store gets the costs of neither model and the benefits of neither — the **straddle penalty**. That is why IKEA's system, mapped on paper for anyone to see, has resisted imitation for decades. The map is public; the **fit** is the moat.

3.3 Static endowments versus dynamic capabilities

Definition 7 (dynamic capabilities) A **static resource endowment** is what a firm holds at a point in time — a patent portfolio, a brand, a factory. A **dynamic capability** is the firm's higher-order capacity to **build, integrate, and reconfigure** its resources and

activities as the environment changes — to sense a shift, seize the opportunity, and transform the activity system accordingly. Static endowments win the present; dynamic capabilities win the **next** environment.

Common pitfall Do not mistake a **stock** for a **capability**. “We have great engineers” names an endowment; “we can repeatedly turn engineering talent into products that fit a market we did not previously serve” names a dynamic capability. Endowments depreciate and can be made obsolete by a shift the firm cannot answer; dynamic capabilities are what let a firm **survive** the shift. The deconstruction of a firm must distinguish the two, because a firm rich in endowments but poor in dynamic capability (late-stage Kodak, Nokia) looks invincible right until the environment turns.

Exercises

7. Deconstruct the activity system of IKEA, Zara, or Ryanair: list 8–10 activities and draw the reinforcing links between at least five pairs. Explain why copying any single activity yields no advantage.
8. Identify, for that firm, one **static endowment** and one **dynamic capability**, and argue which would matter more if the industry were disrupted tomorrow.
9. Describe a concrete “straddle” a rival attempted (or might attempt) against your firm, and explain via the activity map why it left the rival worse off.

4 Competitive positioning

4.1 Three ways to not be like everyone else

A firm that offers the same thing as everyone, at the same price, to the same people, earns the industry-average return — which Chapter 1 showed can be miserable. To earn more, a firm must occupy a **distinct position**. There are three classical ways to think about position, and they answer three different questions: **how do we compete?** (generic strategies), **can we stop competing on the usual terms?** (Blue Ocean), and **who are we actually competing against?** (strategic groups).

4.2 Porter’s generic strategies

Definition 8 (generic strategies) A firm builds advantage on one of two bases — **lower cost** or **differentiation** — across either a **broad** or a **narrow** competitive scope, giving four positions:

- **Cost leadership** — the lowest-cost producer serving a broad market;
- **Differentiation** — a broadly-targeted offer buyers value as unique and pay a premium for;
- **Cost focus** — cost leadership within a narrow segment;
- **Differentiation focus** — differentiation within a narrow segment.

A firm that commits to none — neither low-cost nor genuinely differentiated — is **stuck in the middle** and earns the worst of both.

		Source of advantage	
		lower cost	differentiation
Scope	broad	Cost Leadership	Differentiation
	narrow	Cost Focus	Differentiation Focus

Figure 4: Porter’s generic strategies. Advantage rests on lower cost or differentiation, pursued across a broad or narrow scope. The danger zone is the **middle** — trying to be cheap and premium at once usually achieves neither.

Common pitfall “Stuck in the middle” is the most common positioning failure, and it rarely feels like a choice — it is the **accumulated drift** of saying yes to everything. A cost leader that adds premium features to chase a richer segment, or a differentiator that cuts quality to chase volume, ends up with a cost base too high to win on price and a product too plain to command a premium. Note the boundary condition, though: some firms (Toyota at its peak, certain retailers) appear to achieve both because a superior **activity system** lowers cost **and** raises quality simultaneously. The rule is not “you can never do both” — it is “you cannot do both by **compromising**; only a genuinely better system earns the exception.”

4.3 Blue Ocean: making the competition irrelevant

Definition 9 (the strategy canvas and value curve) The **strategy canvas** (Blue Ocean) plots, on the horizontal axis, the **factors an industry competes on**, and on the vertical axis, the **level** each competitor offers on each factor. Joining a competitor’s points gives its **value curve**. A **red-ocean** strategy competes on the existing factors (curves bunch together); a **blue-ocean** move **reshapes the curve** — applying the **Eliminate–Reduce–Raise–Create** grid to drop and lower some factors while raising and inventing others — to open uncontested space and make the competition irrelevant.




Figure 5: A strategy canvas. Incumbents share a value curve (grey) — they compete on the same factors. The blue-ocean mover (purple) **eliminates and reduces** costly factors buyers do not truly value (lounges, meals) and **raises and creates** new ones (fun, frequency), tracing a divergent curve in uncontested space. The shape, not the height, is the strategy.

4.4 Strategic groups: who you actually fight

Definition 10 (strategic groups) A **strategic group** is a set of firms in an industry pursuing similar strategies along the dimensions that matter (e.g. price level and breadth of range, or quality and geographic reach). A **strategic-group map** plots firms on two such dimensions; the firms in your group are your **direct rivals**, **mobility barriers** (the cost of moving between groups) protect each group, and **empty regions** of the map are candidate uncontested positions — or graveyards where no viable position exists.

Worked example — Reading a disruption on the map. Map the car industry on **price** (vertical) and **technology/software intensity** (horizontal). The legacy groups cluster on the left: budget, mid-market, and luxury marques, each a strategic group with its own customers and mobility barriers. Tesla enters in the empty **high-price, high-software** corner — not competing inside any existing group but defining a **new** one, then forcing every legacy group to migrate rightward toward software, which their activity systems (built around mechanical engineering, not over-the-air updates) are poorly organised to do. The map does two things at once: it shows that Tesla's initial rivals were not “all carmakers” but the empty quadrant it chose, and it predicts **which** incumbents struggle — those whose mobility barrier to the software dimension is highest. A disruption is exactly a new group opening where the map was empty.

Exercises

10. For a chosen industry, classify three named firms by generic strategy, and identify one that is “stuck in the middle” — with the cost and price evidence that shows it.
11. Build a strategy canvas for an industry with 5–7 competing factors. Plot two incumbents, then design a blue-ocean value curve using the Eliminate–Reduce–Raise–Create grid, naming each moved factor.
12. Draw a strategic-group map for that industry on two dimensions, place the firms, mark the mobility barriers, and evaluate a disruption currently in progress: which group does the entrant threaten, and which is structurally least able to respond?

5 Competitive dynamics

5.1 Strategy as a move in a game the rival also plays

Everything so far analysed a firm against a **static** backdrop. But your rival is not a backdrop — it **reacts**. Cut your price and a competitor cuts deeper; enter its market and it retaliates in yours. Competitive dynamics treats strategy as what it is: a sequence of **moves and counter-moves** between thinking opponents. This is where BUS13-4's game theory — payoff matrices, Nash equilibrium, the prisoner's dilemma — stops being an abstraction and becomes the daily reality of pricing, capacity, and entry decisions.

Definition 11 (the levers of competitive dynamics)

- **Signalling** — actions or announcements that communicate intentions to rivals (a pre-announced capacity expansion warns entrants off).
- **Commitment** — an **irreversible** move that constrains your own future options and is therefore **credible**: burning a bridge so the rival knows you cannot retreat.

- **Reputation** — the accumulated track record that makes your signals and threats believable; a firm known to always retaliate need not retaliate as often.
- **Tit-for-tat** — the strategy of cooperating until the rival defects, then retaliating in kind, then forgiving: the robust winner of repeated prisoner’s-dilemma play.
- **First-mover vs fast-follower** — moving first to pre-empt, versus moving second to learn from the pioneer’s costly mistakes.

5.2 Why a threat must be credible to work

A threat or promise only changes a rival’s behaviour if the rival **believes** it, and belief comes from **commitment** and **reputation**, not from words. This is the deep idea: sometimes **reducing your own options** makes you stronger, because it makes a threat credible.

		Firm B	
		B: hold	B: cut
Firm A	A: hold	3 , 3	0 , 4
	A: cut	4 , 0	1 , 1

Figure 6: A pricing prisoner’s dilemma (payoffs: Firm A, Firm B). Both **hold** yields (3, 3); both **cut** yields (1, 1); the dominant-strategy Nash equilibrium is the outlined (cut, cut) — each firm defects fearing the other will. **Tit-for-tat** in repeated play, backed by a credible retaliation **reputation**, is how rational firms sustain the cooperative (hold, hold) outcome instead.

Definition 12 (first-mover advantage and its conditions) First-mover advantages — pre-emption of scarce assets, buyer **switching costs**, **learning-curve** cost leads, and **network effects** — are real only under specific conditions. When they are absent or weak, the **fast-follower** wins: it lets the pioneer spend to educate the market and prove demand, then enters with a better product, a lower cost, and none of the pioneer’s mistakes. The question is never “is it better to move first?” in the abstract — it is “does **this** market reward the first move, given switching costs, network effects, and how fast imitation is?”

5.3 When ignoring a competitor is the right move

Common pitfall The instinct to **respond to every competitive move** destroys value. Retaliation is costly — a price cut to punish a rival hits your **whole** book of business, not just the contested customers — so it is only worth it when the threat is material and the response is proportionate. A competitor can be **rationally ignored** when: it targets a segment you do not value or cannot profitably serve; responding would damage your core more than the rival’s incursion does (a premium brand that price-matches a discounter cheapens itself); or the rival is not viable and will exit unaided. The disciplined question is not “can we respond?” but “**does responding leave us better off than not?**” — and often the answer is no.

Worked example — A premium brand that should not fight on price. A discount entrant undercuts a premium incumbent on price. The incumbent's instinct is to match. But run the dynamics: the entrant targets price-sensitive buyers the premium brand never profitably served; matching prices would (a) signal to the incumbent's **own** loyal customers that the high price was never justified, eroding willingness-to-pay across the whole base, and (b) drag the incumbent into a cost game its differentiated activity system is not built to win. The value-maximising move is to **ignore the price** and instead reinforce the differentiation the discounter cannot match. Ignoring is not passivity — it is the recognition that the only winning move in this particular game is **not to play it on the rival's terms**.

Exercises

13. Take a named rivalry (e.g. two consoles, two airlines, two streaming services). Identify one instance each of signalling, commitment, and reputation, and explain how each shaped the other firm's behaviour.
14. Build a 2×2 payoff matrix for a capacity or pricing decision in that rivalry, find the Nash equilibrium, and explain how repeated play plus tit-for-tat could sustain a better joint outcome.
15. Decide whether the firm is better placed as first-mover or fast-follower in its next product battle, citing the specific conditions (switching costs, network effects, imitation speed) that drive your answer.
16. Find a real case where a firm **correctly ignored** a competitor, and justify the choice with the cost of the response it avoided.

6 Strategic case methodology

6.1 Why the framework is not the hard part

By now you hold a full toolkit. The surprise of real strategy work is that **knowing the frameworks is the easy 20%**. The hard 80% is reasoning well when you have a messy, ambiguous problem, missing data, and an hour before you must recommend. A student who can recite Five Forces but cannot **structure an unfamiliar problem under time pressure** is useless in the room. This chapter teaches the reasoning discipline — MECE, issue trees, hypothesis-led analysis — that turns frameworks into judgement. It is taught as **strategic practice**, not as consulting-interview preparation; the goal is clear thinking, not a job offer.

Definition 13 (MECE) A decomposition is **MECE — Mutually Exclusive, Collectively Exhaustive** — when its parts **do not overlap** (mutually exclusive) and **together cover the whole** (collectively exhaustive). MECE is the quality test for any breakdown: overlap causes double-counting and muddled ownership; gaps cause missed causes. “Raise revenue” decomposed into “new customers / more per existing customer / win back lost customers” is MECE; “new customers / online channel / discounts” is neither.

Definition 14 (the issue tree) An **issue tree** decomposes a central question into sub-questions, **MECE at every level**, branching until each leaf is a question you can actually answer with available data or a cheap test. It turns one overwhelming question (“why

are profits falling?") into a navigable structure where each branch can be confirmed or eliminated.




Figure 7: An issue tree for “why is profit falling?”. The first split — revenue versus cost — is MECE because profit **is** revenue minus cost. Each branch splits again, MECE at every level, until leaves are directly testable. Hypothesis-led work then **prioritises** the branch most likely to carry the answer rather than investigating all equally.

6.2 Hypothesis-led reasoning: answer first, then test

Definition 15 (hypothesis-led reasoning) Hypothesis-led reasoning starts from a **candidate answer** — a stated, falsifiable hypothesis — and then gathers only the evidence that would **confirm or refute it**, rather than collecting all possible data before forming a view. It inverts the naïve order: not “study everything, then conclude” but “conclude provisionally, then test the conclusion.” This is what makes a structured recommendation possible under information and time constraints — you investigate the few things that would change your mind, not the many that would not.

Common pitfall Two failure modes bracket good case reasoning. The first is **boiling the ocean**: gathering data with no hypothesis, drowning in facts, and running out of time before reaching a recommendation — the analyst confuses **thoroughness** with **progress**. The second is **confirmation bias**: forming a hypothesis and then seeking only evidence that flatters it. The discipline that avoids both is to ask, of each hypothesis, “what is the **single cheapest test whose result would change my mind?**” — and to run **that** test first. A hypothesis you are not willing to kill is a prejudice, not an analysis.

Worked example — Diagnosing a sales decline under time pressure. A retailer’s sales are down 15% and you have one day. **Boiling the ocean** would pull every report on every store. Instead, build the issue tree (volume vs price; by region, by channel, by product) and form a **hypothesis**: “the decline is concentrated in physical stores in two regions where a new competitor opened.” That hypothesis is cheap to test — pull same-store sales by region for those two and a control region. If the drop is uniform everywhere, the hypothesis is **refuted in twenty minutes** and you move to the next branch (a pricing or product issue) without having wasted the day. The structure plus the willingness to be wrong fast is what delivers a defensible recommendation by evening.

Exercises

17. Take a real decision (“should this café open a second location?”). Build a MECE issue tree at least two levels deep, and check each split for overlaps and gaps.
18. For a “why are profits falling?” problem in a chosen firm, state a falsifiable hypothesis and name the single cheapest test that would refute it.
19. Produce a one-page, hypothesis-led recommendation for a novel strategic problem, stating your assumptions explicitly and marking which would most change the answer if wrong.

7 Bridges: market structure, unit economics, and ROIC

7.1 When the strategy is brilliant and the spreadsheet says no

A strategy that wins the competitive argument can still **destroy value** — and a strategy the frameworks dislike can **create** it. The frameworks of Chapters 1–6 reason about **position**; they do not, by themselves, tell you whether a position **pays**. For that you must connect strategy down to the two layers beneath it: the **microeconomic market structure** (BUS13-4), which sets what profit is sustainable, and the **unit economics and ROIC** (BUS23-2, BUS23-3), which measure whether the advantage actually earns its keep. The most valuable thing this chapter teaches is the **conflict** — the decision on which the strategic logic and the financial logic point opposite ways.

Definition 16 (market structure as a feasibility constraint) An industry’s **market structure** — perfect competition, monopolistic competition, oligopoly, or monopoly (BUS13-4) — constrains which strategies are feasible and what profit is sustainable. In near-perfect competition, no positioning escapes zero economic profit for long; in an oligopoly, competitive dynamics (Chapter 5) dominate; only where structure permits durable market power can a differentiation premium **persist**. A recommendation that ignores the structure recommends a profit the structure forbids.

Definition 17 (unit economics and ROIC as strategic-health indicators) **Strategic health** is ultimately financial. **LTV/CAC** (BUS23-2) tests whether an advantage converts into customers worth more than they cost; **ROIC** (BUS23-3) tests whether the firm earns more on its invested capital than that capital costs ($\text{ROIC} > \text{WACC}$). An advantage that does not eventually show up as LTV/CAC above its threshold and ROIC above WACC is, financially, not an advantage at all — it is a story.

Financial logic (ROIC vs WACC) →		
Strategy logic ↑	Strategic darling attractive, ROIC < WACC <i>the trap</i>	Win-win attractive, ROIC > WACC <i>invest</i>
	Avoid weak, ROIC < WACC <i>exit</i>	Hidden gem weak, ROIC > WACC <i>harvest quietly</i>

Figure 8: Where strategic and financial logic meet. The diagonal cells agree (invest / exit). The **off-diagonal** cells are where this course earns its keep: the strategically attractive market that destroys capital (top-left — the value trap that sinks many “must-win” expansions) and the unglamorous business that quietly out-earns its cost of capital (bottom-right).

Common pitfall The seductive error is the **strategic darling**: a market everyone agrees is attractive — growing, prestigious, “strategically essential” — that nonetheless earns ROIC below WACC for everyone in it (early streaming, much of e-commerce delivery, parts of space launch). Strategy logic shouts “be there”; financial logic shows every euro invested returns less than it costs. The mature analyst holds **both** numbers at once and names the conflict explicitly, rather than letting the more exciting narrative silently win. The reverse trap also exists — abandoning a dull but cash-generative business because it lacks a glamorous strategic story.

Worked example — When strategy says enter and ROIC says don't. A profitable consumer brand considers entering meal-kit delivery — strategically attractive (fast-growing, adjacent, on-trend). Strategy frameworks are encouraging: a differentiation-focus position looks open. But run the unit economics: CAC is high and rising (Chapter 1’s buyer power — customers multi-home across services), churn is brutal, and the capital tied up in fulfilment depresses ROIC far below the firm’s WACC. The honest recommendation names the conflict: **the position is defensible but the economics are not**, and on this decision the financial logic should win unless the firm can find a structural CAC or churn advantage. Surfacing that conflict — not resolving it by ignoring one side — is the deliverable.

Exercises

- 20.** For a strategic recommendation you have made earlier in this course, name the industry’s market structure and explain how it constrains the profit your recommendation assumes.
- 21.** Take a “must-win” expansion a real firm pursued. Estimate LTV/CAC and whether ROIC plausibly beats WACC. Place it on the strategy-vs-finance quadrant.
- 22.** Construct a decision on which strategy logic and financial logic genuinely conflict, state both verdicts, and argue which should prevail and why.

8 Strategic failure modes

8.1 How good firms lose

Most strategic failure is not a dramatic defeat by a brilliant rival. It is **self-inflicted** and **slow**, and it comes in three recognisable shapes. Learning to name them on a real case — and to state the corrective move — is the diagnostic skill that sets up the S6 capstones (Tech & Platform Strategies and Strategic Allocation), where these failures play out at platform scale and across whole portfolios.

Definition 18 (the three failure modes)

- **Imitation** — copying a rival's **position** without copying the **activity system** that makes it work (Chapter 3), yielding parity at best and the straddle penalty at worst.
- **Strategic drift** — the gradual divergence of a firm's strategy from its changing environment, as small accommodations accumulate until the strategy no longer fits the world it operates in.
- **Unjustified diversification** — entering businesses with no value-creating link to the firm's resources or activities, destroying value through complexity and the loss of focus (often dressed up as “synergy”).




Figure 9: Strategic drift. The environment's demands (purple) rise continuously, while the firm adapts in small, comfortable increments (grey staircase) that each feel reasonable. The accumulating **gap** (dashed) is the drift — invisible quarter to quarter, fatal over years. The corrective move is a discontinuous **transformation**, not another small step.

Common pitfall Strategic drift is dangerous precisely because **every individual step is defensible**. No one decides to fall behind; each year the firm makes a small, sensible accommodation — protect the legacy product, please the existing customer, hit the quarterly number — and the sum of sensible steps is a strategy that no longer fits. By the time the gap is obvious in the financials, the dynamic capabilities (Chapter 3) needed to close it have often atrophied. The warning sign is not bad results — it is **good results from an increasingly outdated model**, the calm before Kodak.

Worked example — Three failures, three corrective moves. **Imitation:** a retailer copies a discounter's low prices but keeps its full-service stores — parity prices, premium costs, the straddle penalty. **Corrective move:** choose a coherent position and rebuild the activity system to fit it, rather than copying a surface. **Strategic drift:** a film-camera leader keeps optimising film as digital rises, each step protecting today's profit — until the environment

has moved entirely. **Corrective move**: a discontinuous reallocation of resources to the new technology **before** the legacy business forces it. **Unjustified diversification**: a beverage firm buys a film studio for “synergy” that never materialises. **Corrective move**: divest and return capital — the value-creating link was never there. Each failure has a **different** corrective move, which is why naming the mode precisely matters.

Exercises

23. Diagnose one named company’s failure as imitation, strategic drift, or unjustified diversification, citing the specific evidence (cost structure, timeline, or acquisition record) that identifies the mode.

24. For that case, state the corrective strategic move and explain why a **smaller** version of it (one more incremental step, a partial copy) would not have worked.

25. Find a firm currently drifting. Identify the small accommodations accumulating into the gap, and propose the discontinuous move that would close it.

9 Integrated strategic diagnosis

9.1 The test of whether the tools have become judgement

Every chapter has handed you one instrument. The capstone plays them together: take a **real** company, run the full external and internal analysis, and produce a **quantified, defensible recommendation** — then defend it before a jury that may include a practitioner who knows the industry from the inside. The jury is the point. A diagnosis that cannot survive an expert’s “that is not how this industry works” was built on frameworks recited, not judgement earned. This is the moment the course has been building toward.

9.2 A protocol for strategic diagnosis

The diagnosis protocol

1. **External analysis.** Run **Five Forces** quantitatively, identify the dominant force, and state explicitly where the framework breaks for this sector (Chapter 1). Map the **strategic groups** and locate the firm and its direct rivals (Chapter 4).
2. **Internal analysis.** Apply **VRIO** to find the firm’s central resource and test its durability (Chapter 2). Deconstruct the **value chain / activity system** and identify the fit that is hard to imitate, distinguishing static endowments from dynamic capabilities (Chapter 3).
3. **Positioning.** Place the firm with the **generic strategies** and the **Blue Ocean canvas**, and assess any disruption in progress (Chapter 4).
4. **Dynamics.** Model the key **rivalry** — signalling, commitment, reputation, first-mover vs fast-follower — and judge when to respond and when to ignore (Chapter 5).
5. **Financial reconciliation.** Link the recommendation to the **market structure** and to **LTV/CAC and ROIC**, and surface explicitly any conflict between strategic and financial logic (Chapter 7).
6. **Failure-mode check.** Test the firm for imitation, drift, and unjustified diversification, and name the corrective move if a mode is present (Chapter 8).
7. **Recommendation.** State a single, **quantified, falsifiable** strategic move, with its assumptions explicit and the evidence that would confirm or refute it (Chapter 6’s hypothesis-led discipline).

Common pitfall The capstone failure is the **framework parade**: running every tool, producing a fat deck, and recommending nothing decisive — five analyses that do not converge on one move. A diagnosis is judged not by how many frameworks appear but by whether they **converge** on a recommendation that is specific, quantified, and falsifiable. “Improve competitiveness” is not a recommendation; “exit segment X within 18 months and redeploy the freed capital into the data-loop capability, because LTV/CAC there is below 1 and the activity-system fit is weakest” is. Before the jury, ask of your own deck: **does every framework I ran change the recommendation? If not, cut it.**

Worked example — From analysis to a single defensible move. A diagnosis of a mid-market grocer might run: Five Forces shows the dominant force is **buyer power** fused with discounter rivalry (margins structurally thin); strategic-group mapping shows the grocer **stuck in the middle** between hard-discounters and premium chains; VRIO finds its one durable asset is a dense network of small urban stores (rare, inimitable by physical scarcity); the value chain shows its cost base is built for a format it has half-abandoned. The frameworks **converge**: the grocer is drifting toward the middle, its real advantage is proximity, and its economics (ROIC below WACC in large-format stores) confirm it. The single move: **commit to a convenience-focus position built on the store network, exit large-format, and redeploy capital** — quantified by the ROIC gap, falsifiable by a pilot. One move, every framework pointing at it, defensible to a grocer who knows the business.

Exercises

26. Produce a complete strategic diagnosis of a real company following the seven-step protocol: external analysis (Five Forces with its boundary, plus strategic groups) and internal analysis (VRIO plus value chain), with each force and resource tied to evidence.
27. Formulate a single quantified, falsifiable recommendation, link it to market structure and to LTV/CAC and ROIC, and state explicitly any conflict between strategic and financial logic.
28. Anticipate the three hardest questions a practitioner on the jury would ask, and prepare an evidence-based answer to each — including the one piece of evidence you most wish you had.

Competitive strategy, in one paragraph. Advantage is not something a firm owns; it is a position it holds against thinking rivals, and the only question that matters is what keeps it in place and for how long. Read the **outside** with Five Forces — quantitatively, to find the dominant force, and honestly, by stating where the framework breaks in platform and network-effect markets. Read the **inside** with VRIO as a durability test and with the value chain as an **activity system** whose fit — not any single part — is the real moat, distinguishing static endowments from the dynamic capabilities that survive a shift. **Position** with generic strategies (never stuck in the middle), the Blue Ocean canvas (reshape the value curve, do not climb it), and strategic-group maps (know who you actually fight and where the map is empty). Treat rivalry as a **game**: signal, commit, build reputation, play tit-for-tat — and know when ignoring a competitor is the optimal move. Reason under pressure with **MECE issue trees** and **hypothesis-led** analysis: conclude provisionally, then run the cheapest test that could prove you wrong. Always reconcile strategy with **market structure** and with **LTV/CAC and ROIC**, and name out loud the decisions where strategic and

financial logic conflict. Watch for the slow, self-inflicted failures — imitation, drift, unjustified diversification — and know each one's corrective move. Then diagnose a real company end to end and defend one quantified, falsifiable recommendation before people who know the truth. Three questions, always: **what holds this advantage in place, who is trying to take it, and what happens when they do?**