

BUS63-6

Launch & Scale

Go-to-Market strategy for the first eighteen months

Building a commercial machine that survives — before optimising one that already runs.

Albert School — 22 hours

Contents

Preface	2
1 Product mindset versus feature mindset	3
2 Customer discovery and demand-side analysis	4
3 Jobs-to-be-Done and Kano: separating differentiation from hygiene	5
3.1 Jobs-to-be-Done: the demand, not the demographic	5
3.2 Kano: which features are differentiators and which are mere hygiene	5
4 Positioning as strategic commitment	6
4.1 The category decision	7
4.2 Repositioning out of direct comparison	7
5 Go-to-market motion selection	7
6 Pricing at launch	8
6.1 Testing price sensitivity before launch	9
7 Distribution strategy and channels	10
8 Acquisition and product-market fit: the metric shift	11
8.1 Acquisition before PMF	11
8.2 Measuring PMF and the metric shift	11
9 Scaling discipline: avoiding premature scaling	12
9.1 Test-and-learn after launch	13
10 Building the first commercial machine	13
10.1 Founder-led sales	13
10.2 The first three commercial hires, by motion	14
10.3 Structuring the pipeline	14
10.4 From founder-led to team-led	14
11 Growth model design	15
12 The integrated go-to-market plan	15
12.1 Coherence is the deliverable	16

Preface

You arrive at this course able to do four things that most founders cannot. From **BUS13-1** you can treat marketing as a set of irreversible choices — *where to play, how to win, how to allocate* — and write a positioning that constrains decisions rather than decorating a deck. From **BUS23-1** you can run a structured sales cycle (SPIN discovery, BANT qualification), diagnose a pipeline stage by stage, and hold a value-based pricing conversation without flinching. From **BUS23-2** you can read the economic architecture of a business and compute its unit economics — CAC, LTV, LTV/CAC, gross margin, contribution margin, break-even — and stress-test them. From **BUS33-1** you can run growth as evidence: AARRR, channel economics, A/B testing, cohort analysis, and Jobs-to-be-Done as a tool for telling structural demand from fashion.

This course does something specific with those tools. Every one of them was built to *optimise an existing commercial machine* — a firm that already has customers, channels, and a price. The entrepreneurial question is the opposite: you hold a validated product and **no commercial machine at all**, and you must build the first one — positioning, motion, price, the first hires, the first hundred customers — under radical uncertainty, with months of runway, knowing that the dominant way early ventures die is not building the wrong product but **scaling the right one too soon**. Optimisation assumes the machine exists. Here you build it.

We organise go-to-market into four moves. **Demand-side discipline** (Chapters 1–3) — the mindset, the discovery process, and the analysis that decide what the market actually wants.

Strategic commitments (Chapters 4–6) — positioning, motion, and price, the three choices that are expensive to reverse. **The commercial machine** (Chapters 7–10) — distribution, acquisition before and after product–market fit, the discipline that governs when to scale, and the founder-led sales engine with its first hires and pipeline. **Integration** (Chapters 11–12) — the growth model that ties it together and the complete GTM plan defended before a jury.

One fictional venture runs through the whole book. **Lumen** is an early-stage *B2B SaaS* founded by two ex-accountants, Maya and Tomás. Its product collects tax and bookkeeping documents from the clients of small accounting firms — the unglamorous, universally hated chase for “the missing receipts” — and nudges those clients until the file is complete. The product works; a dozen friendly firms love it. Lumen has €400k of pre-seed cash, roughly fifteen months of runway, and no answer yet to a single go-to-market question: who exactly is it for, what is it *against*, how is it sold, what does it cost, and how does the tenth customer become the hundredth. Watch how each decision constrains the next — the positioning narrows the motion, the motion dictates the price structure, the price sets the acquisition channels you can afford, and the channels determine who you hire first. That web of constraints *is* go-to-market.

1 Product mindset versus feature mindset

Lumen’s first real argument was not about code. A friendly firm asked for a built-in e-signature; another wanted a payroll module; a third wanted the document chase to also reconcile bank statements. Tomás, an engineer at heart, started building all three. Maya stopped him with a question that organises this whole course: *which of these makes a prospect who has never heard of us decide to buy?* None of them, it turned out. Prospects bought — or didn’t — on one thing: whether Lumen ended the weekly humiliation of chasing clients for paperwork. The three requested features were comforts for customers who had *already* bought.

Definition 1 (Feature mindset vs product mindset) A **feature mindset** asks “what can we add?” and measures progress in shipped functionality. A **product mindset** asks “what demand are we serving, for whom, and what is the smallest thing that serves it?” and measures progress in *demand captured*. In go-to-market, the product mindset is the **demand-side discipline**: it starts from the customer’s job and works back to what to build, charge, and sell — never the reverse.

This is not a slogan; it is a decision rule that changes four concrete choices.

Worked example — How the mindset rewrites four GTM decisions for Lumen.

- **What to build first.** Feature mindset: e-signature, payroll, reconciliation — a broad suite. Product mindset: make the document chase *ruthlessly* complete and effortless, because that is the demand. Build depth on the job, not breadth around it.
- **What to charge for.** Feature mindset: price per module, nickel-and-dime the suite. Product mindset: charge for the outcome — files completed on time — so price tracks the value the job creates (Chapter 6).
- **Who to hire.** Feature mindset: more engineers to ship the backlog. Product mindset: the first hire is whoever shortens the path from a prospect’s pain to a signed contract (Chapter 10).
- **When to scale.** Feature mindset: scale when the roadmap is “done” — it never is. Product mindset: scale when the demand-side evidence says the job is being served repeatably (Chapter 9).

Common pitfall “Customers are asking for it” is the feature mindset’s most seductive disguise. Existing customers ask for features that deepen *their* use; they are not a sample of the market you have not yet won. Building the requested suite can grow churn-resistance among the converted while doing nothing for acquisition — and consume the very runway you needed to find the demand. The product mindset does not ignore requests; it *ranks* them by whether they serve the core demand that wins new customers, not the comfort of those already aboard. The test (Learning Outcome 1) is to take a real venture and trace how the two mindsets would give it *different* answers on build, price, hire, and scale.

2 Customer discovery and demand-side analysis

A product mindset is useless without evidence about the demand it claims to serve. That evidence comes from **customer discovery** — a structured process, not a round of friendly coffees. The distinction matters because the default failure mode of founder interviews is seeking confirmation: you describe your product, the interviewee is polite, you hear “great idea,” and you learn nothing.

Worked example — Two versions of the same Lumen interview. *Confirmation-seeking:* “We’ve built a tool that automates document collection — would you use something like that?” The accountant, being kind, says yes. Maya leaves elated and none the wiser.

Discovery: “Walk me through the last time a client’s file was late. What did you do? How long did it take? What did it cost you?” The accountant describes three reminder emails, a tense phone call, a missed filing deadline, and a €600 penalty she ate to keep the client. Maya leaves with a quantified pain, a workflow, and a willingness-to-pay anchor — *none of which she suggested.*

Definition 2 (The discovery loop) Customer discovery is a repeating loop: (1) **formulate a falsifiable hypothesis** about who has the demand and why — “small firms lose >€500/year to late client files and would pay to stop it”; (2) **design interviews** that could prove it *wrong* — questions about past behaviour, not future intentions; (3) **extract insight** that *changes the GTM plan* — a segment to drop, a price to test, a positioning to rewrite. An interview that cannot change your plan was not discovery; it was reassurance.

Common pitfall Ask about the *past*, not the *future*. “Would you pay €50/month for this?” measures politeness; “What did the last late file cost you, and what have you already tried or paid to fix it?” measures revealed behaviour (the stated-vs-revealed-preference distinction from BUS13-1). The single most common discovery error is treating “that sounds useful” as demand. Useful is not the bar; *already trying to solve it* is.

The output of discovery is not a transcript pile but a set of *insight-to-GTM translations*: each validated insight names what it changes. “Firms blame themselves, not their clients, for late files” → reposition Lumen as protecting the relationship, not policing the client. “The bookkeeper, not the partner, feels the pain” → target the bookkeeper as champion (Chapter 5’s motion). Running this process — hypotheses, interview design, insight extraction that moves the plan — is Learning Outcome 2.

3 Jobs-to-be-Done and Kano: separating differentiation from hygiene

Discovery tells you the demand exists; two frameworks tell you *which parts of the product create it*. You met both in BUS33-1 — here they become launch-priority tools, used from the market side without crossing into owning the product roadmap.

3.1 Jobs-to-be-Done: the demand, not the demographic

Definition 3 (Jobs-to-be-Done (Christensen)) Customers do not buy products; they “hire” them to make progress on a *job* that arises in their lives. The job is stable; the products hired for it change. Specifying the job — functional, emotional, and social dimensions — defines the true competitive set and the real demand, independent of any feature.

Worked example — Lumen's job is not 'collect documents'. The functional job: *get a complete, accurate client file by the deadline*. But discovery surfaced two more dimensions. Emotional: *stop feeling like a nag and stop dreading deadline week*. Social: *look organised and professional to the client, not chaotic*. Lumen's true competitors are therefore not other software — they are the spreadsheet-and-reminder-emails workaround and the firm's own anxiety. That reframing decides positioning (Chapter 4): Lumen sells calm and professionalism, not a file-upload widget.

3.2 Kano: which features are differentiators and which are mere hygiene

Definition 4 (The Kano model) Features fall into categories by how customer satisfaction responds to their presence:

- **Must-haves (hygiene)** — expected; their presence is unnoticed, their absence is fatal (e.g. Lumen must be secure and not lose files).
- **Performance features** — satisfaction rises roughly linearly with how well they are done (e.g. how few steps a client needs to upload).
- **Delighters (differentiators)** — unexpected, create disproportionate satisfaction and demand (e.g. Lumen auto-detects a missing receipt before the client even submits).

Delighters decay into must-haves over time as the market catches up.

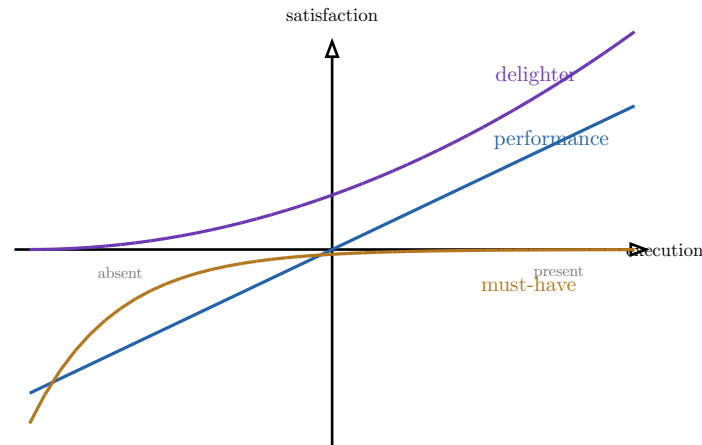


Figure 1: The Kano model. A must-have (orange) buys no satisfaction when present but destroys it when absent; a performance feature (blue) pays back linearly; a delighter (purple) is invisible when absent yet creates outsized demand when present. Launch priority funds must-haves to the threshold, then spends the marginal euro on a delighter — never on polishing a must-have past “good enough”.

Common pitfall The classic launch mistake is over-investing in must-haves and calling it differentiation. Making Lumen’s security “even more secure” wins nothing once it clears the threshold — security is hygiene. Scarce pre-launch effort should clear every must-have to “good enough,” then concentrate on the one or two delighters that actually create demand. Mistaking a must-have for a differentiator is how ventures build a flawless product nobody chooses. The deliverable (Learning Outcome 3) recommends launch priorities *from the market’s perspective* — which features create demand — without claiming supply-side ownership of the roadmap.

4 Positioning as strategic commitment

Discovery and JTBD tell you the demand; positioning is the deliberate choice of *how the market should understand what you are*. BUS13-1 taught positioning as a constraint that survives stress tests. Here we use a sharper instrument built for launch — April Dunford’s *Obviously Awesome* — and confront the single most consequential pre-launch decision hidden inside it: the category.

Definition 5 (Dunford's five components of positioning) Strong positioning is assembled, in order, from five linked pieces:

1. **Competitive alternatives** — what customers would do if you did not exist (for Lumen: chase by email, or a generic file-share).
2. **Unique attributes** — what you have that the alternatives lack (automatic chasing, missing-document detection).
3. **Value** — the benefit those attributes enable (complete files by deadline, no nagging, no penalties).
4. **Target segment** — who cares most about that value (small firms whose partners do the chasing themselves).
5. **Market category** — the frame you place yourself in, which tells the customer what to compare you to and what to expect.

4.1 The category decision

Worked example — Lumen chooses its frame of reference. Lumen could call itself a “*document management tool*.” Instantly it is compared to Dropbox, SharePoint, and a dozen cheaper file-shares — a brutal frame where its automatic chasing looks like a minor feature and its price looks high. Or it can name a category that did not exist on the shelf: a “*client-file completion system for accountants*.” Now the comparison set is “the way you chase files today,” its automation is the defining attribute, and its price is measured against €600 penalties and lost evenings. Same product, same attributes — a different *category* makes them obvious or invisible.

Definition 6 (The category as the most consequential pre-launch choice) The market category sets the customer’s *reference price*, *comparison set*, and *expectations* before they evaluate a single feature. Choosing a crowded category forces you to win on the incumbents’ terms; choosing or creating a category where your unique attributes are the defining ones lets you set the terms. Because everything downstream — pricing anchor, sales narrative, channel — inherits the category, it is the decision hardest to reverse after launch.

4.2 Repositioning out of direct comparison

Common pitfall If you find yourself constantly justifying your price against a cheaper incumbent, the problem is usually the *category*, not the price. The fix is **category repositioning** — moving the frame so the direct comparison no longer applies — not another discount. But repositioning is expensive and confusing if done after a market has filed you under the old category, which is why the category is a launch decision, made deliberately, not a label chosen by default. Constructing the full five-component positioning and defending the category choice — including when and how to reposition — is Learning Outcome 4.

5 Go-to-market motion selection

Positioning decides how the market understands you; the **motion** decides the mechanical path by which a stranger becomes a paying customer. Choosing it wrong is the most expensive ordinary mistake in early-stage GTM, because a motion takes roughly eighteen months to build and prove — so a mismatch is not a tweak, it is a lost runway.

Definition 7 (The three go-to-market motions)

- **Sales-led (SLG)** — humans (founders, then reps) drive the deal: outbound, demos, proposals, negotiation. Fits high price, high complexity.
- **Product-led (PLG)** — the product drives acquisition and conversion: self-serve signup, free trial or freemium, expansion inside the account. Fits low price, low complexity, fast individual value.
- **Marketing-led (MLG)** — content, SEO, community, and brand generate inbound demand that converts with light human touch. Fits mid price, a learning-intensive category, broad audience.

Definition 8 (The four decision criteria) Select a motion against four variables, read together:

1. **Price point** — can a deal’s revenue pay for human selling? Low ACV cannot; it forces PLG or MLG.
2. **Sales complexity** — many stakeholders, customisation, security review → SLG; a single user who “gets it” in minutes → PLG.
3. **Team profile** — do the founders sell, or build? Play to the founding team’s actual edge for the first eighteen months.
4. **Stage** — pre-PMF favours high-touch motions that maximise learning per customer, even if they will not be the scale motion later.

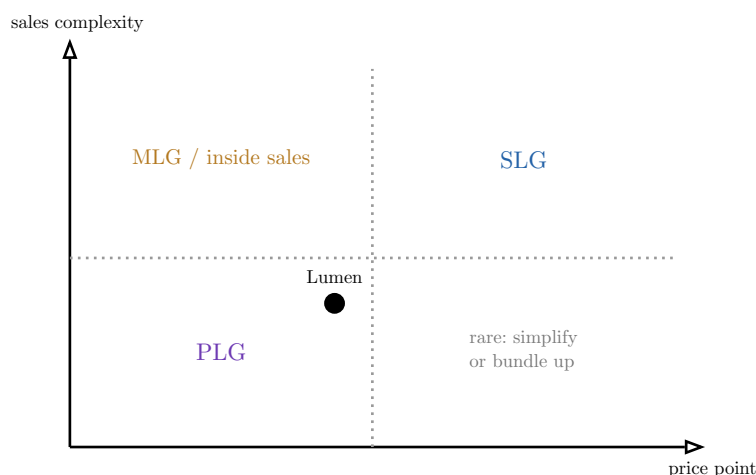


Figure 2: Motion selection by price point and sales complexity. Low price + low complexity → product-led; high price + high complexity → sales-led; the middle band is marketing-led or light inside sales. Lumen sits low-mid on both, near the PLG/MLG seam — a hybrid where founder-led sales teaches the motion before a self-serve layer is added.

Common pitfall The **motion-mismatch failure mode** costs eighteen months. The classic case: a €40/month product sold with a five-call enterprise sales process — the deal can never pay for the selling, and the venture burns its runway proving a motion the economics forbade. The mirror case is just as fatal: a €50,000, twelve-stakeholder enterprise product launched “self-serve,” where no individual can buy and the funnel is silent. Match the motion to the price and complexity *first* the team and stage tell you which version to run *now*. For Lumen — mid price, modest complexity, two founders who can sell — the answer is founder-led sales today (to learn fast), with a product-led self-serve layer added once the message converts reliably. Selecting and justifying a motion against the four criteria, and naming the mismatch it avoids, is Learning Outcome 5.

6 Pricing at launch

Price is the most reversible-looking decision that is in fact deeply constrained: it inherits the category (Chapter 4) and the motion (Chapter 5), and it sets the acquisition channels you can later afford (Chapter 8). BUS13-1 gave you cost-plus vs value-based vs competitive pricing as a strategic signal; BUS23-2 gave you the unit economics any price must satisfy. Here we choose a *structure* and *test* it before launch.

Definition 9 (Launch pricing structures)

- **Penetration** — deliberately low to win share fast; bets on retention and later expansion.
- **Skimming** — high initial price to harvest the most willing buyers, then descend.
- **Freemium** — a free tier for acquisition, paid tiers for value; only works when free users convert or refer enough to pay for themselves.
- **Tiered** — good/better/best packages that segment willingness-to-pay.
- **Value-based** — priced to a fraction of the quantified value delivered.
- **Usage-based** — price scales with consumption, so cost tracks value received.

Worked example — Lumen reasons from value, packages in tiers. Discovery quantified the value: a late file risks a €600 penalty and hours of partner time; a firm with 80 clients faces this several times a year. Value-based logic says Lumen can credibly charge a meaningful fraction of that avoided cost — far above the €10/month a “file-share” frame would allow. Lumen packages it as three **tiers** by number of active clients (a usage proxy that grows with the account), and rejects **freemium**: its buyers are firms, not individuals, and a free tier would attract hobbyists who never convert while teaching competitors the product. Penetration is tempting to win logos fast, but it would anchor the category low — exactly the frame Chapter 4 worked to escape.

6.1 Testing price sensitivity before launch

You do not have to guess. Two methods turn price into evidence.

Definition 10 (Van Westendorp Price Sensitivity Meter) Ask prospects four questions about a product: at what price is it *too cheap* (you’d doubt quality), *a bargain*, *getting expensive*, and *too expensive* (you’d never buy)? Plotting the cumulative curves, their intersections bound an *acceptable price range*: the floor where “too cheap” meets “too expensive,” and the ceiling beyond which resistance dominates. It reveals the *range*, not the single optimal point.

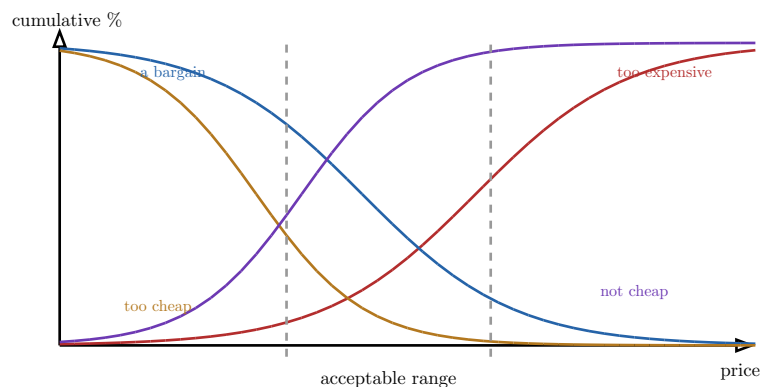


Figure 3: A Van Westendorp price-sensitivity meter. The four cumulative curves intersect to bound an acceptable price range (dashed band): below the floor, buyers doubt the quality; above the ceiling, “too expensive” dominates. The method gives a defensible *interval* to test within, not a single price.

Definition 11 (Willingness-to-pay interviews) Complementing the survey, WTP interviews probe the *reasoning* behind a price: what the prospect compares it to, what budget it comes from, what they already pay for the workaround, and where the “no” becomes a “yes.”

This recovers the value anchor and the reference category — exactly the value-based and package-vs-line-item conversation skills from BUS23-1 — and explains *why* a price works, which a survey number alone cannot.

Common pitfall Survey-stated prices skew high — saying “I’d pay €100” costs nothing. Van Westendorp gives a *range to test*, not a price to set; the real read comes from behaviour in actual sales conversations and from a confirmed alternative the buyer would otherwise choose. Always pair the survey with WTP interviews and live deals, and benchmark conversion by structure (freemium converts low single-digit percentages; sales-led trials far higher). Selecting a structure, testing it with Van Westendorp and WTP, and defending it against at least one plausible alternative is Learning Outcome 6.

7 Distribution strategy and channels

A price the market accepts still needs a *path* to reach the customer. Distribution is the choice of channels through which the product is sold and delivered — and, like the motion, it must fit the product type, the segment, and the competitive context rather than copying whatever a famous company did.

Definition 12 (Distribution fits product, segment, and context) The right channel depends on three things: the **product type** (a self-serve SaaS, a marketplace, a physical DTC good, and a professional service reach customers through structurally different paths); the **segment** (where the buyer already looks for solutions); and the **competitive context** (channels an incumbent has locked up are expensive to contest head-on).

Definition 13 (Build, acquire, or ally a channel) To reach customers through a channel you do not yet own, you can:

- **build** it yourself — your own sales team, website, or storefront: most control and margin, slowest and most capital-hungry;
- **acquire** access — buy placement, a list, or a distributor: fast reach, pays away margin, access can be withdrawn;
- **ally** — partner with someone who already has the relationship: fast and cheap, but the customer relationship is shared and the partner can leave.

Worked example — Lumen's channel question. Lumen’s buyers — small accounting firms — cluster around a few hubs: professional bodies, accounting-software ecosystems, and a handful of influential practice-management consultants. Lumen can **build** direct outreach (founder-led, full control, slow), **acquire** reach by advertising in the professional bodies’ channels (fast, costly, shallow trust), or **ally** with a popular bookkeeping platform whose users are exactly its segment (a partnership that borrows trust and distribution at once, but makes Lumen dependent on the partner). Early on, build founder outreach to learn; in parallel, pursue one alliance — the integration partner — as the scalable path, while avoiding a deep dependency before the economics are proven.

Common pitfall Channel choice is a competitive decision, not a logistics one. Entering through the channel an incumbent dominates means fighting on their ground and their terms; an under-contested channel where *your* segment already gathers is worth more than a bigger channel you cannot win. And every “ally” is a borrowed relationship — convenient until the partner changes terms or builds the feature themselves. Match the channel to product, segment, and context, and decide build/acquire/ally on control versus speed versus cost (Topic: distribution).

8 Acquisition and product–market fit: the metric shift

We now build the engine that brings customers in — and confront the most important phase transition in a venture’s life. **Before product–market fit (PMF)** and **after it**, acquisition is a different activity measured by different numbers. Confusing the two is how ventures scale a machine that is not yet ready.

8.1 Acquisition before PMF

Definition 14 (Pre-PMF acquisition levers) Before PMF the goal is *learning and the first customers*, not efficiency. The levers are deliberately **unscalable**:

- **direct outreach** — founders personally contacting prospects one by one;
- **community** — being present where the segment already gathers;
- **content** — writing the thing your buyer is searching for;
- **referral** — early lovers introducing the next customers.

These maximise *insight per customer* and rarely have payback-positive economics yet — which is correct at this stage.

Worked example — Lumen's first thirty customers are hand-won. Maya does not buy ads. She emails 200 firms personally, joins two accountant forums, publishes a short guide on “ending the deadline-week scramble,” and asks every happy firm for one introduction. It is slow and unscalable — and that is the point: each of the first thirty customers teaches her something that rewrites the positioning, the price, or the product. Spending on paid acquisition now would only buy expensive customers for a message not yet proven.

8.2 Measuring PMF and the metric shift

Definition 15 (Product–market fit and the metrics that bracket it) PMF is the state where the market pulls the product out of your hands: usage deepens, customers stay, and referral becomes organic. **Before PMF** you track *qualitative and engagement signals* — depth of use, retention curves that flatten rather than decay to zero, and the proportion of users who would be “very disappointed” without the product. **After PMF** the question becomes efficiency, so the metrics shift to *CAC, LTV, payback period, and growth rate* — the unit economics from BUS23-2 and BUS33-1, now used to govern scale.

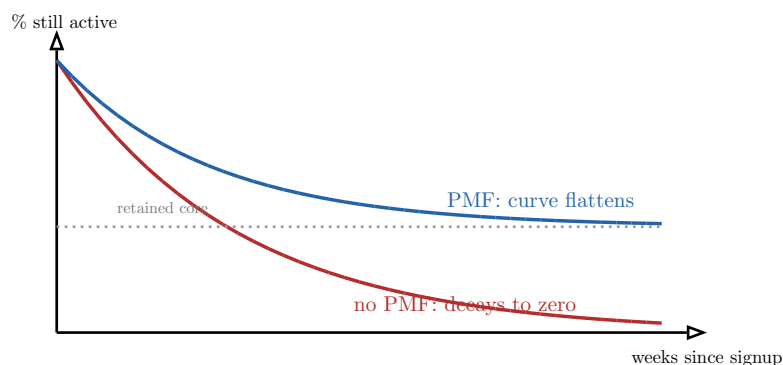


Figure 4: The PMF signal in a retention curve. Without PMF (red) every cohort decays toward zero — a leaky bucket no amount of acquisition can fill. With PMF (blue) retention flattens onto a plateau: a durable core that compounds. The flattening, not the signup count, is the green light to shift to CAC/LTV metrics and scale.

Common pitfall CAC and LTV computed before PMF are *fiction*. With retention still decaying to zero, LTV is unknowable and any “payback” is an artefact of a cohort you will lose. Reporting a flattering LTV/CAC pre-PMF is the number that most often justifies premature scaling (Chapter 9). Use engagement and retention to decide *whether* you have fit; switch to CAC/LTV/payback/growth only *after* the retention curve flattens. Identifying pre-PMF levers and why they differ structurally from post-PMF strategy is Learning Outcome 7; defining and tracking the two metric sets, and what shifts at PMF, is Learning Outcome 8.

9 Scaling discipline: avoiding premature scaling

This is the chapter the whole course is built around. The dominant cause of early-stage death is not a bad product or a weak market — it is **premature scaling**: spending on growth as if the commercial machine were proven when it is not. Discipline here is the difference between a venture that compounds and one that accelerates into a wall.

Definition 16 (Premature scaling) Premature scaling is investing in the inputs of scale — paid acquisition, sales hires, inventory, ops — *before* the evidence shows a repeatable, efficient path from stranger to retained customer. It converts runway into volume at a loss that does not improve with size, because the underlying motion or fit is not yet there.

Worked example — How Lumen could kill itself by succeeding. Suppose Lumen, thrilled by thirty hand-won customers, hires three salespeople, signs a €120k/year ad contract, and triples its plan. But retention has not yet flattened and the message still converts erratically. The reps sell a story that does not stick, churn eats the new logos, CAC balloons because the funnel is not tuned, and fifteen months of runway becomes five. The venture did not fail by doing too little; it failed by scaling an unproven machine — the textbook premature-scaling death.

Definition 17 (Observable signals of premature scaling) Scaling is premature when you are growing spend while these are still true:

- retention curves have **not** flattened (no PMF, Chapter 8);
- CAC is rising, not falling, as you spend more;

- conversion rates differ wildly by rep or channel (the motion is not repeatable);
- the founders no longer understand why deals are won or lost.

Any of these means the machine is not ready to take more fuel.

9.1 Test-and-learn after launch

The alternative to scaling on faith is scaling on evidence — the experimentation discipline of BUS33-1, applied to GTM rather than to ads.

Definition 18 (Test-and-learn on positioning, pricing, and channel) Post-launch iteration treats each GTM element as a hypothesis tested in sequence: state *what was tested*, *what was learned*, *what changed* as a result, and *which assumption is now falsified*. A “test” with no possible falsifying outcome is not a test — it is a roadmap with a rationale attached.

Worked example — One disciplined Lumen iteration. *Tested:* reposition from “document collection tool” to “client-file completion system” on the landing page, against the old version. *Learned:* demo-request rate rose from 4% to 9%; sales calls stopped arguing about Dropbox. *Changed:* the category language propagates to the deck and outreach. *Falsified:* the assumption that buyers frame Lumen as storage — they frame it as deadline risk. One assumption killed, one change shipped, evidence on record.

Common pitfall “Growth is up” is not evidence that scaling is working — growth can rise while the economics rot beneath it (the value-destroying-growth lesson, transposed from finance to GTM). Before adding fuel, confirm the machine is *repeatable* (conversion stable across reps and channels), *efficient* (CAC flat or falling, payback inside tolerance), and *retained* (curves flattened). Judging when evidence is sufficient to shift from exploratory growth to a scalable engine — and naming the signals that would invalidate the decision — is Learning Outcome 9; running disciplined test-and-learn is Learning Outcome 10.

10 Building the first commercial machine

With fit established and discipline in place, you build the machine that turns the motion into a repeatable operation: a founder-led sales protocol, the first commercial hires, a structured pipeline, and a plan to hand the selling from founders to a team. This is where BUS23-1’s sales craft becomes an *organisation*, not a personal skill.

10.1 Founder-led sales

Definition 19 (The founder-led sales protocol) Before any rep is hired, the founders sell — not because it scales but because it *teaches*. A protocol makes that learning systematic: a defined target list, a discovery script (SPIN, from BUS23-1), a consistent demo, an objection log, and a written record of why each deal was won or lost. The output is not just revenue; it is the **repeatable playbook** a rep will later run.

Common pitfall You cannot hire your way out of a sales motion you have not personally proven. A rep handed a vague “go sell it” will fail, and you will wrongly conclude the market

is the problem. The founder's job is to make the sale boringly repeatable *first*, then transfer a script that works — not to outsource discovery of the script to someone with no product knowledge and a quota.

10.2 The first three commercial hires, by motion

Definition 20 (First three hires depend on the motion) Who you hire first is dictated by Chapter 5's motion:

- **Sales-led** → an account executive who can run the founder's playbook, then a sales-development rep to feed pipeline, then a sales engineer or customer-success hire as deals get complex.
- **Product-led** → a growth engineer/marketer to optimise the self-serve funnel, then a product-onboarding owner, then a customer-success lead for expansion.
- **Marketing-led** → a content/SEO owner, then a demand-gen marketer, then an inside-sales rep to convert the inbound.

Hiring an enterprise AE into a product-led motion (or a growth engineer into an enterprise one) is the motion-mismatch error of Chapter 5 reappearing as a payroll mistake.

10.3 Structuring the pipeline

Definition 21 (Pipeline by stage, conversion, and cycle length) A pipeline makes the motion measurable: explicit **stages** (e.g. lead → qualified → demo → proposal → closed), a **conversion rate** between each, and a **cycle length** (time from first touch to close). Together they predict how much top-of-funnel produces how much revenue, and where the funnel leaks — the diagnosis skill from BUS23-1, now used to size and forecast the machine.

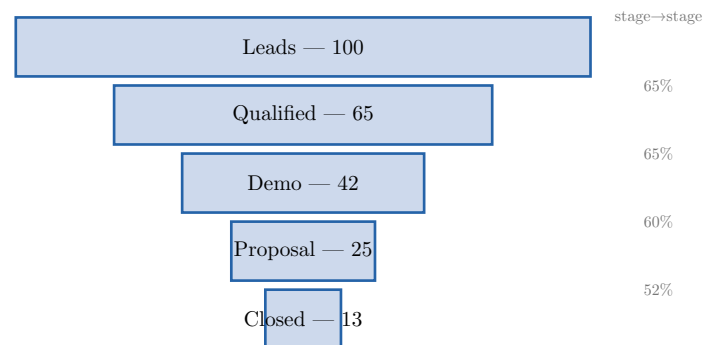


Figure 5: A structured sales pipeline. Each stage carries a count and a stage-to-stage conversion rate; the product of the rates and the cycle length turn raw leads into a revenue forecast and expose the leakiest stage — here the proposal→closed step at 52%, the first place to intervene.

10.4 From founder-led to team-led

Common pitfall The transition from founder-led to team-led sales fails when founders hand over *before* the playbook is written or *stay on* long after reps could run it. The signal to transfer is a documented, repeatable motion with stable conversion that a trained rep can reproduce — not a revenue milestone or founder exhaustion. Authoring the founder-led

protocol, specifying the first three hires by motion, structuring the pipeline with conversion and cycle length, and planning the transition is Learning Outcome 11.

11 Growth model design

A pipeline forecasts sales; a **growth model** explains how the whole business compounds. It connects product usage, acquisition, retention, and revenue into one coherent engine that a stakeholder — an investor, a co-founder, the jury — can stress-test. This is where the AARRR and growth-loop thinking of BUS33-1 is assembled into a single causal picture of the venture.

Definition 22 (The growth model as a connected engine) A growth model links four flows: **acquisition** (how new users arrive), **activation/usage** (how they reach value), **retention** (how many stay, the plateau of Chapter 8), and **revenue** (how retained usage converts to money). It is *coherent* when the outputs of one flow are the inputs of the next — and a **loop** when revenue or usage feeds back into acquisition (referrals, content from usage, expansion), so growth compounds rather than relying on ever-more spend.

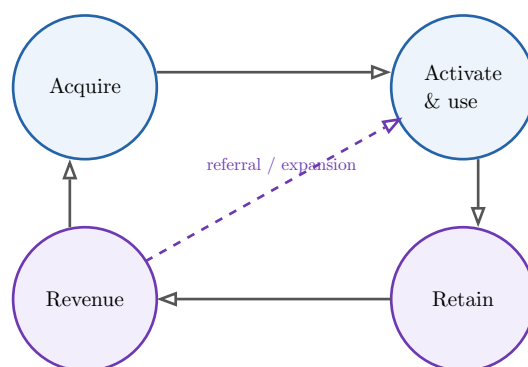


Figure 6: A growth model as a closed loop. Acquisition feeds activation, which feeds retention, which feeds revenue; the dashed feedback edge — referrals and expansion revenue financing the next cohort's acquisition — is what makes the engine *compound* rather than depend on ever-rising spend. The constraint that breaks first (here, retention) caps the whole loop.

Common pitfall A growth model that is just a revenue spreadsheet with an optimistic growth rate is not a model — it is a wish. The test is whether each flow's number is grounded in the evidence built across this course (retention from Chapter 8, conversion from Chapter 10, CAC from the channel economics of BUS33-1) and whether the *loop* is real. Identify the constraint that breaks the loop first at scale, exactly as you learned to in BUS33-1, and you have a model a jury can stress-test rather than admire. Designing this coherent engine is Topic 10's deliverable.

12 The integrated go-to-market plan

Everything converges here. The course terminates in a complete GTM plan for a real business model — B2B SaaS, marketplace, DTC physical product, or professional service — defended before a practitioner jury whose job is to test *coherence across elements* and penalise incoherence. The lesson of the capstone is the lesson of the whole book: the elements are not independent, and a plan of individually fine but mutually contradictory parts is the most common way to fail.

Definition 23 (The integrated GTM plan) A complete plan for a named venture integrating six components:

1. **positioning** — Dunford’s five components with a defended category (Chapter 4);
2. **motion** — selected against the four criteria (Chapter 5);
3. **pricing** — a structure with testable sensitivity (Chapter 6);
4. **a 100-customer acquisition plan** — pre- vs post-PMF levers and channels (Chapters 7–8);
5. **a 12-month commercial structure** — founder-led protocol, first hires, pipeline, transition (Chapter 10);
6. **a growth model** — usage → acquisition → retention → revenue, that stakeholders can stress-test (Chapter 11).

12.1 Coherence is the deliverable

Worked example — Lumen, assembled. The pieces only form a plan if one thread runs through all of them. Lumen’s thread is its *category choice*: “client-file completion system for small accounting firms.” That single commitment drives every component. The **positioning** makes deadline risk, not storage, the frame. The **motion** is founder-led sales with a later self-serve layer, because the buyer is a firm at a mid price, not an individual at a low one. The **price** is value-based and tiered against avoided penalties — a price the “storage” category could never support. The **acquisition plan** hand-wins the first thirty through outreach and the professional-body community, then, only after retention flattens, opens the integration-partner alliance toward 100. The **commercial structure** hires an AE to run Maya’s proven playbook, not a growth engineer for a funnel that does not yet exist. The **growth model** shows referrals and expansion compounding once the retained core is real. Change the category and *every* element moves with it.

Common pitfall The plan fails not when one number is wrong but when the components **contradict** each other: a low penetration price that cannot fund the sales-led motion it is paired with; a product-led acquisition plan bolted onto an enterprise hire plan; a growth model that assumes referral loops the retention data does not support; a 100-customer plan that scales paid spend before PMF. A jury — like a real market — probes exactly these seams. The integration *is* the test (Learning Outcome 12): defend a plan whose positioning, motion, pricing, acquisition, commercial structure, and growth model are mutually consistent, and you have built not a deck but a machine that can survive its first eighteen months.

The arc of this course. Take the optimisation tools the Build stage gave you — marketing strategy (BUS13-1), sales and pipeline craft (BUS23-1), business models and unit economics (BUS23-2), growth analytics (BUS33-1) — and use them to *build a commercial machine from nothing*, not tune one that exists. **Demand-side discipline:** adopt the product mindset, run real customer discovery, and use JTBD and Kano to separate the features that create demand from the hygiene that merely avoids losing it. **Strategic commitments:** write positioning whose category is the most consequential pre-launch choice, select a motion against price/complexity/team/stage to dodge the eighteen-month mismatch, and price by value — tested with Van Westendorp and WTP — not by category default. **The commercial machine:** choose channels by build/acquire/ally, acquire unscalably before PMF and efficiently after it,

govern scaling by the observable signals of premature scaling, and build founder-led sales into a pipeline, hires, and a hand-off. **Integration:** assemble a growth loop that compounds and a single coherent GTM plan — because in a real venture, these were never separate decisions, and the seams between them are where it lives or dies.