

BUS63-10

Financial Due Diligence

Albert School — 22 hours

Contents

Preface	2
1 Quality-of-earnings analysis	3
1.1 Non-recurring items	3
1.2 Aggressive revenue recognition	4
1.3 Pro-forma adjustments: legitimate versus cosmetic	4
2 The buyer-normalised EBITDA bridge	5
2.1 Value-transfer implications	6
3 Transaction-context net debt	6
3.1 Accounting versus contractual definition	7
3.2 Non-obvious debt-like items	7
4 Normalised working capital	7
4.1 Why a single month-end will not do: seasonality	8
4.2 Setting the peg	8
5 Working capital and net debt in the Share Purchase Agreement	8
5.1 The boundary is the battleground	9
5.2 Locked-box versus completion accounts	9
5.3 The financial stakes of a poorly drafted BFR clause	9
6 Tax due diligence	10
6.1 Accounting tax versus cash tax	10
6.2 Latent tax risks	10
6.3 Deferred tax liabilities as an earnings-quality signal	10
7 LBO debt architecture	11
7.1 Financial covenants	11
7.2 Headroom and the binding covenant	11
7.3 Structural adjustments	12
8 Forensic due diligence	12
8.1 Recurring manipulation patterns	12
8.2 Aggressive accounting versus fraud	12
8.3 The structural limits of due diligence	13
9 Due diligence as a management process	13
10 The red-flags memo under time pressure	13
10.1 Triage: spending the hours where they pay	13
10.2 Quantified valuation impacts	14
10.3 Escalate versus close, and the oral defence	14

Preface

By the time you reach this course you can already do almost everything an analyst does to a set of accounts — you have simply been doing it to understand a company, not to defend a price across a negotiating table. From **BUS13-3** you read the three statements as constructed documents and you recognise earnings-quality signals — non-recurring items, aggressive revenue recognition, normalised earnings. From **BUS23-3** you compute ROIC, free cash flow and the Working Capital Requirement (BFR), and you adjust financials for honest comparison. From **BUS33-3** you handle the cost of debt, the tax shield and the WACC. From **BUS43-3** you price a firm by DCF and by multiples, and you already use enterprise value, net debt and normalised EBITDA as the raw inputs of a valuation.

Valuation answered one question: *what is this company worth, given a set of numbers?* Financial Due Diligence (FDD) answers the question that comes logically before it and is asked in anger only when real money is about to change hands: *can those numbers be trusted, and what is the number a buyer should actually underwrite?* The seller hands you an EBITDA, a net-debt figure and a working-capital balance. Each is a claim. FDD is the discipline of testing those claims under a fixed budget of hours and incomplete information, and converting what you find into euros of price.

That last clause is the whole course. FDD is not an audit and not a valuation — it is **judgment under constraint**. You will never have time to check everything, you will never be given every document, and the people answering your questions are usually the people who built the number you are doubting. So the real skills are triage (where to spend the hours), translation (turning a finding into a price impact), and decision (when to escalate, and when to recommend not proceeding at all).

One fictional target runs through the whole book: **Atlas Packaging**, a mid-market manufacturer of food-grade packaging being sold to a private-equity buyer in a leveraged transaction. The seller presents €50m of EBITDA and a clean story. By the end you will have walked it down to the number a buyer should underwrite, reconciled its real net debt, set its working-capital peg, reviewed its tax exposure and its acquisition-debt covenants, and written the red-flags memo that decides the deal.

1 Quality-of-earnings analysis

Start with the suspicion, not the definition. Atlas reports €50m of EBITDA, up smartly from €44m two years ago. Before you believe the trajectory, ask a single question of every euro in that number: *will the buyer still earn it next year, in the ordinary course of business, after the seller has gone?* Quality-of-earnings (QoE) analysis is the systematic asking of that question, line by line, until what remains is the recurring, cash-generative profitability the buyer will actually own.

Definition 1 (Earnings quality) The degree to which reported earnings represent *recurring*, cash-generative profitability that will persist under the buyer's ownership — as opposed to one-off, mis-classified, or accounting-driven amounts that happen to land in the period. High-quality earnings repeat; low-quality earnings flatter a single year.

1.1 Non-recurring items

Worked example — Atlas's flattering year. Digging into Atlas's €50m, you find it includes a €2.0m insurance recovery from a one-off warehouse fire and a €1.5m gain on the sale of a surplus building. Neither will happen again. Strip them out and the recurring base is already €3.5m lower than advertised — before you have looked at anything subtle.

Definition 2 (Non-recurring item) Income or expense not expected to repeat in the ordinary course of business: restructuring charges, litigation settlements, insurance recoveries, gains or losses on asset disposals, one-off consulting or transaction fees. QoE removes them — in *both* directions — to isolate the recurring base.

Common pitfall The trap is direction, not existence. An item can be one-off in *direction* yet recurring in *kind*. A company that books a “non-recurring” restructuring charge every single year has recurring restructuring; treating each year’s instance as a one-off add-back quietly inflates underlying EBITDA. Before you accept an add-back, look back three to five years: if the same “exceptional” line keeps reappearing, it is a cost of doing business.

1.2 Aggressive revenue recognition

Removing one-off items cleans the surface. Aggressive revenue recognition is deeper: it borrows profit from the future, so the reported figure is real on paper but will not recur — or will actively reverse.

Worked example — The quarter that was too good. Atlas’s revenue jumps 9% in the final quarter before the sale. The QoE team pulls monthly shipment and distributor inventory data and finds distributors holding three months of stock against a normal one month. Atlas has been shipping product its customers cannot yet sell — pulling next year’s sales into this year to dress the company for sale. None of that quarter’s bulge is sustainable revenue.

The recurring patterns to test for:

- **Channel stuffing** — shipping more product to distributors than they can sell through, pulling future sales into the current period (Atlas, above).
- **Long-term contracts** — recognising revenue too early on percentage-of-completion contracts by overstating the degree of completion.
- **Vendor financing** — lending the customer the money to buy the product, so the “sale” depends on a collection that may never occur.

Each produces revenue that either will not recur or will reverse. QoE tests whether recognition matches *delivery* and *collectability*, not merely whether an invoice was raised.

1.3 Pro-forma adjustments: legitimate versus cosmetic

The seller will not hand you raw earnings; they will hand you an *adjusted* EBITDA, already grossed up by a list of pro-forma add-backs. Your job is to sort the honest from the cosmetic.

Definition 3 (Pro-forma adjustment) A restatement of historical EBITDA *as if* a change had been in place for the whole period. It is **legitimate** when it reflects a real, evidenced and committed change in the business; it is **cosmetic** when it adds back costs that will in fact continue, or projects savings not yet secured.

Legitimate	Cosmetic
Full-year effect of a signed, closed acquisition	Run-rate of “synergies” not yet realised
Removal of a genuinely closed business line	Add-back of owner compensation that <i>continues</i> post-deal
Annualising a price increase already in force	Add-back of “exceptional” costs that recur every year
Standalone cost the seller bore but the buyer will not	Cost savings “management is confident” it will find

Common pitfall The test is evidence and commitment, not plausibility. “We are confident we can cut €3m of overhead” is a *forecast*, and forecasts belong to the buyer’s business plan, not the historical EBITDA the buyer pays a multiple for. A legitimate add-back can be documented to a signed contract, a board decision, or a completed action; a cosmetic one rests on management’s confidence.

2 The buyer-normalised EBITDA bridge

The seller presents one EBITDA; the buyer must underwrite another. The **EBITDA bridge** is the itemised, signed reconciliation from the first to the second — the single most negotiated artefact of an FDD.

Worked example — Walking Atlas down to the underwritten number. Assemble every QoE finding into one bridge from the seller’s €50m:

Adjustment	€m
Seller-presented EBITDA	50.0
+ Owner’s above-market salary (owner departs)	+ 1.5
– One-off insurance recovery and disposal gain	– 3.5
– Cosmetic synergy add-back, not yet realised	– 2.0
– Channel-stuffing pull-forward, reversed	– 2.0
– Standalone costs the buyer will bear	– 2.0
Buyer-normalised EBITDA	42.0

€50m of story becomes €42m a buyer should pay for. Note the bridge runs both ways: the owner’s excess salary is a *legitimate add-back* (a real cost the buyer will not incur), while the “synergy” add-back is reversed because it is not yet real.

Definition 4 (Buyer-normalised EBITDA bridge) The signed list of adjustments reconciling seller-presented EBITDA to the number the buyer will underwrite:

$$\text{Buyer-normalised EBITDA} = \text{Seller EBITDA} + \sum_i a_i,$$

where each a_i is a documented add-back (positive) or deduction (negative): removal of non-recurring items, reversal of cosmetic add-backs, normalisation of accounting policies, correction of mis-classifications, and recognition of costs the buyer will carry but the seller did not.




Figure 1: The Atlas EBITDA bridge as a waterfall. The first (blue) and last (purple) bars are the seller and buyer-normalised levels; the floating bars are the adjustments — one upward (a real cost the buyer avoids) and four downward. The y-axis is zoomed to the 38–52 band to make the €8m walk visible.

2.1 Value-transfer implications

Why fight line by line over a few million of EBITDA? Because each line is multiplied.

Worked example — Every euro is eight. Atlas is priced at an $8\times$ EBITDA multiple. The seller’s €50m implies an enterprise value of $8 \times 50 = \text{€}400\text{m}$. The buyer’s €42m implies $8 \times 42 = \text{€}336\text{m}$. The €8m walked off the EBITDA is worth $8 \times 8 = \text{€}64\text{m}$ of price — sixteen times the headline EBITDA gap. The bridge is not bookkeeping; it is a €64m argument.

At a multiple m , every unit of *recurring* EBITDA disputed in the bridge moves enterprise value by m units. The bridge is therefore the principal quantified input to price negotiation: each line is a value-transfer claim between buyer and seller, and the analyst who can evidence a line owns its value.

Common pitfall Multiply only *recurring EBITDA* adjustments by the multiple. A one-off cash item — a litigation payout, a deferred bonus, a missed capex catch-up — belongs in *net debt* or *working capital* at its face value, *not* in the EBITDA bridge grossed up by m . Putting a one-off cash cost in the bridge and multiplying it by 8 over-penalises the seller eightfold and is the fastest way to lose credibility across the table.

3 Transaction-context net debt

Valuation taught you a tidy identity: equity value = enterprise value – net debt. In a transaction, “net debt” stops being an accounting figure you read off the balance sheet and becomes a **contractual** one you negotiate — and it is almost always larger than the balance sheet suggests.

Worked example — Atlas's net debt is bigger than it looks. Atlas’s balance sheet shows €120m of borrowings and €20m of cash, so accounting net debt is €100m. But the FDD finds:

a €8m earn-out still owed on a past acquisition; a €12m deficit on the defined-benefit pension; €5m of committed environmental remediation at an old site; €4m of probable litigation; and €6m of the cash trapped in a subsidiary that cannot legally distribute it. Each of these reduces the equity price exactly as debt would. Transaction net debt is $100 + 8 + 12 + 5 + 4 + 6 = €135\text{m}$ — €35m more than the accounts show, straight off the seller's equity cheque.

Definition 5 (Transaction-context net debt) The set of items deducted from enterprise value to reach the equity price, as defined by the sale agreement: reported financial debt, less *freely available* cash, plus all **debt-like** items and less any **cash-like** items the parties agree to treat as such. The accounting definition is the starting point, not the answer.

3.1 Accounting versus contractual definition

The *accounting* definition counts interest-bearing borrowings less cash and equivalents. The *contractual* definition adds everything that is *economically* debt — a fixed or near-fixed claim on the business that exists regardless of future trading — even when it carries no “loan” label.

3.2 Non-obvious debt-like items

- **Earn-outs** — deferred contingent consideration owed for a past acquisition. A fixed claim the buyer inherits.
- **Environmental liabilities** — committed obligations to remediate contamination.
- **Underfunded pensions** — the deficit between defined-benefit obligations and plan assets; a debt to the workforce.
- **Contingent liabilities** — probable obligations from litigation, guarantees, or warranty claims.

Symmetrically, not all cash is a deduction from debt. **Restricted or trapped cash** — minimum operating cash the business needs to function, or cash locked in a jurisdiction by repatriation barriers — is not freely available to the buyer and should not net against debt.

Common pitfall Cash on the balance sheet is *not* automatically “free” cash. The cash a business must hold just to run (float in tills, minimum compensating balances) and cash trapped abroad are unavailable; netting them against debt overstates what the seller delivers and transfers value to the seller. Test every euro of cash for availability, not just existence.

Remark The boundary matters as much as the total. A single accrual — say a deferred bonus pool — can be argued into net debt (debt-like, reduces price) or into working capital (part of the normal operating cycle). The same euro cannot live in both, and where it lands changes the price. We return to this boundary war in Chapter 5.

4 Normalised working capital

The buyer must receive Atlas with a *normal* tankful of operating working capital — enough that the business can run on day one without an emergency cash injection. Neither party should be able to transfer value by quietly draining or stuffing the operating cycle in the weeks around closing. Setting the *normal* level — the **peg** — is the job here. The working capital in question is the operating BFR you built in BUS23-3: receivables plus inventory less payables.

4.1 Why a single month-end will not do: seasonality

Worked example — Atlas swings with the seasons. Atlas supplies packaging to food producers, so it builds inventory ahead of the autumn harvest and runs it down over winter. Its month-end working capital ranges from about €22m in January to €38m in September, averaging €30m across the year. Close the deal in September against the €38m balance and the buyer pays for a seasonal peak that will melt to €22m by winter — handing the seller €8m of pure timing. Close in January and the reverse robbery occurs. A single month-end is the wrong anchor; the *normal* level is the one that strips the season out.




Figure 2: Atlas's monthly working capital over a year. The seasonal swing (blue) runs from a January trough of €22m to a September peak of €38m; the peg (dashed) is the €30m twelve-month average that removes the season. Pricing off any single month-end transfers the gap between that month and the peg.

4.2 Setting the peg

Definition 6 (Working-capital peg) The target (*normal*) level of working capital agreed in the contract, typically the average of normalised monthly working capital over a representative trailing period (commonly twelve months). At completion, actual working capital is compared to the peg: the price is adjusted *up* if actual exceeds the peg (the seller leaves more than normal) and *down* if it falls short.

The peg converts working capital from a timing game into a fixed reference. Set it too low and the seller can deliver the business starved of operating cash and still collect full price; set it too high and the seller is penalised for a level the business never truly needed.

Common pitfall “Normalise” before you average. The trailing months may themselves be distorted — a receivables factoring spike, a one-off inventory build, a stretched-payables window engineered before sale. Averaging raw month-ends bakes the manipulation into the peg. Strip out the abnormal months first, *then* take the average.

5 Working capital and net debt in the Share Purchase Agreement

Everything in the last two chapters meets the lawyers here. The **Share Purchase Agreement** (SPA) is the contract of sale, and its working-capital and net-debt clauses draw the single most

consequential line in the document: the boundary that decides, for every balance-sheet item, whether it counts as net debt, as working capital, or as neither.

5.1 The boundary is the battleground

Worked example — Where does the deferred bonus go?. Atlas owes €3m of accrued staff bonuses at completion. Put it in *net debt* and it reduces the equity price by €3m, full stop. Put it in *working capital* and it merely lowers working capital toward the peg — and if working capital is still above the peg, it may cost the seller nothing at all. The same €3m, two clauses, a €3m swing. Multiply this fight across a dozen accruals and the SPA scope definition is worth more than most of the QoE bridge.

Definition 7 (Scope definition) The clause-by-clause classification of every balance-sheet item into exactly one of: *net debt*, *working capital*, or *neither*. The cardinal rule is **no item in two buckets, and no item in none** — every euro is counted exactly once.

5.2 Locked-box versus completion accounts

Two mechanisms fix *when* the price is settled and who bears the operating result in between.

Locked-box	Completion accounts
Price fixed on a past, “locked” balance-sheet date	Price adjusted to actual figures measured at completion
No post-closing true-up; economic value accrues to the buyer from the locked date	Post-closing true-up of net debt and working capital against the peg
Certainty; lower closing cost; leakage controlled contractually	Accuracy; but post-closing dispute risk over the completion accounts




Figure 3: The two settlement mechanisms on a timeline. Under a locked-box the price is struck on a past locked date and economic value passes to the buyer from then (so the seller is barred from “leakage”); under completion accounts the price is trueed up against the peg shortly after completion.

5.3 The financial stakes of a poorly drafted BFR clause

Common pitfall A vague working-capital (BFR) clause is not a legal nicety — it is a direct leak of price. Three failure modes recur:

- **Double count or gap:** an item lands in both net debt and working capital (the seller is hit twice) or in neither (a euro of liability escapes the price entirely).
- **A mis-set peg:** a peg that ignores true seasonality lets whichever party controls the closing date time it to their advantage.

- **Ambiguous scope:** undefined terms produce post-closing disputes over the completion-accounts true-up, sometimes litigated for years.

For the *buyer*, a loose clause risks acquiring a business starved of operating cash; for the *seller*, it risks an unjustified downward adjustment. The clause is a determinant of net price, not boilerplate.

6 Tax due diligence

Tax due diligence does two jobs at once: it hunts for *hidden liabilities* (tax the buyer would inherit) and it reads the tax position as a *signal* of historical earnings quality. A company's tax filings are a second, independently prepared set of books — and divergences between them and the accounts are informative.

6.1 Accounting tax versus cash tax

Worked example — Atlas pays less tax than it reports. Atlas's income statement shows a €10m tax charge, yet its cash-flow statement shows only €6m of tax actually paid. The €4m gap is deferred — timing differences between when profit is booked for shareholders and when it is taxed. That gap is the thread to pull: it tells you the accounting profit is running ahead of the taxable profit, and you need to know why.

Definition 8 (Accounting tax versus cash tax) The *accounting tax* is the charge in the income statement; the *cash tax* is the amount actually paid. They reconcile through deferred tax and one-off items:

$$\text{Cash tax} = \text{Accounting tax} - \Delta \text{ deferred tax} \pm \text{timing and one-off items} .$$

6.2 Latent tax risks

- **Transfer pricing** — intra-group prices that tax authorities may challenge, generating reassessment exposure (and back-tax plus penalties).
- **Ongoing reassessments** — open audits or live disputes with tax authorities, whose downside must be estimated and treated as debt-like.
- **Derogatory regimes** — reliance on preferential tax treatments (special zones, patent boxes, rulings) that may lapse, be withdrawn, or be disallowed, raising the go-forward effective rate.

Common pitfall A low effective tax rate is not a free gift to the buyer's model. If it rests on a derogatory regime that expires or an aggressive transfer-pricing position that may be reassessed, the *sustainable* rate is higher — which lowers go-forward cash flow — *and* the historical under-payment may crystallise as a back-tax liability that belongs in net debt. One soft tax line can cost twice.

6.3 Deferred tax liabilities as an earnings-quality signal

Definition 9 (Deferred tax liability (DTL)) A liability for tax that will be paid in future periods because taxable income has been deferred relative to accounting income.

A *persistently growing* DTL means accounting income has consistently exceeded taxable income — the company is reporting earnings to its shareholders faster than it is recognising them to the tax authority. Sustained divergence is a classic low-earnings-quality signal: the profit shown to the market is “softer” than the profit shown to the state. So the DTL trend is read twice — as a future cash obligation (it enters net-debt thinking) and as a flag on the quality of the historical earnings the bridge is built on.

7 LBO debt architecture

Atlas is being bought in a **leveraged buy-out** (LBO): the buyer funds most of the price with acquisition debt that the target’s own cash flows must service. That changes the FDD question from “are the earnings real?” to the sharper “are the earnings *robust enough* to service this debt without breaching its terms?” The lenders protect themselves with covenants; you must judge how much room those covenants leave.

7.1 Financial covenants

Definition 10 (Financial covenant) A contractual financial test the borrower must satisfy at each measurement date; a breach is a default trigger that can accelerate the loan or hand control to lenders. The three workhorse covenants:

- *Leverage*: net debt / EBITDA must stay *below* a ceiling.
- *Interest coverage*: EBITDA / interest must stay *above* a floor.
- *Cash-flow* (debt-service coverage): cash available for debt service relative to scheduled service, *above* a floor.

7.2 Headroom and the binding covenant

Definition 11 (Headroom) The margin between a covenant threshold and the borrower’s actual ratio — how far performance can deteriorate before a breach. For a leverage ceiling L and current leverage ℓ , headroom is $L - \ell$, measured in turns of EBITDA.

The **binding** covenant is the one with the least headroom: it will be breached first, so it is the real constraint on the business, regardless of which test management talks about most.

Worked example — The covenant that bites depends on whose EBITDA. Atlas is bought with €210m of acquisition net debt against a leverage ceiling of $L = 5.5 \times$. On the *seller’s* €50m EBITDA, leverage is $210/50 = 4.2 \times$ — a comfortable 1.3 turns of headroom. But the buyer does not own €50m; the bridge said €42m. On the *buyer-normalised* €42m, leverage is $210/42 = 5.0 \times$ — headroom collapses to 0.5 turns. The deal that looked safe on the seller’s number is one bad quarter from a breach on the real one.




Figure 4: Covenant headroom on Atlas’s €210m acquisition debt. Measured on the seller’s €50m EBITDA the leverage bar (blue) stops well short of the 5.5× ceiling; recomputed on the buyer-normalised €42m (purple) it nearly touches it. Headroom falls from 1.3 to 0.5 turns purely by using the right EBITDA.

7.3 Structural adjustments

Where headroom is thin, the package can be re-engineered: loosen covenant levels, re-profile (back-load) amortisation, add an *equity cure* right letting the sponsor inject equity to fix a breach, or simply reduce the leverage by putting in more equity. FDD’s job is to identify the binding constraint and recommend the adjustment that relieves *it* specifically.

Common pitfall Headroom computed on seller-presented EBITDA overstates safety, every time. Recompute every covenant ratio on *buyer-normalised* EBITDA before you judge the package — as Atlas shows, a covenant that looks comfortable on the seller’s number can be effectively in breach on the number the buyer will actually report.

8 Forensic due diligence

QoE assumes the numbers are honestly prepared but optimistically presented. Forensic DD drops that assumption and looks for *deliberate manipulation*. The skill is to recognise the recurring shapes of manipulation, to draw the line between aggressive accounting and fraud, and — most importantly — to stay humble about what DD can and cannot catch.

8.1 Recurring manipulation patterns

Manipulation is not infinitely creative; it recurs in a small set of forms: fictitious or related-party revenue; overstated assets (inventory or receivables that do not exist); capitalising costs that should be expensed (to lift both profit and assets); and round-tripping cash between entities to fabricate sales.

Worked example — The Sino-Forest shape. In the canonical Sino-Forest case, the core assets — vast timber holdings — could not be independently verified to exist at the claimed scale, and revenue flowed through related intermediaries that obscured whether real third-party sales had occurred. The lesson generalises: when the *existence* of an asset or the *independence* of a counterparty cannot be confirmed by a source outside the company, treat the associated earnings as unproven until they can.

8.2 Aggressive accounting versus fraud

Definition 12 (Aggressive accounting versus fraud) *Aggressive accounting* exploits the discretion permitted *within* the accounting framework — optimistic but defensible estimates, early-but-arguable recognition. *Fraud* is misrepresentation *outside* the framework

— fabricated transactions, falsified records, concealed related parties. The line is intent and rule-breaking: staying inside the rules while pushing them, versus stepping outside them.

8.3 The structural limits of due diligence

Common pitfall DD *samples* it does not audit exhaustively. It relies on management-supplied information, it operates inside a fixed time-box, and it cannot — by construction — detect a *collusive* fraud designed specifically to defeat it. A clean DD is evidence that *no red flags were found within the work performed*, not proof that none exist. State this limit explicitly in the memo; promising more than DD can deliver is itself a professional failure.

9 Due diligence as a management process

An FDD is not an open-ended investigation that ends when curiosity is satisfied; it is a **managed engagement** that ends when the hours run out. Running it well is a discipline in its own right.

- **Team structure.** Work is split into workstreams — QoE, net debt, working capital, tax, legal/forensic — each with an owner, and a lead who integrates findings and owns the overall view. Findings cross-check: a tax gap may corroborate a QoE doubt.
- **Time-boxing.** Each workstream gets a fixed budget of hours. Investigation stops when the box is spent unless a finding earns an escalation. The constraint is the point — it forces prioritisation.
- **Escalation discipline.** A finding is escalated when its potential value impact crosses a pre-agreed threshold; otherwise it is closed and *recorded* (a closed finding is still documented, so the decision can be defended later).
- **Questioning management.** Questions are sequenced from open (“walk us through the Q4 revenue”) to specific (“explain these three September shipments to Distributor X”). Crucially, a **non-response** — an evasion, a delay, a data request that is never fulfilled — is itself evidence, and is weighted as such.

Remark Interpreting non-responses is a genuine analytical skill, not paranoia. Honest management can produce the supporting schedule; the systematic inability or unwillingness to produce it, especially around the one number you have flagged, is information. The non-answer often locates the problem more precisely than any answer would.

10 The red-flags memo under time pressure

Everything converges on one deliverable: a **red-flags memo**, produced under time constraint and defended orally before a jury that includes an invited M&A practitioner. The memo is where findings become a recommendation a buyer can act on.

10.1 Triage: spending the hours where they pay

Worked example — Atlas's hour budget. You have 200 investigation-hours left and four open issues: the channel-stuffing question (high likelihood, €16m price impact), the pension deficit (certain, €12m, but already well evidenced), a minor expense-classification quibble (low impact), and the transfer-pricing exposure (uncertain, potentially €20m). You do *not* split the hours evenly. The pension is quantified — close it. The expense quibble is immaterial — close

it. The hours go to the two issues that are both *large* and *unresolved*: channel-stuffing and transfer pricing.

Definition 13 (Risk weight) The product of an issue’s likelihood and its potential value impact, $\text{risk weight} = \text{likelihood} \times \text{impact}$, used to rank issues so investigation hours flow to the highest-weighted first. Issues that are already resolved, or immaterial, are closed regardless of category.




Figure 5: Triage by risk weight ($\text{likelihood} \times \text{impact}$). Hours flow to the shaded top-right: large *and* uncertain. The pension is high-impact but already resolved (grey), so it is closed and recorded; the expense quibble is low on both axes and closed. Channel stuffing and transfer pricing earn the remaining hours.

10.2 Quantified valuation impacts

Every red flag must be stated as a *number*, not a worry, and routed to the right place in the price:

- a recurring **EBITDA** adjustment is multiplied by the deal multiple;
- a one-off or fixed **debt-like** item enters net debt at face value;
- a **working-capital** shortfall against the peg adjusts the price directly.

Common pitfall “We have concerns about Q4 revenue” is not a finding; it is an anxiety. “Q4 contains €2m of channel-stuffed revenue; at 8× that is €16m of price” is a finding. The memo’s job is to convert every concern into a euro figure and a destination (EBITDA bridge, net debt, or working capital). A concern the analyst cannot quantify is a concern the buyer cannot negotiate.

10.3 Escalate versus close, and the oral defence

For each line of investigation the analyst decides to **escalate** — invest further hours, raise it to the deal team, or make it a condition of the deal — or to **close** it, recording the finding and moving on. The decision is justified by the issue’s risk weight and the hours remaining in the box. At the limit, the recommendation is not a price adjustment at all but *do not proceed*: when the red flags are too large, too numerous, or too unverifiable, the professional answer is to walk away.

The memo is then defended orally before the jury, including the practitioner juror, who tests the robustness of each quantified claim and each escalate-versus-close call. The defence stands or falls on the *evidence* behind each number and on the *judgment* of where the hours were spent — the two things this entire course has been training.

The Atlas FDD, end to end. Test earnings quality and build the buyer-normalised EBITDA bridge (€50m $\rightarrow$ €42m); reconcile transaction-context net debt beyond the accounting figure (€100m $\rightarrow$ €135m) and set the normalised working-capital peg (€30m), then fix their boundary precisely in the SPA; read the tax position as both a hidden liability and an earnings-quality signal; recompute the LBO covenants on the underwritten EBITDA, where headroom collapses from 1.3 to 0.5 turns; apply forensic scrutiny while stating honestly what DD cannot catch; and run the whole thing as a time-boxed, escalation-disciplined process that ends in a quantified, defensible red-flags memo — and, when the numbers warrant it, the courage to recommend not proceeding.