

HUM14-3

Philosophy I — The Examined Life

Authentic desire, freedom and responsibility, ethics under pressure

Thinkers as diagnostic instruments to wield, not authorities to memorise.

Albert School — 12 hours

Preface

There is a line, attributed to Socrates at his trial, that gives this module its name: *the unexamined life is not worth living*. It is a strong claim — most lives are unexamined, and most of them seem worth living to those who live them. But strip it of the rhetoric and a usable idea remains: a life run on borrowed goals, unquestioned definitions of success, and desires one has never inspected is a life lived partly by accident. Philosophy, in this course, is the tool for the inspection.

That is the whole orientation of what follows, and it is worth being explicit about it because it differs from what “philosophy” often means. We are not here to collect opinions of famous people so they can be quoted at dinner. We are here to build a small set of **instruments** — frameworks sharp enough that, pointed at a real decision you face, they show you something you could not see before. When we introduce Aristotle, Girard, Spinoza, Epictetus, Sartre, Kant or Mill, the operative question is never “what did this person believe?” but “**what does this framework help me see that the others miss?**” A thinker who shows you nothing new is, for our purposes, useless, however celebrated.

This has a sharp consequence for how you are judged. The sole criterion of evaluation in this module is the **rigour of your argument** — its premises, its inferences, and whether its conclusion actually follows — and nothing else. You will not earn marks for agreeing with Aristotle, nor lose them for demolishing him, provided the demolition is well argued. A conclusion the instructor dislikes, defended rigorously, beats a conclusion the instructor shares, defended badly. We therefore begin, in Chapter 1, with the tool that underlies every later chapter: how to tell a good argument from a bad one. The course then moves through three real tensions — what you think you want versus what you actually want (Chapter 2), freedom versus constraint (Chapter 3), and efficiency versus ethics (Chapter 4) — before turning the instruments on the idea of success itself (Chapter 5) and asking you to build and defend a thesis of your own (Chapter 6).

These same frameworks return in **Philosophy II** next semester, where they are put to work on harder contemporary cases — algorithmic discrimination, whistleblowing, environmental trade-offs — under the pressure of an adversarial jury. What you build here, you will be expected to wield there.

Contents

Preface	2
1 Evaluating philosophical arguments	3
1.1 What an argument is, and what it is not	3
1.2 The three checks	3
1.3 The cardinal rule: separate the conclusion from the argument	4
2 Three frameworks for authentic desire	4
2.1 Aristotle: eudaimonia and the whole life	5
2.2 Girard: desire is mimetic	5
2.3 Spinoza: conatus and the power to act	6
2.4 Using the three together	6
3 Freedom and responsibility	7
3.1 Epictetus: the dichotomy of control	7
3.2 Sartre: radical freedom, responsibility, and bad faith	8
3.3 Applying both to one dilemma	8
4 Three ethical frameworks under pressure	9
4.1 Kant: the categorical imperative	9
4.2 Mill: utilitarianism	9

4.3 Aristotle: virtue ethics	10
4.4 The point of the comparison	10
5 Success versus the good life	11
6 Constructing and defending a thesis	11
6.1 What a rigorous thesis contains	11
6.2 Steelmanning the objection	12
6.3 Defending under pressure	12

1 Evaluating philosophical arguments

Before any framework, the underlying skill. Everything in this module is delivered as argument, and the one thing you are graded on is whether an argument is any good — your own, a classmate’s, or Aristotle’s. So we start by learning to take an argument apart.

1.1 What an argument is, and what it is not

Worked example — two ways of disagreeing. Two students disagree about a course of action. The first says: “I just feel it’s wrong — I can’t explain why, but I do.” The second says: “It’s wrong because it uses a person merely as a tool, and using people as mere tools is wrong.” Only the second has made an **argument**. She has stated a general principle (using people as mere tools is wrong), a claim about the case (this action does that), and drawn a conclusion (so it is wrong). You can now engage her — challenge the principle, or deny that the case fits it. The first student has given you nothing to engage; you can only agree or shrug.

Definition 1 An **argument** is a set of **premises** offered in support of a **conclusion**. The premises are the claims taken as starting points; the conclusion is the claim they are meant to establish. An **inference** is the step that takes you from premises to conclusion. A report of a feeling, an assertion with no support, or a description of how things are felt to be is **not** an argument and cannot be evaluated as one.

The reason we insist on arguments rather than opinions is practical: only an argument can be **examined**. An opinion can only be shared or not shared. An argument has parts, and each part can be inspected and attacked.




Figure 1: The anatomy of an argument. The premises (blue) are the stated starting points; the conclusion (purple) is what they are meant to establish; the inference is the move between them. An argument can fail at any of these three places — a false premise, a faulty inference, or a conclusion that does not in fact follow even when each premise is granted.

1.2 The three checks

Evaluating an argument means running three checks, **independently of whether you like the conclusion**.

1. **Are the premises true, or at least defensible?** An argument built on a false starting point proves nothing, however neatly it proceeds.
2. **Does each inference follow?** Granting the premises, is each step actually licensed, or has something extra been quietly assumed?
3. **Does the conclusion follow from the premises as a whole?** Sometimes every individual step looks fine but a key term has shifted meaning along the way, or the conclusion is simply wider than the premises support.

A useful distinction here is between an argument being **valid** and being **sound**. An argument is valid if the conclusion would follow **supposing** the premises were true; it is sound if it is valid **and** the premises are in fact true. Validity is about the plumbing; soundness is about the plumbing **and** the water.

Worked example — valid but not sound. “All billionaires are happy; Smith is a billionaire; therefore Smith is happy.” The reasoning is **valid** — if both premises held, the conclusion would be forced. But the first premise is false, so the argument is **unsound** and tells us nothing about Smith. Validity is necessary but not sufficient: a watertight inference from a false premise is still worthless.

1.3 The cardinal rule: separate the conclusion from the argument

Common pitfall The single most common error in philosophical evaluation — and the one this module polices hardest — is letting your agreement or disagreement with the **conclusion** decide your verdict on the **argument**. A conclusion you cheer can rest on a terrible argument; a conclusion you find repugnant can rest on an excellent one. Your job is to evaluate the reasoning, not to score the team you support. When you catch yourself calling an argument “good” and notice you only mean “I agree with where it ends up,” start the three checks over.

This is also why, throughout the module, claims are tested **under adversarial pressure**: you state a position and peers attack it. A claim that survives serious objection has earned something; a claim never objected to has merely not yet been tested. Apply the three checks to your own positions as ruthlessly as to others’.

2 Three frameworks for authentic desire

The tension: what you believe you want versus what you actually want.

Worked example — the coveted job. A student is certain she wants a particular prestigious finance job. Pressed on why, the reasons thin out: the work itself does not especially interest her, the hours frighten her, but “it’s what the best people go for.” Does she want the job? In one sense, obviously — she is working hard for it. In another sense it is unclear the want is even hers. This is the gap the chapter lives in, and three frameworks each cut into it differently.

We will not ask which framework is “correct.” We will ask what each **reveals**. A desire that looks healthy through one lens can look borrowed, or diminishing, through another — and noticing that is the entire payoff.

2.1 Aristotle: eudaimonia and the whole life

Worked example — the pleasant evening versus the good life. An evening of effortless pleasure — good food, easy company, nothing demanded — is plainly enjoyable. Yet a life made of nothing but such evenings, year after year, strikes most people as somehow wasted, even if pleasant throughout. That intuition — that a life can be pleasant and still not **good** — is exactly what Aristotle’s framework is built to explain.

Definition 2 For Aristotle, the highest human good is **eudaimonia** — usually translated “flourishing” or “living well.” It is not a feeling or a momentary state but an **activity of the soul over a complete life**, lived in accordance with reason and virtue. A desire is **authentic**, on this view, to the extent that satisfying it contributes to such a flourishing life rather than to fleeting pleasure or external reward.

The diagnostic question Aristotle hands you is therefore: *does acting on this desire move me toward a life that is good taken as a whole?* Eudaimonia is the criterion. The coveted job from our example is to be evaluated not by the thrill of getting the offer but by whether the life it builds is, over decades, a flourishing one.

Remark Note what makes this a **tool** and not a slogan. “Flourishing” sounds vague until you use it as a test that spans a whole life rather than a moment — and that span is precisely what separates it from “whatever feels good now.” Its weakness is the mirror of its strength: it tells you to aim at a flourishing life without fully specifying what one contains.

2.2 Girard: desire is mimetic

Worked example — the toy two children fight over. Put a child alone in a room full of toys and one toy in the child’s hand, and the other toys stay ignored. Bring in a second child who picks up one of the ignored toys, and suddenly the first child wants **that** toy — not because the toy changed, but because someone else now wants it. The desire was copied. Adults are more sophisticated but, Girard argues, no different in structure.

Definition 3 René Girard’s thesis is that desire is **mimetic**: we do not, for the most part, desire objects directly and spontaneously. We **imitate** the desires of others, who serve as **models**. Desire runs through a triangle — **subject, model, object** — and the object is often wanted only because the model wants it.




Figure 2: Girard’s mimetic triangle. The subject appears to desire the object directly (the dashed line), but in fact the desire is routed through a model: the subject imitates the model, and so comes to want what the model wants. The closer the model, the more likely the shared desire turns into **rivalry** over the same object.

The diagnostic value here is **suspicion**. Girard’s question is: *is this desire even mine, or have I borrowed it from a model I am imitating?* The desires that feel most intensely one’s own are often the most thoroughly imitated — and because two people imitating each other end up wanting the same object, mimetic desire breeds **rivalry**. The coveted job, recall, was wanted because “it’s what the best people go for”: a textbook model-driven desire.

Common pitfall Mimetic desire is easy to spot in **other** people and almost invisible in oneself — that asymmetry is the whole difficulty. “My friends are obviously copying each other, but my ambition is authentic” is exactly the report Girard predicts you will give whether or not it is true. The framework is only useful when turned inward, which is uncomfortable.

2.3 Spinoza: conatus and the power to act

Worked example — two promotions. Two people take demanding promotions. A year on, the first is more capable, more connected, more able to make things happen — enlarged by the role. The second is anxious, dependent on the approval the title brings, and less able to act freely than before — diminished by it, though the title is impressive. Same external success; opposite effect on the person. Spinoza gives us a vocabulary for that difference.

Definition 4 For Spinoza, each thing strives to persevere in its own being; this striving is the **conatus**. **Joy** is the felt **increase** of one’s power to act; **sadness** its decrease. A desire is to be evaluated by whether acting on it **increases one’s genuine power to act** — and so one’s joy — or diminishes it.

The diagnostic question is internal: *does acting on this desire enlarge my power to act, or shrink it?* Unlike Aristotle’s, this criterion does not appeal to an external picture of the good life; unlike Girard’s, it does not ask where the desire came from. It asks only what the desire **does to you**. The first promotion above increases the power to act (joy); the second decreases it (sadness), whatever its prestige.

2.4 Using the three together

The point is not to choose one framework but to run all three over the same desire and watch them disagree, because each is blind to what the others see.

Framework	Question it asks of a desire	Its blind spot
Aristotle (eudaimonia)	Does satisfying it contribute to a flourishing life taken as a whole?	Says little about where the desire came from, or its felt effect on you now.
Girard (mimetic desire)	Is the desire even mine, or borrowed from a model I imitate?	Says nothing about whether a desire — borrowed or not — is actually good for you.
Spinoza (conatus)	Does acting on it increase or decrease my power to act?	Ignores both the origin of the desire and its place in a whole life.

Table 1: The three frameworks for authentic desire as competing diagnostic instruments. A single desire — the coveted finance job — can pass one test and fail another: it may be intensely mimetic (Girard suspicious), may shrink the person’s power to act (Spinoza negative), yet conceivably still fit a flourishing life (Aristotle neutral-to-positive). The disagreement is the information.

Remark Girard’s question — *is this desire mine?* — is simply invisible to Aristotle, who takes the desire as given and asks only where it leads. That is why we hold three instruments rather than one: each sees a dimension the others cannot.

3 Freedom and responsibility

The tension: freedom versus constraint. Are we free, given everything that pushes on us — circumstances, other people, our own histories? Two frameworks carve the territory in opposite ways, and applying **both** to one dilemma is the exercise.

3.1 Epictetus: the dichotomy of control

Worked example — the delayed flight. Your flight is cancelled and you will miss an important meeting. One traveller rages at the gate agent, refreshes the app, and is miserable for hours. Another accepts in seconds that the cancellation is not hers to undo, and turns to what **is** hers — rebooking, a calm message to the meeting, using the time. Same event; the difference is entirely in where each directed their effort.

Definition 5 Epictetus, a Stoic, opens from the **dichotomy of control**: some things are **up to us** and some are **not up to us**. Up to us are our judgements, choices, desires and responses. Not up to us are external events, our body, reputation, wealth, and the actions of others. Freedom and tranquillity come from investing effort and identity **only** in what is up to us, and accepting the rest.




Figure 3: Epictetus’s dichotomy of control. The Stoic discipline is to keep effort and self-worth inside the inner circle (what is genuinely up to us) and to accept the outer ring (what is not) rather than staking one’s peace on it. Misery, on this view, is the predictable result of investing oneself in the outer ring.

The diagnostic move is a sorting one: faced with a dilemma, draw the line. What here is genuinely up to me, and what am I merely suffering as if it were? Effort spent on the second is, for Epictetus, wasted, and the source of avoidable misery.

Common pitfall The dichotomy is easy to misread as passivity or as an excuse — “the outcome isn’t up to me, so I’ll do nothing.” That inverts it. What is up to you includes **all your actions**; only the **results**, which depend on the world and others, fall outside. Epictetus counsels full effort on the action and equanimity about the result — not withdrawal.

3.2 Sartre: radical freedom, responsibility, and bad faith

Worked example — “I had no choice”. A manager carries out an order he believes is wrong and explains: “I had no choice — it’s my job, those are the rules.” But he did have a choice: refuse, resign, protest, comply under protest. Each has a cost, and he chose the cheapest. “I had no choice” is false; what is true is “I chose this over the alternatives I disliked more.” Sartre’s whole framework lives in the gap between those two sentences.

Definition 6 For Sartre, the human being is **radically free** and therefore **radically responsible**: we are “condemned to be free.” We cannot **not** choose, since even refusing to choose is itself a choice. **Bad faith** is the self-deception by which a person hides this freedom from themselves — pretending that a role, a rule, a circumstance or one’s “nature” left no choice, in order to escape the responsibility freedom imposes.

The diagnostic question Sartre supplies is an accusation pointed inward: *where am I telling myself “I had no choice” in order to avoid owning a choice I am in fact making?* Every “I can’t” that is really an “I won’t” is a candidate for bad faith.

3.3 Applying both to one dilemma

Epictetus and Sartre pull in opposite directions, which is exactly why holding both is useful.




Figure 4: Two lenses on one dilemma. Epictetus sorts the situation into what is and is not within your control and counsels acceptance of the latter. Sartre refuses any such alibi: within the situation you are always choosing, and “circumstances left me no choice” is bad faith. Used together, Epictetus stops you wasting effort on the genuinely uncontrollable, while Sartre stops you hiding a real choice behind that same boundary.

The risk each guards against is the other’s blind spot. Lean only on Epictetus and you may file a real choice under “not up to me” and call it acceptance — Sartrean bad faith. Lean only on Sartre and you may torment yourself with responsibility for outcomes that genuinely were not yours to command. The exercise is to apply both to the same personal or professional dilemma and draw the line honestly: what here is truly outside my control, and where am I dressing up a choice as a constraint?

4 Three ethical frameworks under pressure

The tension: efficiency versus ethics, and the harder fact that doing right by one standard can mean doing wrong by another. Here we take **one** ethical case and run **three** frameworks over it, forcing each to state out loud **what it sacrifices**. No framework is free; the maturity is in naming the price.

Worked example — the case we will keep returning to. You manage a small team. You discover that quietly laying off one well-liked, loyal employee would let you save the jobs of five others whose roles are at risk, and keep the firm solvent. The one employee has done nothing wrong. Do you do it? Hold this single case in mind; each framework below answers it differently, and each pays a different price for its answer.

4.1 Kant: the categorical imperative

Definition 7 Kant locates rightness in the **principle** of an action, not its consequences. The **categorical imperative** has two formulations we will use: (i) act only on a maxim you could will to become a **universal law**; and (ii) treat humanity, in yourself and others, always also as an **end**, never **merely** as a means.

On the case: sacrificing the loyal employee to save the others treats him **merely as a means** — a lever for an outcome — which the second formulation forbids. Kant’s answer leans firmly against it. **What it sacrifices:** possibly a much better outcome. Because consequences do not enter the judgement, Kantian ethics can require an action that leads to worse results for everyone, the five included.

4.2 Mill: utilitarianism

Definition 8 Mill’s **utilitarianism** judges an action right “in proportion as it tends to promote happiness,” wrong as it tends to produce the reverse — counting **each person’s**

happiness equally and summing across everyone affected. The right act is the one producing the greatest happiness for the greatest number.

On the case: five jobs saved against one lost is, on the face of it, a clear utilitarian gain, so Mill’s framework leans toward the layoff. **What it sacrifices:** the individual. Because only the aggregate matters, utilitarianism can license real harm to one person whenever the sum comes out ahead — the single employee has no standing the total cannot override.

4.3 Aristotle: virtue ethics

Definition 9 **Virtue ethics** asks neither “what rule applies?” nor “what maximises good?” but “**what would a person of good character do here?**” The right action is the one a **virtuous** person, exercising **practical wisdom** (*phronesis*), would choose — characteristically a **mean** between excess and deficiency (e.g. courage between recklessness and cowardice).

On the case: virtue ethics asks what an honest, just, courageous and compassionate manager would actually do — perhaps seeking a fourth option, perhaps being transparent with the team, perhaps something the two rule-based frameworks never raise. **What it sacrifices:** a decision procedure. It gives no formula, and when it is genuinely unclear what the virtuous person would do, it offers little to settle the matter.

4.4 The point of the comparison

Framework	Where it locates rightness	What it sacrifices
Kant (deontology)	In the principle of the act: universalisable maxim; people as ends, not mere means.	Good consequences — it can forbid an act that would produce the best outcome.
Mill (utilitarianism)	In the consequences: greatest happiness for the greatest number, each counted equally.	The individual — it can license real harm to one if the aggregate gains.
Aristotle (virtue ethics)	In the agent’s character: what a person of practical wisdom would do.	A decision procedure — it gives no formula when the virtuous course is unclear.

Table 2: Three ethical frameworks on the same case. Applied to the layoff, they can recommend **different** actions (Kant against, Mill for, Aristotle “it depends on what good character requires here”); and even where two agree, they justify the agreement on incompatible grounds. The required intellectual move is to run all three and state each one’s price — not to adopt one and quietly ignore its cost.

Common pitfall The cheap move — and a common one under pressure — is to pick the framework that happens to endorse the answer you already wanted, and present it as **the** ethical view. That is the opposite of the skill being trained. The discipline is to apply all three honestly, including the one that condemns your preferred answer, and to acknowledge the cost of whichever you finally side with.

5 Success versus the good life

The tension: inherited definitions of success versus living well. We now turn the instruments built above onto a target most people never examine — their own definition of “success.”

Worked example — the inherited scoreboard. Ask someone what a successful life looks like and the answers arrive fast and pre-packaged: a high salary, an impressive title, recognition, the right postcode. Ask **why** these constitute success and the answers thin out. Mostly they were inherited — from family, peers, social media, the ambient culture — and never inspected. This is precisely the unexamined life Socrates warned of, operating at the level of one’s life goals.

Rather than accept the inherited scoreboard, we replace it with two **operational criteria** drawn from earlier chapters — criteria you can actually apply to a decision:

- **Eudaimonia** (Aristotle): does this contribute to a whole life lived well, or only to a momentary or external marker?
- **Power to act** (Spinoza’s conatus): does this increase my **genuine power to act**, or merely my **exposure to the judgement of others**?

Definition 10 The **success / good-life distinction** contrasts **inherited success** — wealth, status, recognition, measured by external markers and others’ judgement — with a **good life** assessed by eudaimonia and power to act. The two can **diverge**: a life high on the inherited scoreboard can be low on flourishing and power to act, and the reverse can hold too.

The method is to take a **real** scenario — a career choice, an ambition, a way of designing one’s life — and ask of it the Spinozist question in particular: does this increase my genuine power to act, or does it mainly increase my exposure to being judged by others? The coveted finance job from Chapter 2 returns here as a candidate; so does any goal pursued chiefly because it looks impressive.

Remark Notice this chapter is not anti-ambition. A demanding, high-status path can **also** increase your power to act and fit a flourishing life — the two need not conflict. The point is to make the inherited definition **answer for itself** against operational criteria, rather than be accepted by default. Sometimes it passes; the examined life does not require the conventional answer to be wrong, only that it be examined.

6 Constructing and defending a thesis

The module culminates in a **defended personal thesis** on what living well means. This is where every earlier tool is put to work at once, and it is judged — like everything here — solely on the rigour of the argument, never on whether the reader shares your conclusion.

6.1 What a rigorous thesis contains

Definition 11 A rigorous thesis for this module has four parts:

1. **A clear claim** about the good life — stated precisely enough that one could argue both for and against it.

2. **Two thinkers in productive tension** — two of the frameworks above mobilised **together** so that their disagreement does real work, rather than one cited as a decorative authority.
3. **A serious counter-argument** — the strongest objection to your claim, stated in its most forceful form.
4. **An answer to that objection**, delivered and then **defended live under peer objection**.

Worked example — a thesis using two thinkers in tension. Claim: “A good career is one that increases your power to act, even at the cost of status.” Here Spinoza (power to act) and Girard (status as mimetic, borrowed desire) are put in **tension** with the inherited success Girard helps unmask — and perhaps against Aristotle, if one argues a flourishing life needs the social goods that status brings. The two thinkers are not ornaments; the argument **runs on** their disagreement. Contrast a non-thesis: “Spinoza said joy is the increase of power to act, and I agree.” That cites an authority and does no work.

6.2 Steelmanning the objection

Common pitfall The most damaging shortcut is the **straw man**: stating a weak, easily-defeated version of the counter-argument so you can knock it down. Graders and peers see it instantly, and it reveals you have not understood the strongest case against yourself. The opposite — the **steel man** — is required: state the objection in the most forceful form you can, **as if you held it**, and only then answer it. A thesis that survives its own strongest objection is worth far more than one that defeats a feeble one.

6.3 Defending under pressure

The live defence reactivates Chapter 1. Under peer objection you must, in real time, run the three checks on the attack — is the objector’s premise true? does their inference follow? — and on your own reply. A good defence is not louder insistence on your conclusion; it is showing precisely **where** the objection fails, or honestly conceding the point and **repairing** the thesis. Conceding a fair point and adjusting is a mark of rigour, not weakness.

The through-line of the module. Chapter 1 gives you the test — premises, inferences, does the conclusion follow — and Chapters 2–5 give you the instruments: three lenses on desire (Aristotle, Girard, Spinoza), two on freedom (Epictetus, Sartre), three on ethics (Kant, Mill, Aristotle), and a way to interrogate “success” itself. Chapter 6 asks you to wield them on a thesis of your own and hold it under fire. Throughout, the thinkers are diagnostic instruments, not authorities, and the only thing being graded is whether your argument is any good. That is what it means, here, to examine a life.