

HUM24-3

Philosophy II — Ethics, Desire & Decision

Holding an ethical position under adversarial pressure

Philosophy as decision-training, not academic discipline.

Albert School — 12 hours

Contents

1 Scope and starting point	2
2 Why we decide badly	2
2.1 System 1 and System 2	2
2.2 Mimetic desire applied to career choices	2
2.3 Spinoza as structural corrective	3
3 Responsibility and the decision chain	3
3.1 The banality of evil as a mechanism	3
4 Jonas: foresight as a moral obligation	3
5 The three frameworks on four high-stakes cases	3
5.1 AI as revealing terrain	4
5.2 Jonas on the contested AI case	4
6 The cost of every position	4
7 The adversarial jury format	4
8 Summary	5

1 Scope and starting point

This course reactivates the Philosophy I toolkit on harder, professionally charged cases. It assumes the frameworks introduced in HUM14-3 and does not re-teach them: authentic desire (Aristotle's eudaimonia, Girard's mimetic desire, Spinoza's conatus), freedom and responsibility (Epictetus's dichotomy of control, Sartre's radical responsibility and bad faith), and the three ethical frameworks under adversarial pressure (Kantian deontology, Millian utilitarianism, Aristotelian virtue ethics).

Here those tools are wielded on cases where decisions are made badly, responsibility diffuses across a chain, and ethical positions must hold under adversarial pressure: algorithmic discrimination in recruitment, whistleblowing, conflict of interest in client advice, environmental decisions, industrial disasters. Philosophy is treated as decision-training rather than as an academic discipline. The course is the Humanities pillar's second philosophy stage, and its target is the ability to hold — and defend — an ethical position in ambiguous organisational contexts where diffusion of responsibility is the default failure mode.

2 Why we decide badly

Real decisions are not made by an ideal rational agent. Three frameworks diagnose why.

2.1 System 1 and System 2

Definition 1 Kahneman distinguishes two modes of thinking. **System 1** is fast, automatic, and effortless; **System 2** is slow, deliberate, and effortful. Most real decisions default to System 1.

The consequence the course insists on: cognitive biases are not rare errors that occasionally corrupt otherwise sound reasoning. They are the **default** of real decision-making. A decision left to System 1 is the normal case, not the exception.

2.2 Mimetic desire applied to career choices

Girard's mimetic desire (introduced in Philosophy I as a framework for authentic desire) is here applied specifically to career choices — consulting, finance, startups. The operative distinction is between **borrowed ambition** (a desire copied from a model or a peer group) and **authentic**

desire (a desire one can claim as one's own). The point is to deconstruct an effective career choice and ask whether it was driven by borrowed desire rather than considered judgment.

2.3 Spinoza as structural corrective

Definition 2 Spinoza, **Ethics** IV.7: a weaker affect can be overcome only by a stronger and clearer affect — not by willpower. An emotion is not defeated by deciding to defeat it, but only by a more powerful and better-understood emotion.

This is the structural corrective to the first two diagnoses: if a decision is driven by a System 1 default or a borrowed desire, the remedy is not an act of will against it, but a stronger and clearer affect that displaces it.

3 Responsibility and the decision chain

3.1 The banality of evil as a mechanism

Definition 3 Arendt's **banality of evil**: harm produced not by hatred or malice but by ordinary functional obedience — people doing their job within a chain, without evil intent.

The course treats this strictly as a **mechanism** to be recognised in oneself — not as an excuse to be invoked. “I was only doing my job” is precisely the mechanism Arendt names; it is therefore not available as a defence. In a decision chain (an industrial disaster, an instance of data manipulation, a discriminatory algorithm), the task is to identify the precise point at which individual responsibility should have activated, and what prevented it from activating.

4 Jonas: foresight as a moral obligation

Definition 4 Jonas's responsibility principle: when consequences are irreversible and operate at large scale, **foresight becomes a moral obligation**. The required foresight is proportional to the technological irreversibility of what is being done.

This is applied to irreversible large-scale consequences: industrial disasters, data manipulation, and discriminatory algorithmic optimisation. Where an action's effects are large and cannot be undone, the obligation to foresee them is itself part of the ethical assessment, alongside the action's intent and its outcomes.

5 The three frameworks on four high-stakes cases

The three Philosophy I ethical frameworks are re-deployed here:

Definition 5 **Kantian deontology** — act only on a maxim you could will as a universal law; treat persons as ends, never merely as means. **Millian utilitarianism** — choose the action producing the greatest aggregate well-being. **Aristotelian virtue ethics** — act as the person of practical wisdom would act, from a stable disposition of character.

They are applied to four high-stakes professional cases:

1. algorithmic discrimination in recruitment;
2. whistleblowing on an illegal employer practice;

3. conflict of interest in client advice;
4. an environmental decision with deferred social impact.

5.1 AI as revealing terrain

The course frames algorithmic decisions as a revealing terrain rather than a new ethical domain: an algorithmic decision makes **explicit and at scale** what human decisions were already doing **implicitly and at small scale**. The algorithm does not introduce the judgment logic; it exposes a tacit logic that was always operating. This reframing is what lets the same three frameworks apply.

5.2 Jonas on the contested AI case

On a contested AI-deployment case specifically, Jonas's responsibility principle is brought into the deliberation alongside the three frameworks. For that case, each framework — Kant, Mill, Aristotle, and Jonas — must be made to say what it **obliges**, what it **permits**, and what it **sacrifices** in that specific AI context.

6 The cost of every position

Definition 6 No ethical position is stated without its assumed cost. Each defended framework must be paired with an explicit acknowledgement of what it sacrifices.

This is a requirement on every position taken in the course, not an optional addition. Defending Kantian deontology means naming the consequences it is willing to accept to keep the rule; defending utilitarianism means naming the individuals or rights it is willing to override for aggregate welfare; defending virtue ethics means naming the determinacy it gives up. A position stated without its cost is incomplete.

Common pitfall Acknowledging a position's cost is **not** abandoning the position. Under jury pressure, the move is to name the limit and hold the position anyway — not to collapse into the objecting framework, and not to caricature it.

7 The adversarial jury format

The course culminates in an adversarial jury format that replicates the conditions of real high-stakes professional deliberation (investment committee, board, advisory mission).

Definition 7 Format. A one-page defended ethical position, an 8-minute presentation, then 12 minutes of targeted objections from two designated opponent peers. The test is coherence under pressure, not a static submission.

Defending a position under this format means: maintaining internal coherence, acknowledging limits without abandoning the position, and responding to targeted objections without collapsing or caricaturing the opposing view.

8 Summary

- Decisions default to System 1; biases are the norm. Girard separates borrowed ambition from authentic desire in career choices; Spinoza (Ethics IV.7) shows affects are displaced by stronger, clearer affects, not by willpower.
- Arendt's banality of evil is a mechanism of ordinary obedience to be recognised in oneself, never invoked as an excuse; in a decision chain, locate where responsibility should have activated.
- Jonas: where consequences are irreversible and large-scale, foresight is a moral obligation proportional to technological irreversibility.
- Kant, Mill, and Aristotle apply to algorithmic discrimination, whistleblowing, conflict of interest, and deferred-impact environmental decisions; AI makes explicit and at scale what human judgment did implicitly. On a contested AI case, add Jonas and state what each framework obliges, permits, and sacrifices.
- Every position is paired with its cost; a position is defended under an adversarial jury (one page, 8-minute presentation, 12 minutes of objections) by holding coherence under pressure.