

HUM24-3

Philosophy II — Ethics, Desire & Decision

Albert School — 12 hours

Preface

In Philosophy I you built a small set of instruments and learned to wield them on your own life: three lenses on desire (Aristotle’s eudaimonia, Girard’s mimetic desire, Spinoza’s conatus), two on freedom (Epictetus’s dichotomy of control, Sartre’s radical responsibility and bad faith), and three on ethics (Kant, Mill, Aristotle) run over a single case until each was forced to name what it sacrifices. The one thing you were graded on was the rigour of the argument, never whether anyone shared your conclusion. This course assumes all of that and does not re-teach it. If “categorical imperative,” “greatest happiness,” “practical wisdom,” “bad faith,” or “mimetic desire” do not yet name working tools for you, go back to Philosophy I before going forward here.

What changes in Philosophy II is the terrain, and the change is not cosmetic. In Philosophy I the cases were yours — your ambition, your dilemma, your definition of success — and you faced them alone. Here the cases are **professional and shared**: algorithmic discrimination in a recruitment pipeline, whistleblowing on an illegal practice, a conflict of interest in client advice, an environmental decision whose harm lands on people far away and years later, an industrial disaster nobody intended. Three things are true of these cases that were not true before. First, the decisions are routinely made **badly** — not by villains, but by ordinary competent people running on cognitive autopilot. Second, responsibility **diffuses**: it is split across a chain of hands such that no single person feels they did it, and so no one does anything. Third, your position will be **attacked** — not by a sympathetic instructor, but by peers assigned to break it.

That last fact sets the form of the course. It treats philosophy as **decision-training**, not as an academic discipline, and it culminates in an **adversarial jury**: you write a one-page defended position, present it for eight minutes, and then absorb twelve minutes of targeted objections from two classmates whose job is to make it fail. This replicates the real rooms where high-stakes decisions are taken under pressure — an investment committee, a board, an advisory mission — where the failure mode is almost never “nobody knew the right framework” and almost always “everybody had a framework and nobody held it when it cost something.” The aim of these twelve hours is narrow and hard: to be the person in that room who can hold an ethical position, name its price, and not collapse when pushed.

The book runs in five movements. Chapter 1 diagnoses **why we decide badly** — the cognitive default, borrowed desire, and the structural reason willpower is the wrong remedy. Chapter 2 follows responsibility as it **diffuses across a chain** (Arendt). Chapter 3 adds the one genuinely new ethical instrument, **Jonas**, for decisions whose consequences are irreversible and large. Chapter 4 re-deploys the three Philosophy I frameworks on the four professional cases, uses AI as the terrain that **reveals** what human judgment was always doing, and makes the cost of every position non-negotiable. Chapter 5 is the jury itself.

Contents

Preface	2
1 Why we decide badly	3
1.1 Kahneman: System 1 is the default, not the exception	3
1.2 Girard, redeployed: borrowed ambition in career choices	4
1.3 Spinoza: why willpower is the wrong remedy	4
2 Responsibility and the decision chain	5
2.1 Arendt: the banality of evil as a mechanism	5
2.2 Locating the point where responsibility should have activated	6
3 Jonas: foresight as a moral obligation	6
4 Three frameworks on four high-stakes cases	8

4.1	The cost of every position is mandatory	8
4.2	AI as revealing terrain	9
4.3	Jonas on the contested AI-deployment case	10
5	The adversarial jury	10
5.1	What defending well actually means	11
5.2	Rehearsing alone, before the room	12

1 Why we decide badly

The starting fact: real decisions are not made by the ideal rational agent that ethics quietly assumes. Before asking what the right choice is, this chapter asks why the actual choice so often goes wrong — because a remedy aimed at the wrong cause does nothing. Three instruments diagnose three distinct failures: the **cognitive** default (Kahneman), the **borrowed** desire (Girard), and the **structural** reason we cannot simply will our way out (Spinoza).

1.1 Kahneman: System 1 is the default, not the exception

Worked example — the hire that felt obviously right. A manager interviews a candidate, feels an immediate click — “great fit, sharp, one of us” — and hires on the strength of it. Pressed later for reasons, she assembles them: confident, articulate, the right background. But the reasons came **after** the verdict, not before it; the decision was made in the first thirty seconds and the reasoning was recruited to defend it. The candidate happened to resemble her, and people who resemble us read as “competent” to a fast, effortless part of the mind that never announces itself as biased.

Definition 1 Kahneman distinguishes two modes of thought. **System 1** is fast, automatic, effortless, and always on: it produces impressions, intuitions, and snap judgements without any sense of effort. **System 2** is slow, deliberate, and effortful: it does the explicit reasoning we like to imagine ourselves doing. The central, uncomfortable claim is that **System 1 is the default driver** of real decisions, and System 2 mostly steps in afterwards to justify what System 1 already decided.

The point the course presses, and the reason this is not just psychology, is a shift in status. The naive picture treats **cognitive biases** as rare faults — occasional glitches that corrupt an otherwise sound reasoning process. Kahneman inverts it: biases are not the exception, they are the **normal operating mode**. The manager above was not malfunctioning; she was running exactly as human cognition runs by default. This matters ethically because it dissolves a comfortable alibi. “I considered it carefully and decided” is, for most decisions, simply false — and an ethics that addresses only the deliberate, considered choice misses where the damage is actually done.




Figure 1: Kahneman’s two systems. In the naive picture, System 2 reasons and then decides. In the realistic picture the course adopts, System 1 makes the decision and System 2 is recruited afterwards to supply reasons — which is why a biased choice can feel, from the inside, like a carefully reasoned one.

1.2 Girard, redeployed: borrowed ambition in career choices

Girard’s mimetic desire entered Philosophy I as a general lens on authentic desire. Here it is aimed at one specific target where it does the most damage: **career choice**.

Worked example — the three-door funnel. A cohort of capable students, asked what they want, converge with striking uniformity on the same three doors — consulting, finance, the funded startup. Pressed individually, few can say what about the work itself draws them; what they can say is who **else** is going for it, and that it is what the strongest students do. The desire is real — they will work brutally hard for it — but it is not clear it is **theirs**. It was caught from the surrounding models, the way the toy in the other child’s hand suddenly becomes the one toy worth having.

Definition 2 Girard’s thesis, recall, is that desire is **mimetic**: we largely do not want objects directly but **imitate** the desires of **models**. Applied to careers, the operative distinction is between **borrowed ambition** — a goal copied from a peer group or a model and mistaken for one’s own — and **authentic desire**, a goal one can actually claim. The diagnostic task is to take a real, effective career choice and ask whether it was driven by borrowed desire rather than by considered judgment.

This is the same move Kahneman makes, one level up. Kahneman shows the individual **decision** runs on autopilot; Girard shows the **goal the decisions serve** can be on autopilot too — adopted by imitation and never inspected. The two compound: a borrowed ambition pursued through System-1 decisions is a life run almost entirely on imported software.

Common pitfall Mimetic desire is trivial to spot in others and nearly invisible in oneself — “my friends are obviously copying each other, but **my** ambition is genuine” is exactly the report Girard predicts you will give whether or not it is true. In the jury, the cheap deflection is to diagnose mimetic desire in the **case** (the reckless trader chasing what his desk chases) while exempting your own position from the same suspicion. The instrument is only honest when turned inward first.

1.3 Spinoza: why willpower is the wrong remedy

Diagnosing a System-1 default or a borrowed desire invites an obvious, and wrong, response: **resolve to do better** — just decide, by force of will, to override the bias and drop the borrowed goal. Spinoza explains structurally why that almost never works, and what does.

Worked example — the resolution that fails by March. Someone recognises that the prestigious path they are on is borrowed and hollowing them out. They resolve, by sheer willpower, to want something else. Within weeks the resolution erodes: the old desire is fed daily by every signal of status and peer approval around them, while the new intention is fed by nothing but a decision made once. The stronger, constantly-reinforced affect wins. Willpower was never going to be enough, and its failure is not a character flaw — it is structurally predictable.

Definition 3 Spinoza, **Ethics** IV, Proposition 7: “An affect cannot be restrained or removed except by an affect opposite to, and stronger than, the affect to be restrained.” A desire or emotion is not overcome by deciding to overcome it — by bare willpower — but only by a **stronger and clearer** affect that displaces it.

This is the **structural corrective** to the first two diagnoses, and it reframes the whole problem. If a decision is driven by a System-1 default or a borrowed desire, the question is not “how do I summon the willpower to resist it?” but “what stronger, clearer affect can I cultivate that **outcompetes** it?” Resistance pits a weak deliberate intention against a strong automatic pull and loses; displacement puts a stronger pull on the other side. The remedy for a borrowed ambition is not to grit one’s teeth against it but to develop a more powerful, better-understood desire that makes the borrowed one quietly lose its grip.



Figure 2: Spinoza’s **Ethics** IV.7 applied. Willpower sets a weak deliberate intention against a strong, constantly-fed affect, and predictably loses (left). The only thing that reliably overcomes an affect is a stronger and clearer affect that displaces it (right). The remedy for a bad desire is therefore a better desire, not more resolve.

2 Responsibility and the decision chain

The shift: in Philosophy I responsibility was **personal** — yours, faced alone, with Sartre forbidding the “I had no choice” alibi. The professional cases add a structural complication Sartre alone cannot handle: harm produced by a **chain** of people, each doing a small, lawful, unremarkable part, none of whom intended it or feels they caused it. This is where responsibility goes to die, and Arendt is the instrument for finding the body.

2.1 Arendt: the banality of evil as a mechanism

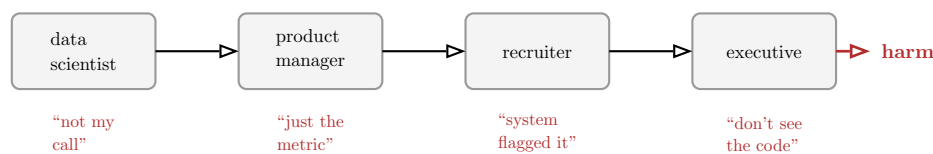
Worked example — the discriminatory pipeline, distributed. A recruitment algorithm systematically screens out qualified women. Trace the chain. A data scientist trained it on the firm’s past hires — “just historical data, not my call what it’s used for.” A product manager shipped it — “I optimise the metric I’m given.” A recruiter trusts its rankings — “the system flagged them, not me.” An executive approved the budget — “I don’t see the code.” Every link is individually reasonable, lawful, and free of any animus toward women. No one hates anyone. And yet the harm is real, systematic, and at scale. Where, exactly, did it happen?

Definition 4 Arendt’s **banality of evil** names harm produced not by hatred, malice, or monstrous intent but by **ordinary functional obedience** — competent people performing their defined role within a system, thoughtlessly, without asking what the whole produces. The evil is “banal” precisely because its agents are not demons but functionaries; the engine of the harm is unremarkable, dutiful role-following.

The course is emphatic on one point, and it is easy to get backwards. The banality of evil is a **mechanism to be recognised in oneself** — never an **excuse to be invoked**. The sentence “I was only doing my job” is not a defence against Arendt’s charge; it is the **exact phrasing of the mechanism she describes**. To say it is to confess the diagnosis, not to escape it. The whole value of the instrument is pre-emptive and self-directed: it lets you catch yourself becoming a thoughtless link before the harm completes, by noticing the tell-tale phrases — “not my call,” “just following the spec,” “above my pay grade” — as they form in your own mouth.

2.2 Locating the point where responsibility should have activated

In a diffuse chain, the productive question is not the unanswerable “who is to blame?” but the precise, operational one the course trains: **at which link should individual responsibility have activated, and what prevented it?**



Each alibi passes responsibility **down the line**; it arrives nowhere.

Figure 3: Diffusion of responsibility across a decision chain. Each link performs a small, lawful, reasonable task and hands off to the next, and each has a ready phrase that relocates responsibility elsewhere. The harm emerges from the chain as a whole while every individual feels exonerated. The analytic task is to mark the link where a thoughtful agent should have stopped and asked what the whole was doing — and to name what (incentive, deference, diffusion itself) prevented it.

Common pitfall Diffusion of responsibility is **the default failure mode** of organisational decisions, not a rare pathology — which is why “responsibility diffuses across a chain” is the through-line of this entire course. Two opposite errors await. One is to let the diffusion stand — “everyone is a bit responsible,” which operationally means no one is. The other is to collapse it onto a single scapegoat and call the problem solved. The discipline is to name the **specific** link and the **specific** mechanism that disabled responsibility there, without either dissolving the accountability or faking a tidy single culprit.

3 Jonas: foresight as a moral obligation

The new instrument. The Philosophy I frameworks were built for decisions whose consequences are roughly contemporaneous and reversible. The professional cases include decisions that are neither — an industrial process, a data manipulation, a discriminatory algorithm running at scale — where the harm is **large** and **cannot be undone**. For these, Hans Jonas adds an obligation the classical frameworks underweight.

Worked example — the deferred, irreversible harm. An engineering team can ship a chemical process, a data pipeline, or an optimisation algorithm now. The immediate effects are benign and the gains are real. But there is a credible, hard-to-quantify path by which the thing causes large, **irreversible** harm years later — a contamination, an entrenched discriminatory feedback loop, a manipulation that erodes something that does not grow back. Classical cost–benefit reasoning, run on the visible near-term numbers, says ship. Jonas says: the very irreversibility and scale of the possible harm changes what you are **obliged to do before acting**.

Definition 5 Jonas’s **responsibility principle** holds that when an action’s consequences are **irreversible** and operate at **large scale**, **foresight itself becomes a moral obligation** — and the required foresight is **proportional to the technological irreversibility** of what is being done. The more powerful and less reversible the technology, the heavier the duty to anticipate its downstream effects **before** acting, giving systematic weight to the worse credible outcome.

The decisive move is that Jonas makes **foresight** part of the ethical assessment, alongside intent (Kant’s home ground) and outcome (Mill’s). For a small, reversible act, failing to foresee a remote harm is a minor matter — you can correct course. For a large, irreversible one, the **failure to look ahead is itself a wrong**, independent of how the gamble turns out. This is the instrument that addresses industrial disasters, data manipulation, and discriminatory algorithmic optimisation specifically: in each, the harm is large and not cleanly reversible, so the duty to foresee is correspondingly heavy.




Figure 4: Jonas’s principle as a proportionality. The obligation to foresee downstream consequences rises with the scale and irreversibility of the action. For trivial, reversible acts the duty is light — you can undo a mistake. For large, irreversible technologies it is heavy, and a failure to look ahead is itself a moral failing, regardless of how the outcome happens to land.

Common pitfall Jonas is not a veto on all powerful technology, and reading him as “never do anything irreversible” caricatures the principle into uselessness. He raises the **burden of foresight** in proportion to the stakes; he does not set it to infinity. In the jury, the disciplined use is to say what **specific** foreseeing the irreversibility obliged — and what it would cost to discharge that duty — not to invoke “irreversibility!” as a conversation-ending trump.

4 Three frameworks on four high-stakes cases

This is the working core of the course. The three Philosophy I ethical frameworks — Kant, Mill, Aristotle — are re-deployed, unchanged, on four professional cases, with Jonas brought in where irreversibility is in play. The frameworks themselves are assumed; what is trained here is their **application under stakes**, and the non-negotiable accounting of **what each costs**.

Definition 6 The three frameworks, in one line each. **Kantian deontology**: act only on a maxim you could will as a universal law, and treat persons always also as ends, never **merely** as means — rightness lies in the principle, not the outcome. **Millian utilitarianism**: choose the action producing the greatest aggregate well-being, each person counted equally — rightness lies in the consequences. **Aristotelian virtue ethics**: act as the person of practical wisdom (*phronesis*) would act, from a stable disposition of good character — rightness lies in the agent.

The four cases on which they are exercised:

1. **algorithmic discrimination in recruitment** — a hiring system that systematically disadvantages a protected group;
2. **whistleblowing on an illegal employer practice** — speaking out at personal cost against an unlawful act by your own organisation;
3. **conflict of interest in client advice** — a recommendation that serves your incentives at the client's expense;
4. **an environmental decision with deferred social impact** — a choice whose benefits are near and local but whose harms are distant in time and fall on others.

4.1 The cost of every position is mandatory

Worked example — the whistleblower, three ways. You discover your employer is doing something illegal that harms outsiders. Should you blow the whistle, at serious risk to your career? **Kant**: the practice cannot be universalised and complicity treats the harmed as mere means — you must not be part of the lie; lean toward disclosure. **Mill**: weigh the aggregate — harm prevented against careers and livelihoods damaged, including your own and colleagues' — and the sum often, though not always, favours disclosure. **Aristotle**: what would a person of integrity and courage actually do — perhaps escalate internally first, perhaps go straight to the regulator, depending on the particulars? Each gives an answer. Each, crucially, also gives up something.

Definition 7 The **cost principle**: no ethical position may be stated without an explicit acknowledgement of **what it sacrifices**. A defended framework is incomplete until its price is named. Defending Kant means naming the **consequences** you will accept to keep the principle intact. Defending Mill means naming the **individual or right** you will override for the aggregate. Defending Aristotle means naming the **determinacy** — the clean decision procedure — you give up.

This requirement is the spine of the whole course, not an add-on. It is what separates a defensible position from a slogan, and it is what the jury will probe first. A student who says “utilitarianism clearly favours shipping the algorithm” has not finished; the position exists only once they add “—and I accept that this overrides the specific people it wrongs, whom the aggregate is allowed to outvote.”

Framework	Where it locates rightness	What it sacrifices (the cost)
Kant (deontology)	In the principle: universalisable maxim; persons as ends, never merely means.	Good consequences — it can forbid the act with the best overall outcome.
Mill (utilitarianism)	In the consequences: greatest aggregate well-being, each counted equally.	The individual — it can license real harm to one if the sum comes out ahead.
Aristotle (virtue ethics)	In the agent: what a person of practical wisdom would characteristically do.	A decision procedure — no formula when the virtuous course is genuinely unclear.
Jonas (responsibility)	In foresight: duty to anticipate, scaled to irreversibility and reach.	Speed and certainty — it can block or delay action under deep uncertainty.

Table 1: The four instruments on any high-stakes case. Applied to one case they can recommend **different** actions, and even where they agree they justify it on incompatible grounds. The rightmost column is mandatory: a position is only complete once its cost is named. Jonas is added to the panel specifically when the case involves irreversible, large-scale consequences.

Common pitfall The cheapest move under pressure — and the one the jury is built to punish — is **framework-shopping**: quietly selecting whichever framework endorses the answer you already wanted, presenting it as **the** ethical view, and never mentioning its cost. This is the exact opposite of the trained skill. Run all three (or four) honestly, **including the one that condemns your preferred answer**, and pay the price of whichever you finally hold.

4.2 AI as revealing terrain

The course gives algorithmic cases a particular framing, and it is worth stating sharply because it is easy to get wrong.

Worked example — the bias that was always there. Before the algorithm, a hundred recruiters across a firm made hiring calls on gut feel — System 1, each slightly biased, each invisible, the bias diffused and deniable across a hundred private judgements. Then the firm trains an algorithm on those past decisions to “remove human bias.” The algorithm reproduces the bias — but now it is **one** rule, **written down**, applied identically to every candidate, auditable, and operating at the scale of the whole pipeline. Nothing new was introduced. The tacit logic of a hundred gut calls was made **explicit and scaled** — and only now is it visible enough to be objected to.

Definition 8 AI as revealing terrain. An algorithmic decision makes **explicit and at scale** what human decisions were already doing **implicitly and at small scale**. The algorithm does not introduce a new ethical domain or a new judgment logic; it **exposes** a tacit logic that was always operating, and its scale and legibility are what make that logic finally available for ethical scrutiny.

This reframing is what licenses the whole chapter: it is **because** the algorithm only externalises pre-existing human judgment that the same three frameworks apply without modification. The lesson runs in both directions. The algorithm is not a new kind of moral agent to which old ethics is irrelevant; and the comfortable human status quo it replaces was not innocent — it was the same logic, merely too diffuse to indict. AI is a **mirror**, and the discomfort it produces is recognition.

4.3 Jonas on the contested AI-deployment case

For a **contested AI-deployment case** specifically — where reasonable people disagree about whether to ship — the deliberation must run the three S1 frameworks **and** Jonas, because algorithmic decisions at scale are exactly the irreversible, far-reaching kind Jonas was built for. The required output is precise: for that case, make each of the four say what it **obliges**, what it **permits**, and what it **sacrifices**.

On the AI case	Obliges	Permits	Sacrifices
Kant	Refuse any design that treats candidates merely as means or that you could not universalise.	Deployment whose rule respects every person as an end.	Aggregate efficiency gains that would require wronging some individual.
Mill	Ship if and only if total well-being is maximised, counting all affected.	Harm to a minority if the net sum is positive.	The overridden individuals and rights the aggregate outweighs.
Aristotle	Act as a person of practical wisdom and integrity would, attentive to context.	Case-by-case judgment rather than a fixed rule.	A clean decision procedure; determinacy under disagreement.
Jonas	Foresee the irreversible large-scale effects before deploying; weight the worst credible outcome.	Deployment once that foresight is genuinely discharged.	Speed and certainty; may block action under deep uncertainty.

Table 2: The four-instrument deliberation on a contested AI deployment. The task is not to pick a winner but to make each framework state, for **this** case, what it obliges, what it merely permits, and what it sacrifices — Jonas included, because deployment at scale is precisely the irreversible, far-reaching kind of action his principle addresses.

5 The adversarial jury

Where it all has to hold. Everything above is rehearsal for one event: a position, written and then defended live against people trying to break it. This is the course’s claim about what ethical competence actually is — not the ability to **state** a view in calm, but the ability to **hold** one under fire without collapsing or caricaturing.

Definition 9 The adversarial jury format. Each student submits a **one-page defended ethical position** on an assigned high-stakes case, presents it for **8 minutes**, and then faces

12 minutes of targeted objections from two designated opponent peers whose explicit role is to find the weakest point and press it. The test is **coherence under pressure**, not a polished static submission.

The format is a deliberate replica of the rooms where consequential decisions are actually made — an investment committee, a board, an advisory mission — where the binding constraint is never eloquence on a calm afternoon but whether a position survives someone competent and motivated trying to dismantle it in real time.




Figure 5: The jury format as a timeline. A one-page written position precedes an 8-minute presentation, followed by 12 minutes in which two designated opponents press targeted objections. Most of the graded time is the part you do not control — by design, since holding a position under hostile questioning is the competence being certified.

5.1 What defending well actually means

A live defence is not louder insistence on your conclusion. It is three specific disciplines, each the answer to a way positions usually fail.

Definition 10 Defending a position under the jury means, concretely:

1. **Maintaining internal coherence** — your answers under fire must stay consistent with each other and with the written position; you cannot quietly switch framework mid-defence to dodge an objection.
2. **Acknowledging limits without abandoning the position** — naming the cost and the hard cases honestly, and then **holding** the position anyway, rather than treating each concession as a reason to collapse.
3. **Responding to objections without collapsing or caricaturing** — engaging the **strongest** form of the objection and answering it, neither folding at the first push nor knocking down a straw version of what the opponent said.

Common pitfall Two symmetrical failures sink most defences, and the jury is calibrated to catch both. The first is **collapse**: an opponent names your position’s cost — “so your Kantian stance lets the five lose their jobs” — and you treat the cost as a refutation and abandon ship. But you **already conceded that cost when you stated the position**; having it repeated is not new information and is not a defeat. The second is **caricature**: you defend by misrepresenting the objection as something flimsier than it is. Both read instantly to a competent jury. Hold the named cost without flinching; engage the real objection, not a weaker stand-in.

Worked example — holding the cost under fire. You defend shipping a beneficial algorithm on utilitarian grounds. An opponent presses: “It demonstrably wrongs this specific subgroup.” Collapse would be: “—you’re right, I withdraw.” Caricature would be: “You just want to ban all technology.” Holding is: “Yes — I acknowledged that cost at the outset: my position overrides that subgroup’s claim for the larger aggregate gain, and here is **why I**

judge the gain to warrant that specific price. If you can show the aggregate doesn't actually come out ahead, you've refuted me; that the cost **exists** I already granted." That is coherence under pressure: the cost named, the position held, the real objection met on its strongest ground.

5.2 Rehearsing alone, before the room

The jury needs opponents, but everything that earns a position is built **beforehand**, on your own. The exercises below rehearse each instrument of the course in the order you will need it on the day; work them in writing, not in your head — a position that was never written down has never actually been tested. Pick **one real case** you could plausibly be assigned (a recruitment algorithm, a whistleblowing call, a conflict of interest, an environmental trade-off) and carry it through all five.

Exercises

- 1. Deconstruct the decision (Kahneman + Girard).** Take a real career or professional choice you have actually made. Write the verdict you reached, then honestly reconstruct the **order**: did the reasons come before the choice or after it? Name one System-1 pull and one **borrowed** desire — a goal you can trace to a model or peer group — that drove it. The test of honesty: the mimetic desire you name must be **your own**, not one you spotted in someone else.
- 2. Find the structural remedy (Spinoza).** For the borrowed desire you just named, write the willpower-only resolution that would predictably fail by March, and then the **stronger, clearer affect** that could actually displace it. State why the second outcompetes the first rather than merely opposing it.
- 3. Locate the link (Arendt).** For your chosen case, draw the chain of hands the harm passes through. Mark the **one** link where a thoughtful agent should have stopped and asked what the whole was producing, and name the **specific** mechanism — incentive, deference, diffusion — that disabled responsibility there. Avoid both failures: do not let “everyone is a bit responsible” stand, and do not invent a single villain.
- 4. Run all four, name the costs.** Make Kant, Mill, Aristotle, and Jonas each state, for your case, what it **obliges**, what it **permits**, and what it **sacrifices** — including the framework that condemns the answer you prefer. If your case is not irreversible and large-scale, say in one line why Jonas adds little here; if it is, say what **specific** foresight he obliges.
- 5. Write and then attack the one page.** State your defended position in a single page: the action, the framework you hold, and — non-negotiably — the cost you accept to hold it. Then switch sides and write the **two strongest** objections an opponent would raise. For each, draft the held answer: not collapse (“I withdraw”), not caricature (“you’d ban everything”), but the cost owned and the real objection met. If your best objection has no held answer, your position is not yet defensible — revise it, not the objection.

The through-line of Philosophy II. You decide badly by default — System 1 runs the choice (Kahneman), borrowed desire sets the goal (Girard), and willpower is the wrong fix because only a stronger, clearer affect displaces an affect (Spinoza). Harm in organisations is produced by **thoughtless functional obedience** diffused across a chain (Arendt): the task is to locate the link where responsibility should have activated, and “I was only doing my job” is the mechanism, not an excuse. Where consequences are irreversible and large, **foresight itself becomes an obligation** proportional to the stakes (Jonas). Kant, Mill,

and Aristotle re-deploy onto algorithmic discrimination, whistleblowing, conflict of interest, and deferred-impact environmental decisions; AI **reveals**, explicit and at scale, the tacit logic human judgment always ran; on a contested AI case, add Jonas and make each framework say what it obliges, permits, and sacrifices. Every position carries its cost, stated up front — and is then **held** before an adversarial jury (one page, 8 minutes, 12 minutes of objections) by keeping coherence, owning limits without collapsing, and meeting the real objection rather than a caricature.