

HUM64-2

Public Policy & Institutional Analysis

Albert School — 12 hours

Preface

This is the capstone of the Law & Regulation thread. In *Law Fundamentals & GDPR* (HUM24-2) you learned to read an individual rule as institutional architecture — designed by humans, encoding particular values, always incomplete. In *Digital Threats & Cybersecurity* (HUM34-3) you watched that architecture meet an adversary and learned to analyse the governance of a whole domain, naming who regulates, who enforces, and where the gaps are. This course takes the final step up the ladder of abstraction: from reading rules and domains to reasoning about the **institutional architecture of governance itself** — why institutions exist at all, how a regulatory agency balances independence against accountability, and why a policy that succeeds in one country fails when copied into another.

Three intellectual tools run through the whole module. The first is **institutional economics** in Douglass North’s sense: institutions are the “rules of the game” — formal and informal — that structure incentives, and they fit together into systems that cannot be dismantled piece by piece. The second is **regulatory-design theory** (in the tradition of Cass Sunstein): the choice of policy instrument is itself a design decision with predictable consequences. The third is a disciplined **evaluation framework** — effectiveness, efficiency, equity, legitimacy — applied not impressionistically but with the causal-inference methods you met in *Labour Economics & Inequalities* (BUS33-4): the counterfactual, natural experiments, and difference-in-differences.

The deliverable skill is concrete and professional. By the end you should be able to pick up a government white paper or a regulatory impact assessment, read it as a **political document** — asking what it assumes, whose interests it serves, and which alternatives it quietly excludes — and then write a structured **policy brief** that diagnoses a real problem, lays three options against the four criteria, recommends one and defends it against its own trade-offs, sequences the implementation, and states how you would know later whether it worked.

A word on why this matters for people heading into business rather than government. At S6 this course runs alongside Strategic Allocation, Corporate Finance Leadership, and Sustainable Futures. The thesis that ties it to them is simple: for any firm, the institutional environment is not fixed scenery — it is a **strategic variable**. A company that understands how a rule is made, who shaped it, and how it will be evaluated can comply earlier, shape the next rule, and turn a constraint into an advantage. Treating regulation as someone else’s problem is itself a strategic choice, and usually a losing one.

The order of the book is deliberate. Chapter 1 defines public policy and the policy cycle and maps the actors. Chapter 2 classifies the policy instruments. Chapter 3 turns to institutions and state models — North, complementarity, and why transplants fail. Chapter 4 examines regulatory-agency design and capture, with AI governance as the running case. Chapter 5 builds the four-criteria evaluation on a causal-inference spine. Chapter 6 uses an institutional lens to dissect a named policy failure. Chapter 7 puts the firm inside the regulation. Chapter 8 assembles everything into the policy brief.

Contents

Preface	2
1 Public policy, the policy cycle, and the actors	3
1.1 What public policy is — and is not	3
1.2 The policy cycle	4
1.3 The actors: visible versus real influence	5
2 Policy instruments	6

2.1	The five instrument families	6
2.2	Choosing the right instrument	7
3	Institutions and state models	8
3.1	Institutions in North's sense	8
3.2	State models as institutional architectures	8
3.3	Institutional complementarity and why transplants fail	9
4	Regulatory-agency design and capture	9
4.1	The three design dimensions	10
4.2	Regulatory capture	10
4.3	The running case: the institutional architecture of AI governance	11
5	Four-criteria policy evaluation	12
5.1	The four criteria	12
5.2	The causal-inference spine	12
6	Policy failure analysis with an institutional lens	13
6.1	The diagnostic	13
6.2	Connecting institutional design to cross-pillar decisions	14
7	Companies inside regulation	15
7.1	The compliance-to-shaping spectrum	15
7.2	Lobbying as a legitimate strategic tool — and its limits	15
8	Producing the policy brief	16
8.1	Reading a white paper as a political document	16
8.2	The structure of a policy brief	17

1 Public policy, the policy cycle, and the actors

Start with the thing itself. A government decides that too few homes are being built. Over several years it commissions studies, a minister announces a target, parliament debates and passes a planning law, civil servants write the implementing regulations, local authorities apply them, and — eventually — someone asks whether more homes actually got built. That whole arc, not the law alone, is the **public policy**. The law is one frozen frame; the policy is the film.

1.1 What public policy is — and is not

Definition 1 A **public policy** is a deliberate course of action (or inaction) adopted by public authorities to address a problem of collective concern, pursued through a combination of instruments and sustained over time. It is broader than any single law and is defined by its **intended effect on a problem**, not by its legal form.

Two distinctions do most of the work, and students blur them constantly.

Worked example — law, policy, strategy — three answers to the same housing shortage.

- A **law** is a binding norm: “buildings over 4 storeys require permit X.” It is an **instrument**, often one piece of a policy.
- A **public policy** is the whole problem-directed effort: the housing policy might combine that planning law with tax incentives, public construction subsidies, and an information campaign — aimed at the outcome “more affordable homes.”
- A **corporate strategy** is a **private** actor’s plan to achieve a **private** goal (profit, market share): a developer deciding which cities to build in.

Policy pursues a **public** goal through **public authority** and is accountable to **citizens**; strategy pursues a **private** goal through **private resources** and is accountable to **shareholders**. The same housing shortage looks completely different through each lens.

Common pitfall Do not equate “public policy” with “law.” Much of what a policy does is **non-legal** — spending, persuasion, the design of a default. And a law on the books that changes nothing in the world is a failed policy, however valid it is legally. Judge a policy by its effect on the problem; judge a law by its validity. Conflating the two makes you mistake the passing of a statute for the solving of a problem — the single most common error in policy commentary.

1.2 The policy cycle

Policy is easier to analyse if you break the film into stages. The **policy cycle** is the standard analytical scaffold — idealised (real policy loops, skips, and stalls) but indispensable for locating where an actor has influence and where a policy went wrong.

Definition 2 The **policy cycle** has five stages:

1. **Agenda-setting** — a problem comes to be seen as a public matter requiring action (and many do not).
2. **Formulation** — options are designed and analysed; instruments are chosen.
3. **Adoption** (decision) — an option is formally enacted by those with authority.
4. **Implementation** — the decision is turned into action by administrations and field-level actors.
5. **Evaluation** — the effects are assessed against the goals, feeding back into the agenda.




Figure 1: The five-stage policy cycle. The loop matters: **evaluation** feeds back into **agenda-setting**, so policy is iterative, not a one-shot. The cycle is a map, not a law of nature — real policies stall at adoption, skip evaluation entirely, or never make the agenda at all. Its analytical value is to let you ask, of any actor, “at which stage does this one actually have power?”

Worked example — the cycle running — the DMA/DSA negotiation. The EU’s Digital Markets Act and Digital Services Act trace the cycle cleanly. **Agenda-setting:** years of concern about dominant platforms (amplified by think tanks, press investigations, and competition cases) turned “platform power” into a problem demanding action. **Formulation:** the European Commission drafted the two regulations, choosing *ex ante* obligations on “gatekeepers” (DMA) and due-process duties on platforms (DSA) over the slower case-by-case competition route. **Adoption:** the Parliament and Council negotiated and amended the texts — the stage where lobbying was most intense — and adopted them in 2022. **Implementation:** the Commission designated gatekeepers and very large platforms and began enforcement from 2023–24. **Evaluation:** built-in review clauses will ask whether contestability and safety actually improved. Watch how the centre of gravity of influence **moves** from stage to stage.

1.3 The actors: visible versus real influence

A policy is made by a crowd, not a person, and the crowd is not who the textbook diagram says. The analytical pay-off of this chapter is learning to separate **visible** influence (who formally decides) from **real** influence (who shapes what gets decided).

Definition 3 The principal policy actors:

- **Governments / executives** — set direction, draft, and implement; usually the formal centre of gravity.
- **Parliaments / legislatures** — deliberate, amend, and adopt; confer democratic legitimacy.
- **Regulatory agencies** — specialised bodies that implement and enforce, often with delegated rule-making power (the subject of Chapter 4).

- **Think tanks** — supply ideas, framings, and draft language, shaping the menu of options before any vote.
- **Lobbies / interest groups** — firms, industry associations, NGOs, and unions that press for outcomes, strongest at formulation and adoption.
- **Media** — shape which problems reach the agenda and how they are framed.
- **Civil society** — citizens, movements, and the public whose attention and consent bound what is feasible.




Figure 2: Where each actor’s influence concentrates along the cycle. The lesson is that influence is **stage-specific**: the media and civil society dominate which problems reach the agenda, think tanks and lobbies shape the menu at formulation, and agencies own implementation. An actor invisible at the vote may have decided the outcome earlier by controlling the menu of options.

Common pitfall The most consequential influence is often the **least visible**. By the time a bill is voted, the genuinely contested choices — which options were even put on the table — were settled earlier, at agenda-setting and formulation, by actors who never appear in the official record. When you map a policy, do not stop at “who signed it.” Ask **who set the agenda, who wrote the first draft, and which options never made the menu**. That is where real influence hides.

2 Policy instruments

Once a government has decided to act, it must choose **how**. The set of tools is small and learnable, and — this is the central claim of regulatory-design theory — the choice of tool is not a technicality. It encodes assumptions about why people behave as they do and distributes costs and freedoms in predictable ways.

Worked example — five ways to cut smoking. Suppose the goal is fewer smokers. A government could: **ban** sales to minors (regulation); raise **tobacco taxes** so cigarettes cost more (taxation); fund **free cessation clinics** (public spending); run a graphic **anti-smoking campaign** (information); or require plain packaging and put cigarettes out of sight behind the counter so the easy default is not to buy (a nudge). Each can “work,” but each assumes a different reason people smoke, hits a different group hardest, and meets different political resistance. Choosing among them is the policy.

2.1 The five instrument families

Definition 4 The main families of policy instrument:

- **Regulation** — binding rules backed by sanction (bans, standards, licences, mandates). Direct and forceful; needs monitoring and enforcement capacity.
- **Taxation** (and subsidies) — changing prices to discourage or encourage behaviour (carbon tax, tobacco duty; or tax credits and subsidies on the other side).

- **Public spending** — the state provides or funds something directly (public housing, infrastructure, services).
- **Information** — changing behaviour by changing beliefs (campaigns, mandatory labels, disclosure rules).
- **Nudges** — restructuring the choice environment so the desired option is the default or the easy one, **without** removing freedom of choice (automatic enrolment, default options, salience).

Remark These map roughly onto a classic typology of governing “by the stick, the carrot, and the sermon” — regulation is the stick, taxes and spending are carrots (and sticks), information and nudges are the sermon (and its quiet cousin). Nudges are the newest family, drawn from behavioural economics: they accept that people are predictably irrational and work **with** that, rather than assuming the rational agent who responds to prices and facts.

2.2 Choosing the right instrument

The professional question is not “which tool is best?” but “under which conditions is each tool appropriate?” The answer turns on political, economic, and institutional conditions.

Instrument	Appropriate when...	Weak / risky when...
Regulation	behaviour must stop hard and fast; harm is severe; a clear bright line is needed.	the state cannot monitor or enforce; the rule is easy to evade; one size misfits.
Taxation	behaviour responds to price; you want flexibility in how targets are met; revenue is useful.	demand is price-inelastic; the burden falls regressively; the price signal is politically toxic.
Public spending	the market under-provides a good; direct provision is the only credible route.	fiscal space is tight; the state delivers inefficiently; it crowds out private supply.
Information	people lack facts and would act on them; trust in the source is high.	the problem is not ignorance but incentives or habit; the audience distrusts the messenger.
Nudges	a default or framing can steer choices cheaply, preserving freedom.	the stakes are high and demand a hard rule; the nudge is seen as manipulation.

Table 1: Matching instrument to conditions. The right-hand column matters most: every instrument has a **failure mode**. A tax on an inelastic good raises revenue without changing behaviour; an information campaign against a problem of incentives wastes money; a nudge where the harm is severe is an abdication. Naming the failure mode is the analytical skill.

Common pitfall Three traps recur. First, **instrument fixation** — reaching for regulation (or lately, nudges) reflexively regardless of fit. Second, ignoring **state capacity**: a rule the administration cannot monitor or enforce is theatre, and in weak-capacity settings a simple tax may beat an elegant regulation that exists only on paper. Third, ignoring **distribution**:

a flat carbon tax cuts emissions but hits poor households hardest — the equity criterion of Chapter 5 is in play from the moment you pick the tool. Instruments are usually **combined**, not chosen singly; the art is the mix.

3 Institutions and state models

Why do the same instruments produce wildly different results in different countries? Because instruments operate inside **institutions**, and institutions differ. This chapter is the theoretical heart of the course: Douglass North’s institutional economics, the main state models, and the reason copying a successful policy across borders so often fails.

3.1 Institutions in North’s sense

Worked example — the rules of the game. Two football teams of equal skill will play very differently under different rules — change the offside rule and the whole pattern of play shifts, though no player changed. North’s insight is that economies and societies are like this: the **institutions** (the rules) shape what the **players** (firms, citizens, officials) find worthwhile to do. Two countries with the same resources and technology can diverge for decades because their rules reward different things — investment versus expropriation, building versus rent-seeking.

Definition 5 In Douglass North’s institutional economics, **institutions** are the humanly devised constraints that structure interaction — the “rules of the game.” They are:

- **formal** — constitutions, laws, property rights, contracts; and
- **informal** — norms, conventions, trust, customs.

Institutions shape incentives, and therefore economic and political outcomes; **organisations** (firms, parties, agencies) are the players who act within them. Crucially, institutions are **path-dependent**: today’s rules are constrained by yesterday’s, which is why they change slowly and unevenly.

3.2 State models as institutional architectures

States are not all built the same way. Three reference models capture distinct institutional architectures for how a state relates to the economy.

Definition 6 Three state models:

- **Liberal state** — the state secures property rights, enforces contracts, and polices competition, then largely leaves allocation to markets. Intervention is the exception. (Archetype: the US/UK tradition.)
- **Developmental state** — the state actively steers economic development, picking sectors, coordinating investment, and disciplining firms toward national goals. (Archetype: post-war Japan, South Korea.)
- **Regulatory state** — the state governs less by owning or directing and more by **rule-making and oversight** through specialised agencies; it sets and enforces the rules of markets it does not run. (Archetype: the modern EU.)




Figure 3: The three state models on a spectrum of how directly the state steers the economy. The liberal state intervenes least and the developmental state most; the regulatory state is distinctive not for **how much** it intervenes but for **how** — through arms-length agencies and rules rather than ownership or direction. These are ideal types: real states mix them, and the mix is itself an institutional choice.

3.3 Institutional complementarity and why transplants fail

Here is the course’s most important single idea. Institutions are not a buffet from which you pick the best dish from each cuisine. They fit together into systems, and a piece that works beautifully in one system can fail or backfire in another.

Definition 7 Institutional complementarity: the effectiveness of one institution depends on the presence of others, so institutions form **coherent systems** rather than independent parts. A reform that imports one institution into a system lacking its complements often underperforms or fails — the imported piece had been doing its work **in concert** with institutions left behind.

Worked example — the transplant that rejected — shock therapy without the complements. After 1991 several post-Soviet economies rapidly privatised state firms, expecting the private-ownership institution that drives Western productivity. But private property does its work **in concert with** enforceable contract law, functioning courts, capital markets, and corporate-governance norms. Transplanted into a setting lacking those complements, mass privatisation produced not competitive markets but asset-stripping and oligarchic concentration. The same institution — private ownership — yielded near-opposite results because its complements were absent. Contrast a more sequenced transition (e.g. Poland’s) where institution-building accompanied privatisation.

Common pitfall Beware “best-practice” reasoning: “country X’s independent central bank / flexible labour market / one-stop business registry works, so let us adopt it.” A policy’s success is a property of the **system**, not the part. Before recommending a transplant, ask: **what complements made this work where it came from, and do they exist here?** This is the institutional-complementarity test, and it is the heart of comparative policy analysis. It is also why this course insists every comparative claim be anchored in a **concrete case**, not an abstraction.

4 Regulatory-agency design and capture

The regulatory state governs through **agencies** — central banks, competition authorities, data-protection regulators, financial supervisors. How an agency is designed determines whether it serves the public interest or the industry it oversees. This chapter examines the three design dimensions, the central pathology (capture), and the running case of AI governance.

4.1 The three design dimensions

Worked example — why we tie the regulator's hands — and the central bank. Imagine a central bank that took orders from the finance minister. Before every election the minister would want low interest rates to boost the economy, storing up inflation for later. Knowing this, markets would never trust the currency. So many countries deliberately make the central bank **independent** — insulated from the election cycle — precisely so it can do an unpopular job credibly. But independence raises its own problem: an unelected, unaccountable body wielding huge power. The whole design problem is balancing these.

Definition 8 A regulatory agency is designed along three dimensions:

- **Independence** — insulation from day-to-day political and industry pressure (fixed terms, protected budgets, removal only for cause), so it can act on expertise and the long view.
- **Mandate** — the scope and objective the law assigns it: what it may regulate, with what goals, using what powers. A clear, bounded mandate is the source of its authority and its limit.
- **Accountability** — the mechanisms that answer “who guards the guardian?”: reporting to parliament, judicial review of decisions, transparency, appeal rights, published reasons.




Figure 4: The regulatory-agency design triangle. The three corners are in tension: more independence without matching accountability produces an unanswerable power; a sweeping mandate without independence invites political abuse; accountability mechanisms that bite too hard erode the independence that justified the agency in the first place. Good design holds all three in balance — and the right balance depends on the sector.

4.2 Regulatory capture

Definition 9 **Regulatory capture** occurs when an agency comes to serve the interests of the industry it is meant to regulate rather than the public. It arises through information asymmetry (the industry knows more), the **revolving door** (staff move between agency and industry), concentrated industry incentives versus diffuse public ones, and dependence on the regulated for data and cooperation. Capture can be overt (corruption) or — more commonly and more dangerously — **cognitive**: the regulator sincerely adopts the industry’s worldview.

Worked example — a named case — financial regulation before 2008. In the run-up to the 2008 crisis, financial supervisors in several countries had grown deferential to the

institutions they oversaw: they relied on banks' own risk models, accepted the industry's framing that complex products dispersed rather than concentrated risk, and faced a revolving door between regulator and bank. This was largely **cognitive** capture — few bribes, but a shared worldview that under-rated the danger. The design lesson drove later reforms: stronger mandates (macro- prudential oversight), more independence, and harder accountability.


Common pitfall Capture is rarely a briefcase of cash. The dangerous form is **cognitive**: the regulator, immersed in the industry's data and people, sincerely comes to see the world as the industry does. So “are the regulators honest?” is the wrong test. Ask instead: **does the design create structural distance** — independent expertise, outside data sources, revolving-door cooling-off periods, transparency that lets third parties contest decisions? Capture is an institutional-design failure, not a moral one, and the redesign is institutional too.

4.3 The running case: the institutional architecture of AI governance

The EU AI Act (previewed in HUM24-2) is this course's named case for agency design, because its institutions are being built right now and its accountability gaps are visible and contested.

Definition 10 The AI Act's emerging institutional architecture has several layers:

- the **AI Office** (within the European Commission) — supervises general-purpose AI models and coordinates enforcement EU-wide;
- **national competent authorities** — designated bodies in each member state that supervise and enforce within their territory;
- **cross-border enforcement** mechanisms — coordination so that a model deployed across the single market is governed coherently;
- **algorithmic-accountability redress** — the channels through which an individual harmed by an AI system (a wrongful automated decision) can obtain explanation, contestation, and remedy.



dashed seam = cross-border coordination & the accountability gap

Figure 5: The AI Act's institutional architecture: the EU-level AI Office over national competent authorities, with redress flowing up from citizens and deployers. The red dashed line marks the seam where the hardest problems live — **cross-border consistency** (will 27 authorities enforce alike?) and the **accountability gap**: an individual harmed by an opaque model may struggle to get a meaningful explanation or remedy. Naming a seam like this is the analytical move the course asks for.

Remark Naming a gap is only half the analysis; the other half is proposing **one specific institutional-redesign** that would close it. A worked example for the AI case: the redress layer is weak because an individual rarely has the information or standing to contest an automated decision effectively. A concrete redesign would be a **mandatory, independent algorithmic-audit and redress body** with the power to demand model documentation, conduct audits on complaint, and order remedies — closing the gap between a paper “right to explanation” and an enforceable one. Any good redesign proposal must (a) name the gap, (b) specify the institution, and (c) say how it changes incentives.

5 Four-criteria policy evaluation

A policy that **exists** is not a policy that **works**. Evaluation is the discipline of judging a real policy against explicit criteria using credible evidence. This course uses four criteria — effectiveness, efficiency, equity, legitimacy — and insists that the first two be assessed with the **causal-inference** methods you learned in BUS33-4, not with before-and-after eyeballing.

5.1 The four criteria

Definition 11 The four evaluation criteria:

- **Effectiveness** — did the policy achieve its stated objective? (Did emissions / smoking / unemployment actually fall **because of** it?)
- **Efficiency** — did it achieve the objective at acceptable cost? (Benefits versus costs; was there a cheaper route to the same result?)
- **Equity** — how are the gains and burdens distributed? (Who pays, who benefits; is the distribution fair?)
- **Legitimacy** — was it adopted and implemented through accepted, accountable procedures, and is it seen as rightful by those it binds?

Common pitfall The four criteria routinely **conflict**, and an evaluation that reports only the one flattering to its author’s prior is propaganda. A carbon tax can be effective (emissions fall) and efficient (cheap per tonne) yet inequitable (regressive) and illegitimate (seen as imposed) — the gilets jaunes protests were exactly this collision. The job is not to find a policy that maxes all four but to make the **trade-offs explicit** and defend a judgement about them. A recommendation that hides its trade-offs is not yet a policy analysis.

5.2 The causal-inference spine

The hard part of **effectiveness** is the word “because.” Outcomes change for many reasons; isolating the policy’s own contribution requires answering a counterfactual.

Definition 12 The **counterfactual problem**: to know a policy’s effect we must compare what happened to **what would have happened without it** — a state of the world we never observe. Credible evaluation is the art of estimating that missing counterfactual. Key tools (from BUS33-4):

- **Natural experiment** — exploit a situation where a policy applied to one group and not to a comparable other, as if by chance.

- **Difference-in-differences (DiD)** — compare the **change** in the treated group to the **change** in an untreated control group, differencing out shared trends.

Worked example — did the minimum wage cost jobs? — Card & Krueger, revisited. You met this in BUS33-4. When New Jersey raised its minimum wage and neighbouring Pennsylvania did not, Card and Krueger compared the **change** in fast-food employment on each side of the border. Naïve before-and-after on New Jersey alone would confuse the policy with everything else moving in the economy. Differencing against Pennsylvania — the control — isolated the policy’s own effect, and found no employment drop. That is difference-in-differences doing the evaluator’s core job: building a credible counterfactual.




Figure 6: Difference-in-differences. The treated group’s **actual** path (red) is compared not to its own past but to its **counterfactual** (orange dashed) — what its trend would have been, inferred from the control group’s change (blue). The policy effect is the gap between actual and counterfactual **after** treatment. The method rests on the **parallel-trends assumption**: that absent the policy, treated and control would have moved together. Naming that identifying assumption — and arguing it is plausible — is what separates a causal claim from a correlation.

Common pitfall The seductive error in evaluation is the **before-and-after** comparison: “emissions fell after the tax, so the tax worked.” It ignores everything else that changed (a recession, a technology shift, a warm winter). Always ask: **compared to what counterfactual?** And whenever you assert a causal effect via DiD or a natural experiment, **state the identifying assumption** (parallel trends; as-good-as-random assignment) and say why it is credible here. An effect claim without its identifying assumption is not yet evidence.

6 Policy failure analysis with an institutional lens

Some policies fail, and the institutional lens explains **why** better than blaming the people involved. The diagnostic move is to look past the surface (“bad luck,” “poor execution”) to the **institutional design** that made failure likely: misaligned incentives and governance gaps. The pay-off is a **redesign** — the institutional change that would have prevented it.

6.1 The diagnostic

Definition 13 An **institutional-lens failure analysis** asks, of a named failure:

1. **Misaligned incentives** — who was rewarded for what, and did those rewards pull against the policy’s goal?

2. **Governance gaps** — what was unmonitored, unaccountable, or nobody’s job?
3. **The redesign** — what change to the rules (the incentives or the oversight) would have realigned behaviour with the goal?

The discipline is to locate the failure in the **design**, not (only) in the individuals — because individuals change but a bad design keeps producing failure.

Worked example — a named failure — subprime mortgages and the 2008 crisis.

Misaligned incentives: mortgage originators were paid per loan issued and bore none of the default risk (they sold the loans on), so they were rewarded for volume, not quality; rating agencies were paid by the issuers whose products they rated. **Governance gaps:** securitised risk fell between national regulators; nobody owned systemic risk; supervisors relied on the banks’ own models (the capture of Chapter 4). **Redesign:** “skin in the game” rules forcing originators to retain a slice of risk; a macro-prudential authority explicitly mandated to watch system-wide risk; rating-agency conflict-of-interest reform. Each redesign targets a specific misalignment — that is what makes it a remedy rather than a wish.

Common pitfall Resist two lazy explanations. “Bad people” (greed, incompetence) is rarely the root cause — swap the individuals and a bad design fails again. “Bad luck” (an unforeseeable shock) ignores that good institutions are built to be **robust** to shocks. The institutional analyst asks what **structure** made the failure likely and what **redesign** removes it. A failure analysis that ends in “they should have tried harder” has not done the work.

6.2 Connecting institutional design to cross-pillar decisions

The reason this belongs in a business curriculum: a specific institutional design **constrains or enables** decisions in Strategy, Finance, and Operations. This is the cross-pillar link that makes institutional analysis a business skill, not only a civic one.

Institutional design	What it constrains / enables	Cross-pillar decision
Central-bank independence	credible, stable monetary policy; predictable inflation and rates.	Finance: cost of capital, hedging, the discount rate in any valuation.
Data-protection authority mandate	how aggressively personal data can be used and monetised.	Strategy: data-product design, market entry, M&A due diligence.
Carbon-market architecture	the price and predictability of emitting carbon.	Operations & Strategy: plant siting, technology choice, supply-chain design.

Table 2: Institutions as strategic variables. The lesson: an institutional design is not background — it sets the constraints inside which Strategy, Finance, and Operations decisions are made. A firm that reads the institution can anticipate the constraint instead of being surprised by it. This is the bridge to the next chapter, where the firm becomes an actor **inside** the regulation.

7 Companies inside regulation

Firms are not passive recipients of rules. They occupy a spectrum of postures toward regulation, and the more sophisticated postures turn the institutional environment into a source of competitive advantage. This chapter — the most directly strategic — shows how, and treats lobbying honestly as a legitimate (and bounded) strategic tool.

7.1 The compliance-to-shaping spectrum

Worked example — three carmakers and the same emissions rule. Tighter emissions limits are coming. Firm A waits until the rule is final, then scrambles to comply at maximum cost (reactive compliance). Firm B reads the direction of travel early, retools ahead of the deadline, and is ready when rivals are still scrambling (anticipation). Firm C goes to the table while the rule is **being written**, shares technical data, and helps shape a standard it can already meet — one its less-prepared rivals cannot (rule-building participation). Same rule, three postures, three very different competitive outcomes.

Definition 14 Firms' postures toward regulation, from passive to active:

- **Reactive compliance** — meet the rule once it binds, at whatever cost; treat regulation as an external shock.
- **Anticipation** — read regulatory trends early and adapt ahead of the deadline, converting lead time into advantage.
- **Rule-building participation** — engage in the **formulation** of rules (consultations, standard-setting, expert input, lobbying) to shape the rule itself.




Figure 7: The posture spectrum. Moving rightward, the firm shifts from absorbing rules as shocks to **shaping the environment** in which it competes. The advantage is not only lower compliance cost: a firm that helped write a standard it already meets has raised a barrier its rivals must now climb. The institutional environment has become a competitive variable.

7.2 Lobbying as a legitimate strategic tool — and its limits

Definition 15 **Lobbying** is the organised effort to influence policy decisions by providing information, arguments, and pressure to decision-makers. It is a **legitimate** part of democratic policy-making — legislators genuinely need expertise that firms hold — **provided** it is transparent and bounded by rules (registers, disclosure, conflict-of-interest limits). It becomes illegitimate when it captures the process: secret, one-sided, or substituting private interest for the public one.

Common pitfall Two symmetric errors. The cynical one: “lobbying is just legalised corruption” — this ignores that policy-makers need real information industry possesses, and that NGOs and unions lobby too. The naïve one: “the consultation was open, so the outcome is fair” — this ignores that concentrated industry interests routinely out-resource diffuse

public ones, so an open process can still be lopsided. The honest analyst asks **who was in the room, who was not, and whether the process was transparent enough to contest** — the visible-versus-real-influence test of Chapter 1, applied to the firm’s own conduct.

Remark **Converting constraints into advantage** is the strategic synthesis of the course. Mechanisms: a firm that complies early can market compliance as quality or trust (a privacy-first brand under GDPR); a firm that helped set a standard meets it cheaply while rivals retool; a firm that anticipates a ban exits a doomed product line before its competitors. The constraint is the same for everyone in the industry — the **advantage** goes to whoever reads the institution best. That is why this course treats the institutional environment as a strategic variable, not a cost centre.

8 Producing the policy brief

Everything converges here. The capstone deliverable is a **policy brief**: a structured, evidence-based document that diagnoses a real problem, weighs options against the four criteria, recommends one, and defends it honestly. Producing it has two halves — **reading** the source documents critically, and **writing** the brief.

8.1 Reading a white paper as a political document

Worked example — what the impact assessment does not say. A government’s regulatory impact assessment recommends Option B. Read closely, it **assumes** firms will comply voluntarily (an embedded assumption), its cost-benefit counts industry compliance costs carefully but treats diffuse public benefits vaguely (whose interests are served), and it considers only three options — never a simple tax that would have hurt a politically connected sector (the alternative excluded). None of this is visible if you read it as a neutral technical report. Read as a **political document**, the recommendation’s real architecture appears.

Definition 16 To read a white paper or regulatory impact assessment **as a political document**, interrogate three things:

- **Embedded assumptions** — what must be true for the argument to hold (about behaviour, compliance, costs), and are those assumptions stated or smuggled?
- **Interests served** — who gains and who loses under the recommendation, and whose framing of the problem has been adopted?
- **Alternatives excluded** — which options were never seriously considered, and why might they have been left off the menu?

Common pitfall A government document is **not** a neutral technical report; it is an argument made by actors with interests, however expert and sincere. The naïve reader takes its framing and option set as given — and so re-derives its recommendation. The critical reader reconstructs the **menu of options that was excluded** and the **assumptions that were smuggled in**, because that is where the real decision was made (Chapter 1’s agenda-setting and formulation, seen from the inside of the paper).

8.2 The structure of a policy brief

Definition 17 A complete **policy brief** contains, in order:

1. **Quantified diagnosis** — the problem stated with evidence and numbers, not adjectives: how big, who is affected, what trend.
2. **Three options**, each assessed with explicit **cost-benefit against the four criteria** (effectiveness, efficiency, equity, legitimacy).
3. **Reasoned recommendation** — one option chosen, with the reasoning shown, including **why** it beats the others on the criteria that matter most here.
4. **Sequenced implementation plan** — who does what, in what order, with what resources.
5. **Monitoring indicators** — the metrics that will show, during implementation, whether it is on track.
6. **Ex post evaluation criteria** — how, and against what counterfactual, success will be judged **after** the fact.
7. **Honest account of limitations** — the trade-offs accepted, the assumptions relied on, and what could go wrong.




Figure 8: The policy-brief structure as a pipeline from diagnosis to ex post evaluation. The spine of the document is the move from step 1 to step 3: a **quantified** diagnosis feeds **three options scored on the four criteria**, which forces the recommendation to be **defended against explicit trade-offs** rather than asserted. Steps 4–6 make it actionable and — crucially — **checkable later**, closing the policy cycle of Chapter 1 by building evaluation in from the start.

Common pitfall The brief's two most common failures are **advocacy** and **evasion**. Advocacy presents one option dressed up, with two straw alternatives — the reader sees the rigging. Evasion recommends nothing clearly, hiding behind “more research is needed.” A real brief commits to a recommendation **and** states honestly what it costs and how you would know if you were wrong. The honest account of limitations is not weakness; it is the mark of an analyst the decision-maker can trust. Remember the brief is delivered before a jury including a practitioner — vagueness and one-sidedness are exactly what they are trained to catch.

The through-line of the module. You moved from defining **public policy** and the **policy cycle**, learning to separate visible from real influence among the actors (Ch. 1); to

the **instruments** and the conditions under which each fits (Ch. 2); to the **institutions** behind the instruments — North’s rules of the game, the liberal / developmental / regulatory state models, and the institutional-complementarity reason transplants fail (Ch. 3). On that base you analysed **regulatory-agency design** — independence, mandate, accountability, the pathology of capture, and the live case of AI governance (Ch. 4); built a **four-criteria evaluation** on a **causal-inference** spine, refusing before-and-after for the counterfactual (Ch. 5); used an **institutional lens** to diagnose a named **failure** and connect institutional design to cross-pillar decisions (Ch. 6); placed the **firm inside the regulation**, from reactive compliance to rule-building and constraint-to-advantage (Ch. 7); and finally assembled the **policy brief**, reading sources as political documents and defending a recommendation against its trade-offs (Ch. 8). The discipline is constant: read policy as a political artifact, respect the counterfactual, make trade-offs explicit, and treat the institutional environment as a strategic variable.